

# 06linearregression

July 14, 2019

```
In [2]: import warnings
        warnings.filterwarnings("ignore")
        from sklearn.datasets import load_boston
        from random import seed
        from random import randrange
        from csv import reader
        from math import sqrt
        from sklearn import preprocessing
        import pandas as pd
        import numpy as np
        import matplotlib.pyplot as plt
        from prettytable import PrettyTable
        from sklearn import preprocessing
        from sklearn.metrics import mean_squared_error
        import seaborn as sns
        from sklearn.linear_model import SGDRegressor
        from sklearn.linear_model import LinearRegression
        import sklearn.cross_validation
```

```
In [24]: boston = load_boston()
         print(boston.data.shape)
         print(boston.target.shape)
```

```
(506, 13)
(506,)
```

```
In [4]: print(boston.feature_names)
```

```
['CRIM' 'ZN' 'INDUS' 'CHAS' 'NOX' 'RM' 'AGE' 'DIS' 'RAD' 'TAX' 'PTRATIO'
 'B' 'LSTAT']
```

```
In [74]: X = boston.data
         Y = boston.target
         print(X)
         print("--"*50)
         print(Y)
```

```
[6.3200e-03 1.8000e+01 2.3100e+00 ... 1.5300e+01 3.9690e+02 4.9800e+00]
[2.7310e-02 0.0000e+00 7.0700e+00 ... 1.7800e+01 3.9690e+02 9.1400e+00]
[2.7290e-02 0.0000e+00 7.0700e+00 ... 1.7800e+01 3.9283e+02 4.0300e+00]
...
[6.0760e-02 0.0000e+00 1.1930e+01 ... 2.1000e+01 3.9690e+02 5.6400e+00]
[1.0959e-01 0.0000e+00 1.1930e+01 ... 2.1000e+01 3.9345e+02 6.4800e+00]
[4.7410e-02 0.0000e+00 1.1930e+01 ... 2.1000e+01 3.9690e+02 7.8800e+00]]
```

---

```
[24. 21.6 34.7 33.4 36.2 28.7 22.9 27.1 16.5 18.9 15. 18.9 21.7 20.4
18.2 19.9 23.1 17.5 20.2 18.2 13.6 19.6 15.2 14.5 15.6 13.9 16.6 14.8
18.4 21. 12.7 14.5 13.2 13.1 13.5 18.9 20. 21. 24.7 30.8 34.9 26.6
25.3 24.7 21.2 19.3 20. 16.6 14.4 19.4 19.7 20.5 25. 23.4 18.9 35.4
24.7 31.6 23.3 19.6 18.7 16. 22.2 25. 33. 23.5 19.4 22. 17.4 20.9
24.2 21.7 22.8 23.4 24.1 21.4 20. 20.8 21.2 20.3 28. 23.9 24.8 22.9
23.9 26.6 22.5 22.2 23.6 28.7 22.6 22. 22.9 25. 20.6 28.4 21.4 38.7
43.8 33.2 27.5 26.5 18.6 19.3 20.1 19.5 19.5 20.4 19.8 19.4 21.7 22.8
18.8 18.7 18.5 18.3 21.2 19.2 20.4 19.3 22. 20.3 20.5 17.3 18.8 21.4
15.7 16.2 18. 14.3 19.2 19.6 23. 18.4 15.6 18.1 17.4 17.1 13.3 17.8
14. 14.4 13.4 15.6 11.8 13.8 15.6 14.6 17.8 15.4 21.5 19.6 15.3 19.4
17. 15.6 13.1 41.3 24.3 23.3 27. 50. 50. 50. 22.7 25. 50. 23.8
23.8 22.3 17.4 19.1 23.1 23.6 22.6 29.4 23.2 24.6 29.9 37.2 39.8 36.2
37.9 32.5 26.4 29.6 50. 32. 29.8 34.9 37. 30.5 36.4 31.1 29.1 50.
33.3 30.3 34.6 34.9 32.9 24.1 42.3 48.5 50. 22.6 24.4 22.5 24.4 20.
21.7 19.3 22.4 28.1 23.7 25. 23.3 28.7 21.5 23. 26.7 21.7 27.5 30.1
44.8 50. 37.6 31.6 46.7 31.5 24.3 31.7 41.7 48.3 29. 24. 25.1 31.5
23.7 23.3 22. 20.1 22.2 23.7 17.6 18.5 24.3 20.5 24.5 26.2 24.4 24.8
29.6 42.8 21.9 20.9 44. 50. 36. 30.1 33.8 43.1 48.8 31. 36.5 22.8
30.7 50. 43.5 20.7 21.1 25.2 24.4 35.2 32.4 32. 33.2 33.1 29.1 35.1
45.4 35.4 46. 50. 32.2 22. 20.1 23.2 22.3 24.8 28.5 37.3 27.9 23.9
21.7 28.6 27.1 20.3 22.5 29. 24.8 22. 26.4 33.1 36.1 28.4 33.4 28.2
22.8 20.3 16.1 22.1 19.4 21.6 23.8 16.2 17.8 19.8 23.1 21. 23.8 23.1
20.4 18.5 25. 24.6 23. 22.2 19.3 22.6 19.8 17.1 19.4 22.2 20.7 21.1
19.5 18.5 20.6 19. 18.7 32.7 16.5 23.9 31.2 17.5 17.2 23.1 24.5 26.6
22.9 24.1 18.6 30.1 18.2 20.6 17.8 21.7 22.7 22.6 25. 19.9 20.8 16.8
21.9 27.5 21.9 23.1 50. 50. 50. 50. 50. 13.8 13.8 15. 13.9 13.3
13.1 10.2 10.4 10.9 11.3 12.3 8.8 7.2 10.5 7.4 10.2 11.5 15.1 23.2
9.7 13.8 12.7 13.1 12.5 8.5 5. 6.3 5.6 7.2 12.1 8.3 8.5 5.
11.9 27.9 17.2 27.5 15. 17.2 17.9 16.3 7. 7.2 7.5 10.4 8.8 8.4
16.7 14.2 20.8 13.4 11.7 8.3 10.2 10.9 11. 9.5 14.5 14.1 16.1 14.3
11.7 13.4 9.6 8.7 8.4 12.8 10.5 17.1 18.4 15.4 10.8 11.8 14.9 12.6
14.1 13. 13.4 15.2 16.1 17.8 14.9 14.1 12.7 13.5 14.9 20. 16.4 17.7
19.5 20.2 21.4 19.9 19. 19.1 19.1 20.1 19.9 19.6 23.2 29.8 13.8 13.3
16.7 12. 14.6 21.4 23. 23.7 25. 21.8 20.6 21.2 19.1 20.6 15.2 7.
8.1 13.6 20.1 21.8 24.5 23.1 19.7 18.3 21.2 17.5 16.8 22.4 20.6 23.9
22. 11.9]
```

```
In [75]: df=pd.DataFrame(X)
```

```
X=df.as_matrix()
df.head()
```

```
Out [75]:
```

	0	1	2	3	4	5	6	7	8	9	10	\
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	

	11	12
0	396.90	4.98
1	396.90	9.14
2	392.83	4.03
3	394.63	2.94
4	396.90	5.33

```
In [76]: X_train, X_test, Y_train, Y_test = sklearn.cross_validation.train_test_split(X, Y, test_size=0.2)
print(X_train.shape)
print(X_test.shape)
print(Y_train.shape)
print(Y_test.shape)
```

```
(354, 13)
(152, 13)
(354,)
(152,)
```

```
In [77]: scaler = preprocessing.StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

```
In [78]: df_train=pd.DataFrame(X_train)
df_train['PRICE']=Y_train

df_train.head()
```

```
Out [78]:
```

	0	1	2	3	4	5	6	\
0	-0.334330	-0.499618	1.069608	-0.251124	1.645428	0.233772	0.969882	
1	0.478536	-0.499618	1.069608	-0.251124	1.113435	-0.149715	0.383159	
2	0.277128	-0.499618	1.069608	-0.251124	-0.168580	0.653301	0.270733	
3	-0.414300	3.445319	-1.442682	-0.251124	-1.293614	1.372699	-1.591321	
4	-0.397576	-0.499618	2.504352	-0.251124	0.502952	-1.215116	0.896102	

	7	8	9	10	11	12	PRICE
0	-0.900522	1.654486	1.538813	0.810913	-3.463820	1.611369	11.8
1	-0.926152	1.654486	1.538813	0.810913	-2.872888	1.265636	11.0
2	-0.241993	1.654486	1.538813	0.810913	0.389957	-0.671032	23.7

```
3  2.387078 -0.527917 -1.061095 -0.265106  0.421447 -1.082820  35.4
4 -0.982361 -0.642780  1.804713  0.764129  0.412198  0.779361  15.2
```

In [79]: *# code source: [https://medium.com/@haydar\\_ai/learning-data-science-day-9-linear-regres](https://medium.com/@haydar_ai/learning-data-science-day-9-linear-regres)*

```
lm = LinearRegression()
lm.fit(X_train, Y_train)


Y_pred = lm.predict(X_test)

plt.scatter(Y_test, Y_pred)
plt.xlabel("Prices:  $Y_i$ ")
plt.ylabel("Predicted prices:  $\hat{Y}_i$ ")
plt.title("Prices vs Predicted prices:  $Y_i$  vs  $\hat{Y}_i$ ")
plt.show()
```




```
In [80]: delta_y = Y_test - Y_pred;

sns.set_style('whitegrid')
sns.kdeplot(np.array(delta_y), bw=0.5)
plt.show()
```



```
In [81]: sns.set_style('whitegrid')
sns.kdeplot(np.array(Y_pred), bw=0.5)
plt.show()
```



```
In [82]: #SGD implementation for linear regression
```

```
W,B,iteration,rate,k = np.zeros(shape=(1,13)),0,700,0.1,25 #intialise W and B to zero
```

```
while iteration>=0 :
```

```
    w,b,delta_w,delta_b = W,B,np.zeros(shape=(1,13)),0
```

```
    data = df_train.sample(25) #sampling random k=batch size=20 data
```

```
    x = np.array(data.drop('PRICE',axis=1))
```

```
    y = np.array(data['PRICE'])
```

```
    for i in range(25):
```

```
        delta_w += (-2)*x[i]*(y[i]-(np.dot(w,x[i])+b)) #partial differentiation wrt
```

```
        delta_b += (-2)*(y[i]-(np.dot(w,x[i])+b)) #partial differentiation wrt
```

```
    W=(w-rate*(delta_w)/k)
```

```
    B=(b-rate*(delta_b)/k)
```

```
    rate = rate/1.01
```

```
    iteration-=1
```

```
print(W)
```

```
print(B)
```

```
[[ -1.2487875   0.87689079 -0.45965487  0.20235369 -1.42100563  2.61437601
  -0.22784103 -2.70649964  2.22489369 -1.38188117 -2.10348261  1.15048756
  -3.46084742]]
[22.50079696]
```

```
In [83]: #prediction on x_test
```

```
    #https://www.geeksforgeeks.org/numpy-asscalar-in-python/
```

```
    y_predic=[]
```

```
    for i in range(len(X_test)):
```

```
        val=np.dot(W,X_test[i])+B #val= wTx+b
```

```
        y_predic.append(np.asscalar(val))
```

```
In [84]: #Scatter plot of actual price vs predicted price
```


```
plt.scatter(Y_test,y_predic)
```

```
plt.xlabel('Actual price')
```

```
plt.ylabel('Predictd price')
```

```
plt.title('Actual price vs Predicted price')
```

```
plt.show()
```



```
In [85]: MSE_lr=mean_squared_error(Y_test,y_predic)
         print('mean squared error =',MSE_lr)
```

```
mean squared error = 30.78503402128002
```

```
In [86]: #SGD regression sklearn implementation
```

```
model=SGDRegressor(learning_rate='constant',eta0=0.01,penalty=None,n_iter=100,max_iter=1000)
model.fit(X_train,Y_train)
Y_pred_sgd=model.predict(X_test)
```

```
#Scatter plot of actual price vs predicted price
```

```
plt.scatter(Y_test,Y_pred_sgd)
plt.xlabel('Actual price')
plt.ylabel('Predictd price')
plt.title('Actual price vs Predicted price')
plt.show()
```



```
In [89]: MSE_sgd=mean_squared_error(Y_test,Y_pred_sgd)
         print('mean squared error =',MSE_sgd)
```

mean squared error = 29.958075584142815

```
In [88]: #Comparison between weights obtained from own implementation and weights obtained from
         from prettytable import PrettyTable
         x = PrettyTable()
         x.field_names=['Weight vector manual','Weight vector SGD sklearn']
         weight_sgd=model.coef_
         for i in range(13):
             x.add_row([W[0][i],weight_sgd[i]])
         print(x)
```

Weight vector manual	Weight vector SGD sklearn
-1.2487875005565663	-1.367634066486797
0.8768907910707392	0.6232126548856614
-0.45965487392711285	-0.034321602974485256
0.20235368792480632	0.40482018147864196
-1.4210056337505508	-1.4762829533157247
2.6143760112285666	3.3679058605944086
-0.22784103355767868	-0.22049142711468647



-2.7064996405269692	-3.008014875734558	
2.2248936947136384	2.912155404513144	
-1.3818811742142798	-2.1673612468913763	
-2.1034826093058006	-2.3343970840414117	
1.1504875571567152	0.9213535450650273	
-3.4608474150234865	-3.7125274213122017	
+-----+		

### 0.0.1 Conclusions

Weight vectors and MSE values are almost same in Custom SGD implementation and Sklearn implementation