

Keras On MNIST

In [1]:

```
# if you keras is not using tensorflow as backend set "KERAS_BACKEND=tensorflow" use this command
%matplotlib inline
import warnings
warnings.filterwarnings("ignore")

from keras.utils import np_utils
from keras.datasets import mnist
import seaborn as sns
from keras.initializers import RandomNormal
```

Using TensorFlow backend.

In [66]:

```
from keras.layers.normalization import BatchNormalization
from keras.layers import Dropout
from prettytable import PrettyTable
```

In [8]:

```
%matplotlib notebook
import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    fig.canvas.draw()
```

In [9]:

```
# the data, shuffled and split between train and test sets
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

In [10]:

```
print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d, %d)"%(X_train.shape[1], X_train.shape[2]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d, %d)"%(X_test.shape[1], X_test.shape[2]))
```

Number of training examples : 60000 and each image is of shape (28, 28)
Number of training examples : 10000 and each image is of shape (28, 28)

In [11]:

```
# if you observe the input shape its 2 dimensional vector
# for each image we have a (28*28) vector
# we will convert the (28*28) vector into single dimensional vector of 1 * 784

X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])
```

In [12]:

```
# after converting the input images from 3d to 2d vectors
```

```
print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d)"%(X_train.shape[1]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d)"%(X_test.shape[1]))
```

Number of training examples : 60000 and each image is of shape (784)
 Number of training examples : 10000 and each image is of shape (784)

In [13]:

```
# An example data point
print(X_train[0])
```

```
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  3  18  18  18 126 136 175  26 166 255
247 127  0  0  0  0  0  0  0  0  0  0  0  0  30  36  94 154
170 253 253 253 253 253 225 172 253 242 195  64  0  0  0  0  0  0
  0  0  0  0  0  49 238 253 253 253 253 253 253 253 251  93  82
 82  56  39  0  0  0  0  0  0  0  0  0  0  0  18 219 253
253 253 253 253 198 182 247 241  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  80 156 107 253 253 205  11  0  43 154
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0 14  1 154 253  90  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0 139 253 190  2  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0 11 190 253  70  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  35 241
225 160 108  1  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  81 240 253 253 119  25  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  45 186 253 253 150  27  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  16  93 252 253 187
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  249 253 249  64  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  46 130 183 253
253 207  2  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  39 148 229 253 253 253 250 182  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  24 114 221 253 253 253
253 201  78  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  23  66 213 253 253 253 253 198  81  2  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  18 171 219 253 253 253 253 195
 80  9  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 55 172 226 253 253 253 244 133  11  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0 136 253 253 253 212 135 132  16
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0]
```

In [14]:

```
# if we observe the above matrix each cell is having a value between 0-255
# before we move to apply machine learning algorithms lets try to normalize the data
# X => (X - Xmin)/(Xmax-Xmin) = X/255

X_train = X_train/255
X_test = X_test/255
```

In [15]:

```
# example data point after normlizing
print(X_train[0])
```

[illegible]

2 Layered Architecture

Relu + Adam

In [17]:

```
from keras.models import Sequential
from keras.layers import Dense, Activation
```

In [18]:

```
# some model parameters

output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128
nb_epoch = 30
```

In [25]:

```
# start building a model
model_relu_2 = Sequential()
model_relu_2.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.074, seed=None)))
model_relu_2.add(Dense(212, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.097, seed=None)) )
model_relu_2.add(Dense(output_dim, activation='softmax'))

print(model_relu_2.summary())
```

Layer (type)	Output Shape	Param #
dense_10 (Dense)	(None, 364)	285740
dense_11 (Dense)	(None, 212)	77380
dense_12 (Dense)	(None, 10)	2130
Total params: 365,250		
Trainable params: 365,250		
Non-trainable params: 0		
None		

In [26]:

```
model_relu_2.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu_2.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/30
60000/60000 [=====] - 6s 102us/step - loss: 0.2389 - acc: 0.9287 - val_loss: 0.1275 - val_acc: 0.9625
Epoch 2/30
60000/60000 [=====] - 5s 91us/step - loss: 0.0899 - acc: 0.9726 - val_loss: 0.0824 - val_acc: 0.9746
Epoch 3/30
60000/60000 [=====] - 5s 91us/step - loss: 0.0589 - acc: 0.9822 - val_loss: 0.0706 - val_acc: 0.9772
Epoch 4/30
60000/60000 [=====] - 5s 87us/step - loss: 0.0386 - acc: 0.9880 - val_loss: 0.0718 - val_acc: 0.9787
Epoch 5/30
60000/60000 [=====] - 6s 96us/step - loss: 0.0286 - acc: 0.9908 - val_loss: 0.0853 - val_acc: 0.9759
Epoch 6/30
60000/60000 [=====] - 5s 90us/step - loss: 0.0224 - acc: 0.9929 -
```

```
val_loss: 0.0732 - val_acc: 0.9787
Epoch 7/30
60000/60000 [=====] - 5s 90us/step - loss: 0.0166 - acc: 0.9950 -
val_loss: 0.0829 - val_acc: 0.9759
Epoch 8/30
60000/60000 [=====] - 6s 94us/step - loss: 0.0127 - acc: 0.9960 -
val_loss: 0.0957 - val_acc: 0.9752
Epoch 9/30
60000/60000 [=====] - 7s 110us/step - loss: 0.0180 - acc: 0.9939 -
val_loss: 0.0787 - val_acc: 0.9791
Epoch 10/30
60000/60000 [=====] - 6s 98us/step - loss: 0.0121 - acc: 0.9959 -
val_loss: 0.0784 - val_acc: 0.9804
Epoch 11/30
60000/60000 [=====] - 6s 97us/step - loss: 0.0090 - acc: 0.9970 -
val_loss: 0.0804 - val_acc: 0.9799
Epoch 12/30
60000/60000 [=====] - 6s 99us/step - loss: 0.0099 - acc: 0.9968 -
val_loss: 0.0889 - val_acc: 0.9788
Epoch 13/30
60000/60000 [=====] - 6s 97us/step - loss: 0.0110 - acc: 0.9960 -
val_loss: 0.0941 - val_acc: 0.9776
Epoch 14/30
60000/60000 [=====] - 6s 105us/step - loss: 0.0105 - acc: 0.9964 -
val_loss: 0.0901 - val_acc: 0.9785
Epoch 15/30
60000/60000 [=====] - 6s 93us/step - loss: 0.0071 - acc: 0.9978 -
val_loss: 0.0992 - val_acc: 0.9787
Epoch 16/30
60000/60000 [=====] - 6s 93us/step - loss: 0.0070 - acc: 0.9979 -
val_loss: 0.0950 - val_acc: 0.9774
Epoch 17/30
60000/60000 [=====] - 5s 91us/step - loss: 0.0078 - acc: 0.9975 -
val_loss: 0.0826 - val_acc: 0.9829
Epoch 18/30
60000/60000 [=====] - 5s 91us/step - loss: 0.0053 - acc: 0.9982 -
val_loss: 0.1057 - val_acc: 0.9796
Epoch 19/30
60000/60000 [=====] - 6s 102us/step - loss: 0.0100 - acc: 0.9967 -
val_loss: 0.0944 - val_acc: 0.9796
Epoch 20/30
60000/60000 [=====] - 6s 94us/step - loss: 0.0074 - acc: 0.9974 -
val_loss: 0.0907 - val_acc: 0.9818
Epoch 21/30
60000/60000 [=====] - 5s 91us/step - loss: 0.0047 - acc: 0.9984 -
val_loss: 0.0971 - val_acc: 0.9818
Epoch 22/30
60000/60000 [=====] - 5s 91us/step - loss: 0.0065 - acc: 0.9977 -
val_loss: 0.0913 - val_acc: 0.9824
Epoch 23/30
60000/60000 [=====] - 5s 91us/step - loss: 0.0091 - acc: 0.9972 -
val_loss: 0.0844 - val_acc: 0.9822
Epoch 24/30
60000/60000 [=====] - 5s 91us/step - loss: 0.0048 - acc: 0.9985 -
val_loss: 0.0964 - val_acc: 0.9818
Epoch 25/30
60000/60000 [=====] - 5s 91us/step - loss: 0.0052 - acc: 0.9983 -
val_loss: 0.0951 - val_acc: 0.9819
Epoch 26/30
60000/60000 [=====] - 5s 91us/step - loss: 0.0071 - acc: 0.9979 -
val_loss: 0.1003 - val_acc: 0.9788
Epoch 27/30
60000/60000 [=====] - 5s 91us/step - loss: 0.0055 - acc: 0.9982 -
val_loss: 0.1001 - val_acc: 0.9821
Epoch 28/30
60000/60000 [=====] - 5s 91us/step - loss: 0.0011 - acc: 0.9996 -
val_loss: 0.0857 - val_acc: 0.9832
Epoch 29/30
60000/60000 [=====] - 6s 93us/step - loss: 0.0064 - acc: 0.9980 -
val_loss: 0.1391 - val_acc: 0.9754
Epoch 30/30
60000/60000 [=====] - 6s 105us/step - loss: 0.0101 - acc: 0.9970 -
val_loss: 0.0956 - val_acc: 0.9817
```

```

score_relu_2 = model_relu_2.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score_relu_2[0])
print('Test accuracy:', score_relu_2[1])


fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.09564156731741441
Test accuracy: 0.9817



In [29]:

```

w_after = model_relu_2.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



Batch normalization on hidden layers + Relu + Adam

In [30]:

```
model_batch_2 = Sequential()

model_batch_2.add(Dense(532, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.061, seed=None)))
model_batch_2.add(BatchNormalization())

model_batch_2.add(Dense(384, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.072, seed=None)))
model_batch_2.add(BatchNormalization())

model_batch_2.add(Dense(output_dim, activation='softmax'))

model_batch_2.summary()
```

Layer (type)	Output Shape	Param #
dense_13 (Dense)	(None, 532)	417620
batch_normalization_7 (Batch Normalization)	(None, 532)	2128
dense_14 (Dense)	(None, 384)	204672
batch_normalization_8 (Batch Normalization)	(None, 384)	1536
dense_15 (Dense)	(None, 10)	3850
Total params: 629,806		
Trainable params: 627,974		
Non-trainable params: 1,832		

In [31]:

```
model_batch_2.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_batch_2.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```


Train on 60000 samples, validate on 10000 samples

Epoch 1/30

60000/60000 [=====] - 11s 189us/step - loss: 0.1812 - acc: 0.9441 - val_loss: 0.1002 - val_acc: 0.9688

Epoch 2/30

60000/60000 [=====] - 10s 164us/step - loss: 0.0654 - acc: 0.9801 - val_loss: 0.0824 - val_acc: 0.9750

Epoch 3/30

60000/60000 [=====] - 10s 159us/step - loss: 0.0439 - acc: 0.9862 - val_loss: 0.0861 - val_acc: 0.9724

Epoch 4/30

60000/60000 [=====] - 10s 164us/step - loss: 0.0322 - acc: 0.9898 - val_loss: 0.0849 - val_acc: 0.9738

Epoch 5/30

60000/60000 [=====] - 10s 164us/step - loss: 0.0228 - acc: 0.9928 - val_loss: 0.0911 - val_acc: 0.9728

Epoch 6/30

60000/60000 [=====] - 11s 179us/step - loss: 0.0197 - acc: 0.9933 - val_loss: 0.0681 - val_acc: 0.9807

Epoch 7/30

60000/60000 [=====] - 10s 164us/step - loss: 0.0199 - acc: 0.9936 - val_loss: 0.0902 - val_acc: 0.9789

Epoch 8/30

60000/60000 [=====] - 10s 161us/step - loss: 0.0161 - acc: 0.9949 - val_loss: 0.0852 - val_acc: 0.9764

Epoch 9/30

60000/60000 [=====] - 10s 161us/step - loss: 0.0150 - acc: 0.9950 - val_loss: 0.0714 - val_acc: 0.9814

Epoch 10/30

60000/60000 [=====] - 10s 163us/step - loss: 0.0152 - acc: 0.9948 - val_loss: 0.0916 - val_acc: 0.9765

Epoch 11/30

60000/60000 [=====] - 10s 173us/step - loss: 0.0149 - acc: 0.9947 - val_loss: 0.0905 - val_acc: 0.9782

Epoch 12/30

60000/60000 [=====] - 11s 176us/step - loss: 0.0095 - acc: 0.9969 - val_loss: 0.0769 - val_acc: 0.9804

Epoch 13/30

60000/60000 [=====] - 10s 167us/step - loss: 0.0116 - acc: 0.9959 - val_loss: 0.0948 - val_acc: 0.9754

Epoch 14/30

60000/60000 [=====] - 10s 164us/step - loss: 0.0134 - acc: 0.9956 - val_loss: 0.0922 - val_acc: 0.9777

Epoch 15/30

60000/60000 [=====] - 10s 165us/step - loss: 0.0102 - acc: 0.9962 - val_loss: 0.0774 - val_acc: 0.9818

Epoch 16/30

60000/60000 [=====] - 10s 166us/step - loss: 0.0095 - acc: 0.9970 - val_loss: 0.0856 - val_acc: 0.9800

Epoch 17/30

60000/60000 [=====] - 10s 166us/step - loss: 0.0087 - acc: 0.9970 - val_loss: 0.0781 - val_acc: 0.9815

Epoch 18/30

60000/60000 [=====] - 11s 179us/step - loss: 0.0083 - acc: 0.9973 - val_loss: 0.0858 - val_acc: 0.9808

Epoch 19/30

60000/60000 [=====] - 10s 162us/step - loss: 0.0072 - acc: 0.9976 - val_loss: 0.0760 - val_acc: 0.9822

Epoch 20/30

60000/60000 [=====] - 10s 164us/step - loss: 0.0078 - acc: 0.9973 - val_loss: 0.0833 - val_acc: 0.9823

Epoch 21/30

60000/60000 [=====] - 10s 165us/step - loss: 0.0072 - acc: 0.9978 - val_loss: 0.0905 - val_acc: 0.9803

Epoch 22/30

60000/60000 [=====] - 10s 162us/step - loss: 0.0074 - acc: 0.9975 - val_loss: 0.0815 - val_acc: 0.9813

Epoch 23/30

60000/60000 [=====] - 10s 165us/step - loss: 0.0044 - acc: 0.9985 - val_loss: 0.0792 - val_acc: 0.9809

Epoch 24/30

60000/60000 [=====] - 11s 176us/step - loss: 0.0063 - acc: 0.9979 - val_loss: 0.0859 - val_acc: 0.9808

Epoch 25/30

60000/60000 [=====] - 10s 167us/step - loss: 0.0054 - acc: 0.9984 - val_loss: 0.0796 - val_acc: 0.9820

```

Epoch 26/30
60000/60000 [=====] - 10s 165us/step - loss: 0.0052 - acc: 0.9985 - val_loss: 0.0906 - val_acc: 0.9824
Epoch 27/30
60000/60000 [=====] - 10s 168us/step - loss: 0.0094 - acc: 0.9969 - val_loss: 0.0905 - val_acc: 0.9814
Epoch 28/30
60000/60000 [=====] - 10s 162us/step - loss: 0.0054 - acc: 0.9983 - val_loss: 0.0793 - val_acc: 0.9825
Epoch 29/30
60000/60000 [=====] - 10s 164us/step - loss: 0.0051 - acc: 0.9984 - val_loss: 0.0840 - val_acc: 0.9813
Epoch 30/30
60000/60000 [=====] - 11s 181us/step - loss: 0.0038 - acc: 0.9987 - val_loss: 0.0907 - val_acc: 0.9804

```

In [32]:

```

score_batch_2 = model_batch_2.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score_batch_2[0])
print('Test accuracy:', score_batch_2[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))


vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

```

Test score: 0.09071999986333348
Test accuracy: 0.9804

```



In [33]:

```

w_after = model_batch_2.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

```


```

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



Dropout + Relu + Adam

In [34]:

```

model_drop_2 = Sequential()

model_drop_2.add(Dense(532, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.061, seed=None)))
model_drop_2.add(BatchNormalization())
model_drop_2.add(Dropout(0.5))

model_drop_2.add(Dense(384, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.072, seed=None)))
model_drop_2.add(BatchNormalization())
model_drop_2.add(Dropout(0.5))

model_drop_2.add(Dense(output_dim, activation='softmax'))

model_drop_2.summary()

```

Layer (type)	Output Shape	Param #
dense_16 (Dense)	(None, 532)	417620

batch_normalization_9 (Batch Normalization)	(None, 532)	2128
dropout_7 (Dropout)	(None, 532)	0
dense_17 (Dense)	(None, 384)	204672
batch_normalization_10 (Batch Normalization)	(None, 384)	1536
dropout_8 (Dropout)	(None, 384)	0
dense_18 (Dense)	(None, 10)	3850
=====		
Total params: 629,806		
Trainable params: 627,974		
Non-trainable params: 1,832		

In [35]:

```
model_drop_2.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_drop_2.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/30
60000/60000 [=====] - 13s 215us/step - loss: 0.4029 - acc: 0.8781 - val_loss: 0.1331 - val_acc: 0.9583
Epoch 2/30
60000/60000 [=====] - 11s 182us/step - loss: 0.1890 - acc: 0.9423 - val_loss: 0.0954 - val_acc: 0.9706
Epoch 3/30
60000/60000 [=====] - 11s 181us/step - loss: 0.1462 - acc: 0.9547 - val_loss: 0.0829 - val_acc: 0.9731
Epoch 4/30
60000/60000 [=====] - 12s 194us/step - loss: 0.1256 - acc: 0.9605 - val_loss: 0.0772 - val_acc: 0.9757
Epoch 5/30
60000/60000 [=====] - 11s 181us/step - loss: 0.1101 - acc: 0.9656 - val_loss: 0.0702 - val_acc: 0.9787
Epoch 6/30
60000/60000 [=====] - 11s 182us/step - loss: 0.0987 - acc: 0.9698 - val_loss: 0.0701 - val_acc: 0.9791
Epoch 7/30
60000/60000 [=====] - 11s 184us/step - loss: 0.0910 - acc: 0.9712 - val_loss: 0.0678 - val_acc: 0.9793
Epoch 8/30
60000/60000 [=====] - 11s 182us/step - loss: 0.0872 - acc: 0.9717 - val_loss: 0.0648 - val_acc: 0.9789
Epoch 9/30
60000/60000 [=====] - 12s 194us/step - loss: 0.0812 - acc: 0.9742 - val_loss: 0.0609 - val_acc: 0.9811
Epoch 10/30
60000/60000 [=====] - 11s 181us/step - loss: 0.0758 - acc: 0.9758 - val_loss: 0.0581 - val_acc: 0.9814
Epoch 11/30
60000/60000 [=====] - 11s 185us/step - loss: 0.0673 - acc: 0.9784 - val_loss: 0.0573 - val_acc: 0.9826
Epoch 12/30
60000/60000 [=====] - 11s 190us/step - loss: 0.0668 - acc: 0.9778 - val_loss: 0.0605 - val_acc: 0.9826
Epoch 13/30
60000/60000 [=====] - 11s 184us/step - loss: 0.0647 - acc: 0.9797 - val_loss: 0.0559 - val_acc: 0.9827
Epoch 14/30
60000/60000 [=====] - 11s 183us/step - loss: 0.0619 - acc: 0.9799 - val_loss: 0.0556 - val_acc: 0.9830
Epoch 15/30
60000/60000 [=====] - 12s 197us/step - loss: 0.0566 - acc: 0.9816 - val_loss: 0.0577 - val_acc: 0.9825
Epoch 16/30
60000/60000 [=====] - 11s 184us/step - loss: 0.0558 - acc: 0.9820 - val_loss: 0.0589 - val_acc: 0.9823
Epoch 17/30
60000/60000 [=====] - 12s 196us/step - loss: 0.0536 - acc: 0.9823 - val_loss: 0.0536 - acc: 0.9823 - val_acc: 0.9823
```

```

oss: 0.0548 - val_acc: 0.9843
Epoch 18/30
60000/60000 [=====] - 11s 182us/step - loss: 0.0505 - acc: 0.9835 - val_l
oss: 0.0547 - val_acc: 0.9843
Epoch 19/30
60000/60000 [=====] - 12s 192us/step - loss: 0.0491 - acc: 0.9842 - val_l
oss: 0.0545 - val_acc: 0.9832
Epoch 20/30
60000/60000 [=====] - 11s 188us/step - loss: 0.0498 - acc: 0.9841 - val_l
oss: 0.0504 - val_acc: 0.9853loss: 0
Epoch 21/30
60000/60000 [=====] - 11s 182us/step - loss: 0.0449 - acc: 0.9849 - val_l
oss: 0.0524 - val_acc: 0.9856
Epoch 22/30
60000/60000 [=====] - 10s 174us/step - loss: 0.0455 - acc: 0.9848 - val_l
oss: 0.0541 - val_acc: 0.9847
Epoch 23/30
60000/60000 [=====] - 11s 186us/step - loss: 0.0422 - acc: 0.9859 - val_l
oss: 0.0566 - val_acc: 0.9833
Epoch 24/30
60000/60000 [=====] - 11s 184us/step - loss: 0.0409 - acc: 0.9862 - val_l
oss: 0.0566 - val_acc: 0.9826
Epoch 25/30
60000/60000 [=====] - 11s 184us/step - loss: 0.0394 - acc: 0.9868 - val_l
oss: 0.0582 - val_acc: 0.9836
Epoch 26/30
60000/60000 [=====] - 11s 179us/step - loss: 0.0379 - acc: 0.9876 - val_l
oss: 0.0568 - val_acc: 0.9842
Epoch 27/30
60000/60000 [=====] - 11s 177us/step - loss: 0.0383 - acc: 0.9871 - val_l
oss: 0.0596 - val_acc: 0.9836
Epoch 28/30
60000/60000 [=====] - 11s 177us/step - loss: 0.0357 - acc: 0.9880 - val_l
oss: 0.0588 - val_acc: 0.9829
Epoch 29/30
60000/60000 [=====] - 10s 174us/step - loss: 0.0355 - acc: 0.9879 - val_l
oss: 0.0546 - val_acc: 0.9841
Epoch 30/30
60000/60000 [=====] - 10s 175us/step - loss: 0.0360 - acc: 0.9883 - val_l
oss: 0.0528 - val_acc: 0.9848

```

In [38]:

```

score_drop_2 = model_drop_2.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score_drop_2[0])
print('Test accuracy:', score_drop_2[1])


fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')


# list of epoch numbers
x = list(range(1,nb_epoch+1))

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.052775146638258594
Test accuracy: 0.9848





In [39]:


```
w_after = model_drop_2.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



3 Layered Architecture

Relu + Adam

In [40]:

```
# start building a model
model_relu_3 = Sequential()
model_relu_3.add(Dense(621, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.056, seed=None)))
model_relu_3.add(Dense(333, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.077, seed=None)))
model_relu_3.add(Dense(161, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.111, seed=None)))
model_relu_3.add(Dense(output_dim, activation='softmax'))

print(model_relu_3.summary())
```

Layer (type)	Output Shape	Param #
dense_19 (Dense)	(None, 621)	487485
dense_20 (Dense)	(None, 333)	207126
dense_21 (Dense)	(None, 161)	53774
dense_22 (Dense)	(None, 10)	1620
Total params: 750,005		
Trainable params: 750,005		
Non-trainable params: 0		
None		

In [41]:

```
model_relu_3.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu_3.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/30
60000/60000 [=====] - 11s 188us/step - loss: 0.2033 - acc: 0.9383 - val_loss: 0.0901 - val_acc: 0.9723
Epoch 2/30
60000/60000 [=====] - 11s 176us/step - loss: 0.0784 - acc: 0.9760 - val_loss: 0.0931 - val_acc: 0.9683
Epoch 3/30
60000/60000 [=====] - 10s 160us/step - loss: 0.0526 - acc: 0.9834 - val_loss: 0.0876 - val_acc: 0.9718
Epoch 4/30
60000/60000 [=====] - 10s 162us/step - loss: 0.0381 - acc: 0.9877 - val_loss: 0.0757 - val_acc: 0.9782
Epoch 5/30
60000/60000 [=====] - 10s 162us/step - loss: 0.0282 - acc: 0.9908 - val_loss: 0.0794 - val_acc: 0.9781
Epoch 6/30
60000/60000 [=====] - 10s 162us/step - loss: 0.0255 - acc: 0.9922 - val_loss: 0.0896 - val_acc: 0.9774
Epoch 7/30
60000/60000 [=====] - 10s 162us/step - loss: 0.0207 - acc: 0.9928 - val_loss: 0.0795 - val_acc: 0.9794
Epoch 8/30
60000/60000 [=====] - 11s 178us/step - loss: 0.0192 - acc: 0.9937 - val_loss: 0.0967 - val_acc: 0.9755
Epoch 9/30
60000/60000 [=====] - 10s 162us/step - loss: 0.0199 - acc: 0.9936 - val_loss: 0.0768 - val_acc: 0.9807
Epoch 10/30
60000/60000 [=====] - 10s 163us/step - loss: 0.0185 - acc: 0.9939 - val_loss: 0.0780 - val_acc: 0.9798
Epoch 11/30
```

```

60000/60000 [=====] - 10s 164us/step - loss: 0.0137 - acc: 0.9953 - val_1
oss: 0.0792 - val_acc: 0.9804
Epoch 12/30
60000/60000 [=====] - 10s 163us/step - loss: 0.0146 - acc: 0.9957 - val_1
oss: 0.1051 - val_acc: 0.9760
Epoch 13/30
60000/60000 [=====] - 10s 162us/step - loss: 0.0130 - acc: 0.9957 - val_1
oss: 0.0877 - val_acc: 0.9805
Epoch 14/30
60000/60000 [=====] - 11s 178us/step - loss: 0.0117 - acc: 0.9963 - val_1
oss: 0.0854 - val_acc: 0.9808
Epoch 15/30
60000/60000 [=====] - 10s 163us/step - loss: 0.0114 - acc: 0.9965 - val_1
oss: 0.0852 - val_acc: 0.9824
Epoch 16/30
60000/60000 [=====] - 10s 172us/step - loss: 0.0108 - acc: 0.9967 - val_1
oss: 0.0961 - val_acc: 0.9800
Epoch 17/30
60000/60000 [=====] - 11s 178us/step - loss: 0.0150 - acc: 0.9955 - val_1
oss: 0.1109 - val_acc: 0.9759
Epoch 18/30
60000/60000 [=====] - 10s 168us/step - loss: 0.0079 - acc: 0.9975 - val_1
oss: 0.1035 - val_acc: 0.9774
Epoch 19/30
60000/60000 [=====] - 10s 164us/step - loss: 0.0104 - acc: 0.9967 - val_1
oss: 0.0814 - val_acc: 0.9827
Epoch 20/30
60000/60000 [=====] - 11s 181us/step - loss: 0.0118 - acc: 0.9961 - val_1
oss: 0.0868 - val_acc: 0.9811
Epoch 21/30
60000/60000 [=====] - 10s 164us/step - loss: 0.0054 - acc: 0.9983 - val_1
oss: 0.0895 - val_acc: 0.9827
Epoch 22/30
60000/60000 [=====] - 10s 162us/step - loss: 0.0110 - acc: 0.9968 - val_1
oss: 0.1167 - val_acc: 0.9769
Epoch 23/30
60000/60000 [=====] - 10s 164us/step - loss: 0.0094 - acc: 0.9971 - val_1
oss: 0.1086 - val_acc: 0.9783
Epoch 24/30
60000/60000 [=====] - 10s 162us/step - loss: 0.0074 - acc: 0.9977 - val_1
oss: 0.1158 - val_acc: 0.9793
Epoch 25/30
60000/60000 [=====] - 10s 162us/step - loss: 0.0076 - acc: 0.9979 - val_1
oss: 0.0930 - val_acc: 0.9823
Epoch 26/30
60000/60000 [=====] - 11s 179us/step - loss: 0.0090 - acc: 0.9976 - val_1
oss: 0.1045 - val_acc: 0.9810
Epoch 27/30
60000/60000 [=====] - 10s 162us/step - loss: 0.0059 - acc: 0.9982 - val_1
oss: 0.0920 - val_acc: 0.9821
Epoch 28/30
60000/60000 [=====] - 10s 162us/step - loss: 0.0110 - acc: 0.9971 - val_1
oss: 0.0996 - val_acc: 0.9795
Epoch 29/30
60000/60000 [=====] - 10s 165us/step - loss: 0.0074 - acc: 0.9981 - val_1
oss: 0.0939 - val_acc: 0.9814
Epoch 30/30
60000/60000 [=====] - 10s 164us/step - loss: 0.0056 - acc: 0.9984 - val_1
oss: 0.0951 - val_acc: 0.9816

```

In [42]:

```

score_relu_3 = model_relu_3.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score_relu_3[0])
print('Test accuracy:', score_relu_3[1])


fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```


Test score: 0.09514001322890968
Test accuracy: 0.9816



In [43]:


```
w_after = model_relu_3.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```





Batch normalization on hidden layer + Relu + Adam

In [45]:

```
model_batch_3 = Sequential()

model_batch_3.add(Dense(589, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.058, seed=None)))
model_batch_3.add(BatchNormalization())

model_batch_3.add(Dense(423, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.068, seed=None)) )
model_batch_3.add(BatchNormalization())

model_batch_3.add(Dense(272, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.085, seed=None)) )
model_batch_3.add(BatchNormalization())

model_batch_3.add(Dense(output_dim, activation='softmax'))

model_batch_3.summary()
```

Layer (type)	Output Shape	Param #
dense_27 (Dense)	(None, 589)	462365
batch_normalization_14 (Batch Normalization)	(None, 589)	2356
dense_28 (Dense)	(None, 423)	249570
batch_normalization_15 (Batch Normalization)	(None, 423)	1692
dense_29 (Dense)	(None, 272)	115328
batch_normalization_16 (Batch Normalization)	(None, 272)	1088
dense_30 (Dense)	(None, 10)	2730
Total params: 835,129		
Trainable params: 832,561		
Non-trainable params: 2,568		

In [46]:

```
model_batch_3.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_batch_3.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/30
60000/60000 [=====] - 16s 259us/step - loss: 0.1742 - acc: 0.9471 - val_loss: 0.0962 - val_acc: 0.9699
Epoch 2/30
```

```
60000/60000 [=====] - 14s 232us/step - loss: 0.0643 - acc: 0.9801 - val_l
oss: 0.0826 - val_acc: 0.9744
Epoch 3/30
60000/60000 [=====] - 13s 209us/step - loss: 0.0463 - acc: 0.9851 - val_l
oss: 0.0828 - val_acc: 0.9734
Epoch 4/30
60000/60000 [=====] - 13s 212us/step - loss: 0.0340 - acc: 0.9891 - val_l
oss: 0.1240 - val_acc: 0.9660
Epoch 5/30
60000/60000 [=====] - 12s 207us/step - loss: 0.0301 - acc: 0.9894 - val_l
oss: 0.0848 - val_acc: 0.9748
Epoch 6/30
60000/60000 [=====] - 12s 207us/step - loss: 0.0232 - acc: 0.9922 - val_l
oss: 0.0811 - val_acc: 0.9769
Epoch 7/30
60000/60000 [=====] - 14s 229us/step - loss: 0.0217 - acc: 0.9924 - val_l
oss: 0.0772 - val_acc: 0.9780
Epoch 8/30
60000/60000 [=====] - 13s 219us/step - loss: 0.0181 - acc: 0.9937 - val_l
oss: 0.0784 - val_acc: 0.9796
Epoch 9/30
60000/60000 [=====] - 13s 219us/step - loss: 0.0180 - acc: 0.9936 - val_l
oss: 0.0938 - val_acc: 0.9760
Epoch 10/30
60000/60000 [=====] - 13s 209us/step - loss: 0.0182 - acc: 0.9943 - val_l
oss: 0.0753 - val_acc: 0.9783
Epoch 11/30
60000/60000 [=====] - 13s 222us/step - loss: 0.0160 - acc: 0.9946 - val_l
oss: 0.0831 - val_acc: 0.9792
Epoch 12/30
60000/60000 [=====] - 13s 218us/step - loss: 0.0144 - acc: 0.9950 - val_l
oss: 0.0808 - val_acc: 0.9794
Epoch 13/30
60000/60000 [=====] - 13s 218us/step - loss: 0.0114 - acc: 0.9965 - val_l
oss: 0.0697 - val_acc: 0.9823
Epoch 14/30
60000/60000 [=====] - 13s 220us/step - loss: 0.0101 - acc: 0.9964 - val_l
oss: 0.0919 - val_acc: 0.9772
Epoch 15/30
60000/60000 [=====] - 13s 218us/step - loss: 0.0132 - acc: 0.9956 - val_l
oss: 0.0850 - val_acc: 0.9783
Epoch 16/30
60000/60000 [=====] - 14s 232us/step - loss: 0.0142 - acc: 0.9954 - val_l
oss: 0.0739 - val_acc: 0.9828
Epoch 17/30
60000/60000 [=====] - 13s 220us/step - loss: 0.0083 - acc: 0.9970 - val_l
oss: 0.0735 - val_acc: 0.9814
Epoch 18/30
60000/60000 [=====] - 13s 220us/step - loss: 0.0090 - acc: 0.9972 - val_l
oss: 0.0784 - val_acc: 0.9825
Epoch 19/30
60000/60000 [=====] - 13s 224us/step - loss: 0.0112 - acc: 0.9963 - val_l
oss: 0.0916 - val_acc: 0.9792
Epoch 20/30
60000/60000 [=====] - 14s 233us/step - loss: 0.0082 - acc: 0.9972 - val_l
oss: 0.0852 - val_acc: 0.9801
Epoch 21/30
60000/60000 [=====] - 13s 220us/step - loss: 0.0062 - acc: 0.9979 - val_l
oss: 0.0767 - val_acc: 0.9827
Epoch 22/30
60000/60000 [=====] - 13s 220us/step - loss: 0.0095 - acc: 0.9969 - val_l
oss: 0.0883 - val_acc: 0.9801
Epoch 23/30
60000/60000 [=====] - 13s 222us/step - loss: 0.0096 - acc: 0.9970 - val_l
oss: 0.0802 - val_acc: 0.9805
Epoch 24/30
60000/60000 [=====] - 14s 226us/step - loss: 0.0080 - acc: 0.9975 - val_l
oss: 0.0832 - val_acc: 0.9815
Epoch 25/30
60000/60000 [=====] - 13s 223us/step - loss: 0.0069 - acc: 0.9978 - val_l
oss: 0.0863 - val_acc: 0.9811
Epoch 26/30
60000/60000 [=====] - 13s 215us/step - loss: 0.0053 - acc: 0.9982 - val_l
oss: 0.0802 - val_acc: 0.9815
Epoch 27/30
60000/60000 [=====] - 13s 225us/step - loss: 0.0068 - acc: 0.9978 - val_l
oss: 0.0896 - val_acc: 0.9805
```

```
Epoch 28/30
60000/60000 [=====] - 13s 220us/step - loss: 0.0075 - acc: 0.9975 - val_1
oss: 0.0816 - val_acc: 0.9811
Epoch 29/30
60000/60000 [=====] - 14s 236us/step - loss: 0.0060 - acc: 0.9978 - val_1
oss: 0.0850 - val_acc: 0.9834
Epoch 30/30
60000/60000 [=====] - 13s 220us/step - loss: 0.0070 - acc: 0.9976 - val_1
oss: 0.0864 - val_acc: 0.9817
```

In [47]:


```
score_batch_3 = model_batch_3.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score_batch_3[0])
print('Test accuracy:', score_batch_3[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

```
Test score: 0.08639855843480118
Test accuracy: 0.9817
```



In [48]:


```
w_after = model_batch_3.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')
```

```
plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



Drop + Relu + Adam

In [49]:

```
model_drop_3 = Sequential()

model_drop_3.add(Dense(401, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.070, seed=None)))
model_drop_3.add(BatchNormalization())
model_drop_3.add(Dropout(0.5))

model_drop_3.add(Dense(219, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.095, seed=None)))
model_drop_3.add(BatchNormalization())
model_drop_3.add(Dropout(0.5))

model_drop_3.add(Dense(121, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.128, seed=None)))
model_drop_3.add(BatchNormalization())
model_drop_3.add(Dropout(0.5))

model_drop_3.add(Dense(output_dim, activation='softmax'))

model_drop_3.summary()
```

Layer (type)	Output Shape	Param #
dense_31 (Dense)	(None, 401)	314785
batch_normalization_17 (Batch Normalization)	(None, 401)	1604

batch_normalization_17 (Batch Normalization)	(None, 401)	0
dropout_9 (Dropout)	(None, 401)	0
dense_32 (Dense)	(None, 219)	88038
batch_normalization_18 (Batch Normalization)	(None, 219)	876
dropout_10 (Dropout)	(None, 219)	0
dense_33 (Dense)	(None, 121)	26620
batch_normalization_19 (Batch Normalization)	(None, 121)	484
dropout_11 (Dropout)	(None, 121)	0
dense_34 (Dense)	(None, 10)	1220
=====		
Total params: 433,627		
Trainable params: 432,145		
Non-trainable params: 1,482		
=====		

In [50]:

```
model_drop_3.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_drop_3.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/30
60000/60000 [=====] - 14s 227us/step - loss: 0.7051 - acc: 0.7825 - val_loss: 0.1937 - val_acc: 0.9398
Epoch 2/30
60000/60000 [=====] - 11s 190us/step - loss: 0.2958 - acc: 0.9122 - val_loss: 0.1423 - val_acc: 0.9571
Epoch 3/30
60000/60000 [=====] - 11s 180us/step - loss: 0.2296 - acc: 0.9335 - val_loss: 0.1228 - val_acc: 0.9622
Epoch 4/30
60000/60000 [=====] - 11s 178us/step - loss: 0.1852 - acc: 0.9450 - val_loss: 0.1040 - val_acc: 0.9688
Epoch 5/30
60000/60000 [=====] - 11s 176us/step - loss: 0.1623 - acc: 0.9525 - val_loss: 0.0963 - val_acc: 0.9722
Epoch 6/30
60000/60000 [=====] - 11s 183us/step - loss: 0.1507 - acc: 0.9566 - val_loss: 0.0854 - val_acc: 0.9741
Epoch 7/30
60000/60000 [=====] - 11s 177us/step - loss: 0.1365 - acc: 0.9603 - val_loss: 0.0850 - val_acc: 0.9747
Epoch 8/30
60000/60000 [=====] - 11s 184us/step - loss: 0.1241 - acc: 0.9631 - val_loss: 0.0837 - val_acc: 0.9755
Epoch 9/30
60000/60000 [=====] - 11s 175us/step - loss: 0.1148 - acc: 0.9662 - val_loss: 0.0748 - val_acc: 0.9773
Epoch 10/30
60000/60000 [=====] - 11s 185us/step - loss: 0.1088 - acc: 0.9678 - val_loss: 0.0690 - val_acc: 0.9804
Epoch 11/30
60000/60000 [=====] - 11s 184us/step - loss: 0.1023 - acc: 0.9699 - val_loss: 0.0676 - val_acc: 0.9799
Epoch 12/30
60000/60000 [=====] - 11s 186us/step - loss: 0.1000 - acc: 0.9695 - val_loss: 0.0669 - val_acc: 0.9815
Epoch 13/30
60000/60000 [=====] - 11s 190us/step - loss: 0.0952 - acc: 0.9709 - val_loss: 0.0672 - val_acc: 0.9805
Epoch 14/30
60000/60000 [=====] - 11s 184us/step - loss: 0.0869 - acc: 0.9739 - val_loss: 0.0658 - val_acc: 0.9814
Epoch 15/30
60000/60000 [=====] - 11s 183us/step - loss: 0.0851 - acc: 0.9742 - val_loss: 0.0635 - val_acc: 0.9813
```

```

Epoch 16/30
60000/60000 [=====] - 11s 184us/step - loss: 0.0839 - acc: 0.9744 - val_loss: 0.0600 - val_acc: 0.9824
Epoch 17/30
60000/60000 [=====] - 11s 184us/step - loss: 0.0766 - acc: 0.9762 - val_loss: 0.0601 - val_acc: 0.9830
Epoch 18/30
60000/60000 [=====] - 12s 192us/step - loss: 0.0753 - acc: 0.9770 - val_loss: 0.0612 - val_acc: 0.9824
Epoch 19/30
60000/60000 [=====] - 11s 187us/step - loss: 0.0725 - acc: 0.9784 - val_loss: 0.0645 - val_acc: 0.9820
Epoch 20/30
60000/60000 [=====] - 11s 184us/step - loss: 0.0694 - acc: 0.9793 - val_loss: 0.0642 - val_acc: 0.9827
Epoch 21/30
60000/60000 [=====] - 11s 184us/step - loss: 0.0671 - acc: 0.9801 - val_loss: 0.0635 - val_acc: 0.9824
Epoch 22/30
60000/60000 [=====] - 11s 186us/step - loss: 0.0668 - acc: 0.9801 - val_loss: 0.0587 - val_acc: 0.9833
Epoch 23/30
60000/60000 [=====] - 11s 184us/step - loss: 0.0659 - acc: 0.9804 - val_loss: 0.0577 - val_acc: 0.9849
Epoch 24/30
60000/60000 [=====] - 12s 193us/step - loss: 0.0597 - acc: 0.9821 - val_loss: 0.0594 - val_acc: 0.9840
Epoch 25/30
60000/60000 [=====] - 11s 184us/step - loss: 0.0617 - acc: 0.9809 - val_loss: 0.0612 - val_acc: 0.9832
Epoch 26/30
60000/60000 [=====] - 11s 188us/step - loss: 0.0583 - acc: 0.9824 - val_loss: 0.0598 - val_acc: 0.9841
Epoch 27/30
60000/60000 [=====] - 11s 185us/step - loss: 0.0581 - acc: 0.9819 - val_loss: 0.0609 - val_acc: 0.9839
Epoch 28/30
60000/60000 [=====] - 11s 185us/step - loss: 0.0595 - acc: 0.9824 - val_loss: 0.0593 - val_acc: 0.9839
Epoch 29/30
60000/60000 [=====] - 12s 192us/step - loss: 0.0540 - acc: 0.9836 - val_loss: 0.0568 - val_acc: 0.9848
Epoch 30/30
60000/60000 [=====] - 11s 187us/step - loss: 0.0544 - acc: 0.9832 - val_loss: 0.0568 - val_acc: 0.9858

```

In [51]:

```

score_drop_3 = model_drop_3.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score_drop_3[0])
print('Test accuracy:', score_drop_3[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')


# list of epoch numbers
x = list(range(1, nb_epoch+1))


vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.05677158319847949

Test accuracy: 0.9858





In [52]:


```
w_after = model_drop_3.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



5 Layered Architecture

Relu + Adam

In [53]:

```
# start building a model
model_relu_5 = Sequential()
model_relu_5.add(Dense(697, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.053, seed=None)))
model_relu_5.add(Dense(550, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.060, seed=None)))
model_relu_5.add(Dense(462, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.065, seed=None)))
model_relu_5.add(Dense(322, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.078, seed=None)))
model_relu_5.add(Dense(199, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.100, seed=None)))
model_relu_5.add(Dense(output_dim, activation='softmax'))

print(model_relu_5.summary())
```

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_35 (Dense)	(None, 697)	547145
dense_36 (Dense)	(None, 550)	383900
dense_37 (Dense)	(None, 462)	254562
dense_38 (Dense)	(None, 322)	149086
dense_39 (Dense)	(None, 199)	64277
dense_40 (Dense)	(None, 10)	2000
=====	=====	=====
Total params: 1,400,970		
Trainable params: 1,400,970		
Non-trainable params: 0		
None		

In [54]:

```
model_relu_5.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu_5.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/30
60000/60000 [=====] - 20s 336us/step - loss: 0.2144 - acc: 0.9343 - val_loss: 0.1240 - val_acc: 0.9648
Epoch 2/30
60000/60000 [=====] - 18s 302us/step - loss: 0.0875 - acc: 0.9726 - val_loss: 0.0997 - val_acc: 0.9691
Epoch 3/30
60000/60000 [=====] - 18s 295us/step - loss: 0.0587 - acc: 0.9813 - val_loss: 0.0940 - val_acc: 0.9722
Epoch 4/30
60000/60000 [=====] - 17s 288us/step - loss: 0.0467 - acc: 0.9851 - val_loss: 0.0828 - val_acc: 0.9745
Epoch 5/30
60000/60000 [=====] - 19s 324us/step - loss: 0.0387 - acc: 0.9877 - val_loss: 0.1106 - val_acc: 0.9712
Epoch 6/30
60000/60000 [=====] - 18s 300us/step - loss: 0.0337 - acc: 0.9900 - val_loss: 0.0836 - val_acc: 0.9748
```

```
Epoch 7/30
60000/60000 [=====] - 18s 292us/step - loss: 0.0305 - acc: 0.9906 - val_1
oss: 0.0880 - val_acc: 0.9793
Epoch 8/30
60000/60000 [=====] - 18s 308us/step - loss: 0.0281 - acc: 0.9917 - val_1
oss: 0.1000 - val_acc: 0.9753
Epoch 9/30
60000/60000 [=====] - 18s 297us/step - loss: 0.0243 - acc: 0.9925 - val_1
oss: 0.1016 - val_acc: 0.9757
Epoch 10/30
60000/60000 [=====] - 17s 287us/step - loss: 0.0201 - acc: 0.9939 - val_1
oss: 0.0913 - val_acc: 0.9784
Epoch 11/30
60000/60000 [=====] - 17s 290us/step - loss: 0.0197 - acc: 0.9939 - val_1
oss: 0.0755 - val_acc: 0.9827
Epoch 12/30
60000/60000 [=====] - 18s 302us/step - loss: 0.0215 - acc: 0.9936 - val_1
oss: 0.0827 - val_acc: 0.9807
Epoch 13/30
60000/60000 [=====] - 17s 286us/step - loss: 0.0154 - acc: 0.9951 - val_1
oss: 0.1097 - val_acc: 0.9790
Epoch 14/30
60000/60000 [=====] - 18s 292us/step - loss: 0.0141 - acc: 0.9962 - val_1
oss: 0.0779 - val_acc: 0.9821
Epoch 15/30
60000/60000 [=====] - 19s 310us/step - loss: 0.0178 - acc: 0.9950 - val_1
oss: 0.0935 - val_acc: 0.9790
Epoch 16/30
60000/60000 [=====] - 18s 292us/step - loss: 0.0153 - acc: 0.9956 - val_1
oss: 0.0952 - val_acc: 0.9795
Epoch 17/30
60000/60000 [=====] - 18s 306us/step - loss: 0.0127 - acc: 0.9964 - val_1
oss: 0.0932 - val_acc: 0.9799
Epoch 18/30
60000/60000 [=====] - 18s 306us/step - loss: 0.0172 - acc: 0.9954 - val_1
oss: 0.0849 - val_acc: 0.9815
Epoch 19/30
60000/60000 [=====] - 18s 298us/step - loss: 0.0145 - acc: 0.9960 - val_1
oss: 0.1123 - val_acc: 0.9778
Epoch 20/30
60000/60000 [=====] - 18s 296us/step - loss: 0.0147 - acc: 0.9967 - val_1
oss: 0.1041 - val_acc: 0.9806
Epoch 21/30
60000/60000 [=====] - 18s 293us/step - loss: 0.0097 - acc: 0.9972 - val_1
oss: 0.0843 - val_acc: 0.9835
Epoch 22/30
60000/60000 [=====] - 18s 302us/step - loss: 0.0124 - acc: 0.9968 - val_1
oss: 0.0881 - val_acc: 0.9824
Epoch 23/30
60000/60000 [=====] - 17s 290us/step - loss: 0.0105 - acc: 0.9970 - val_1
oss: 0.1012 - val_acc: 0.9819
Epoch 24/30
60000/60000 [=====] - 18s 292us/step - loss: 0.0127 - acc: 0.9967 - val_1
oss: 0.0963 - val_acc: 0.9805
Epoch 25/30
60000/60000 [=====] - 19s 313us/step - loss: 0.0083 - acc: 0.9977 - val_1
oss: 0.1070 - val_acc: 0.9827
Epoch 26/30
60000/60000 [=====] - 17s 287us/step - loss: 0.0108 - acc: 0.9970 - val_1
oss: 0.1149 - val_acc: 0.9806
Epoch 27/30
60000/60000 [=====] - 17s 288us/step - loss: 0.0124 - acc: 0.9970 - val_1
oss: 0.0957 - val_acc: 0.9822
Epoch 28/30
60000/60000 [=====] - 17s 288us/step - loss: 0.0095 - acc: 0.9976 - val_1
oss: 0.0884 - val_acc: 0.9837
Epoch 29/30
60000/60000 [=====] - 18s 304us/step - loss: 0.0137 - acc: 0.9965 - val_1
oss: 0.0969 - val_acc: 0.9824
Epoch 30/30
60000/60000 [=====] - 18s 293us/step - loss: 0.0107 - acc: 0.9974 - val_1
oss: 0.0969 - val_acc: 0.9801
```

In [55]:

```
score_relu_5 = model_relu_5.evaluate(X_test, Y_test, verbose=0)
```

```

print('Test score:', score_relu_5[0])
print('Test accuracy:', score_relu_5[1])


fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.09693762213855957
Test accuracy: 0.9801



In [56]:

```

w_after = model_relu_5.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



Batch Normalization on Hidden Layers + Relu + Adam

In [57]:

```
model_batch_5 = Sequential()

model_batch_5.add(Dense(667, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.054, seed=None)))
model_batch_5.add(BatchNormalization())

model_batch_5.add(Dense(579, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.058, seed=None)) )
model_batch_5.add(BatchNormalization())

model_batch_5.add(Dense(499, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.063, seed=None)) )
model_batch_5.add(BatchNormalization())

model_batch_5.add(Dense(349, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.075, seed=None)) )
model_batch_5.add(BatchNormalization())

model_batch_5.add(Dense(205, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.098, seed=None)) )
model_batch_5.add(BatchNormalization())

model_batch_5.add(Dense(output_dim, activation='softmax'))

model_batch_5.summary()
```

Layer (type)	Output Shape	Param #
dense_41 (Dense)	(None, 667)	523595
batch_normalization_20 (Batch Normalization)	(None, 667)	2668
dense_42 (Dense)	(None, 579)	386772
batch_normalization_21 (Batch Normalization)	(None, 579)	2316
dense_43 (Dense)	(None, 499)	289420
batch_normalization_22 (Batch Normalization)	(None, 499)	1996

dense_44 (Dense)	(None, 349)	174500
batch_normalization_23 (Batch Normalization)	(None, 349)	1396
dense_45 (Dense)	(None, 205)	71750
batch_normalization_24 (Batch Normalization)	(None, 205)	820
dense_46 (Dense)	(None, 10)	2060
=====		
Total params: 1,457,293		
Trainable params: 1,452,695		
Non-trainable params: 4,598		

In [58]:

```
model_batch_5.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_batch_5.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/30
60000/60000 [=====] - 28s 466us/step - loss: 0.1868 - acc: 0.9430 - val_loss: 0.1172 - val_acc: 0.9652
Epoch 2/30
60000/60000 [=====] - 24s 401us/step - loss: 0.0786 - acc: 0.9756 - val_loss: 0.0948 - val_acc: 0.9700
Epoch 3/30
60000/60000 [=====] - 25s 418us/step - loss: 0.0562 - acc: 0.9816 - val_loss: 0.0929 - val_acc: 0.9730
Epoch 4/30
60000/60000 [=====] - 24s 401us/step - loss: 0.0463 - acc: 0.9853 - val_loss: 0.0935 - val_acc: 0.9725
Epoch 5/30
60000/60000 [=====] - 26s 427us/step - loss: 0.0406 - acc: 0.9869 - val_loss: 0.0850 - val_acc: 0.9742
Epoch 6/30
60000/60000 [=====] - 24s 402us/step - loss: 0.0339 - acc: 0.9886 - val_loss: 0.0818 - val_acc: 0.9769
Epoch 7/30
60000/60000 [=====] - 24s 403us/step - loss: 0.0326 - acc: 0.9893 - val_loss: 0.0851 - val_acc: 0.9754
Epoch 8/30
60000/60000 [=====] - 25s 418us/step - loss: 0.0264 - acc: 0.9912 - val_loss: 0.0864 - val_acc: 0.9775
Epoch 9/30
60000/60000 [=====] - 24s 399us/step - loss: 0.0253 - acc: 0.9919 - val_loss: 0.0801 - val_acc: 0.9782
Epoch 10/30
60000/60000 [=====] - 25s 414us/step - loss: 0.0233 - acc: 0.9920 - val_loss: 0.0914 - val_acc: 0.9743
Epoch 11/30
60000/60000 [=====] - 24s 400us/step - loss: 0.0239 - acc: 0.9923 - val_loss: 0.0754 - val_acc: 0.9794
Epoch 12/30
60000/60000 [=====] - 24s 404us/step - loss: 0.0201 - acc: 0.9932 - val_loss: 0.0921 - val_acc: 0.9759
Epoch 13/30
60000/60000 [=====] - 25s 415us/step - loss: 0.0224 - acc: 0.9924 - val_loss: 0.0864 - val_acc: 0.9785
Epoch 14/30
60000/60000 [=====] - 24s 400us/step - loss: 0.0158 - acc: 0.9947 - val_loss: 0.0756 - val_acc: 0.9798
Epoch 15/30
60000/60000 [=====] - 25s 418us/step - loss: 0.0183 - acc: 0.9938 - val_loss: 0.0750 - val_acc: 0.9809
Epoch 16/30
60000/60000 [=====] - 24s 402us/step - loss: 0.0169 - acc: 0.9947 - val_loss: 0.0672 - val_acc: 0.9831
Epoch 17/30
60000/60000 [=====] - 24s 405us/step - loss: 0.0139 - acc: 0.9957 - val_loss: 0.0648 - val_acc: 0.9832
Epoch 18/30
```

```

60000/60000 [=====] - 25s 418us/step - loss: 0.0136 - acc: 0.9954 - val_loss: 0.0740 - val_acc: 0.9828
Epoch 19/30
60000/60000 [=====] - 24s 406us/step - loss: 0.0135 - acc: 0.9957 - val_loss: 0.0809 - val_acc: 0.9819
Epoch 20/30
60000/60000 [=====] - 25s 416us/step - loss: 0.0158 - acc: 0.9952 - val_loss: 0.0721 - val_acc: 0.9813
Epoch 21/30
60000/60000 [=====] - 24s 406us/step - loss: 0.0135 - acc: 0.9955 - val_loss: 0.0798 - val_acc: 0.9808
Epoch 22/30
60000/60000 [=====] - 25s 412us/step - loss: 0.0094 - acc: 0.9969 - val_loss: 0.0690 - val_acc: 0.9835
Epoch 23/30
60000/60000 [=====] - 25s 409us/step - loss: 0.0130 - acc: 0.9956 - val_loss: 0.0845 - val_acc: 0.9805
Epoch 24/30
60000/60000 [=====] - 24s 400us/step - loss: 0.0108 - acc: 0.9964 - val_loss: 0.0743 - val_acc: 0.9826
Epoch 25/30
60000/60000 [=====] - 25s 416us/step - loss: 0.0081 - acc: 0.9973 - val_loss: 0.0753 - val_acc: 0.9825
Epoch 26/30
60000/60000 [=====] - 24s 404us/step - loss: 0.0092 - acc: 0.9970 - val_loss: 0.0675 - val_acc: 0.9826
Epoch 27/30
60000/60000 [=====] - 25s 418us/step - loss: 0.0101 - acc: 0.9966 - val_loss: 0.0846 - val_acc: 0.9821
Epoch 28/30
60000/60000 [=====] - 24s 404us/step - loss: 0.0118 - acc: 0.9962 - val_loss: 0.0752 - val_acc: 0.9821
Epoch 29/30
60000/60000 [=====] - 25s 408us/step - loss: 0.0075 - acc: 0.9978 - val_loss: 0.0731 - val_acc: 0.9828
Epoch 30/30
60000/60000 [=====] - 25s 421us/step - loss: 0.0081 - acc: 0.9976 - val_loss: 0.0669 - val_acc: 0.9851

```

In [59]:

```

score_batch_5 = model_batch_5.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score_batch_5[0])
print('Test accuracy:', score_batch_5[1])


fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

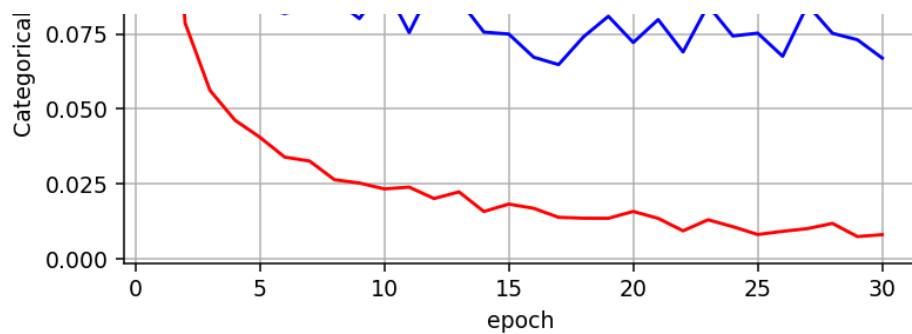
# list of epoch numbers
x = list(range(1, nb_epoch+1))

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.06694545334176073
Test accuracy: 0.9851





In [60]:


```
w_after = model_batch_5.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



Dropout + Relu + Adam

In [61]:

```

model_drop_5 = Sequential()

model_drop_5.add(Dense(609, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.057, seed=None)))
model_drop_5.add(BatchNormalization())
model_drop_5.add(Dropout(0.5))

model_drop_5.add(Dense(599, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.057, seed=None)) )
model_drop_5.add(BatchNormalization())
model_drop_5.add(Dropout(0.5))

model_drop_5.add(Dense(410, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.069, seed=None)) )
model_drop_5.add(BatchNormalization())
model_drop_5.add(Dropout(0.5))

model_drop_5.add(Dense(232, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.092, seed=None)) )
model_drop_5.add(BatchNormalization())
model_drop_5.add(Dropout(0.5))

model_drop_5.add(Dense(144, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.117, seed=None)) )
model_drop_5.add(BatchNormalization())
model_drop_5.add(Dropout(0.5))

model_drop_5.add(Dense(output_dim, activation='softmax'))

model_drop_5.summary()

```

Layer (type)	Output Shape	Param #
dense_47 (Dense)	(None, 609)	478065
batch_normalization_25 (Batch Normalization)	(None, 609)	2436
dropout_12 (Dropout)	(None, 609)	0
dense_48 (Dense)	(None, 599)	365390
batch_normalization_26 (Batch Normalization)	(None, 599)	2396
dropout_13 (Dropout)	(None, 599)	0
dense_49 (Dense)	(None, 410)	246000
batch_normalization_27 (Batch Normalization)	(None, 410)	1640
dropout_14 (Dropout)	(None, 410)	0
dense_50 (Dense)	(None, 232)	95352
batch_normalization_28 (Batch Normalization)	(None, 232)	928
dropout_15 (Dropout)	(None, 232)	0
dense_51 (Dense)	(None, 144)	33552
batch_normalization_29 (Batch Normalization)	(None, 144)	576
dropout_16 (Dropout)	(None, 144)	0
dense_52 (Dense)	(None, 10)	1450
Total params: 1,227,785		
Trainable params: 1,223,797		
Non-trainable params: 3,988		

In [62]:

```

model_drop_5.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

```



```

oss: 0.0583 - val_acc: 0.9839
Epoch 25/30
60000/60000 [=====] - 24s 403us/step - loss: 0.0616 - acc: 0.9829 - val_1
oss: 0.0564 - val_acc: 0.9848
Epoch 26/30
60000/60000 [=====] - 25s 408us/step - loss: 0.0633 - acc: 0.9815 - val_1
oss: 0.0620 - val_acc: 0.9840
Epoch 27/30
60000/60000 [=====] - 25s 416us/step - loss: 0.0620 - acc: 0.9823 - val_1
oss: 0.0594 - val_acc: 0.9843
Epoch 28/30
60000/60000 [=====] - 24s 405us/step - loss: 0.0574 - acc: 0.9836 - val_1
oss: 0.0573 - val_acc: 0.9852
Epoch 29/30
60000/60000 [=====] - 25s 417us/step - loss: 0.0549 - acc: 0.9839 - val_1
oss: 0.0573 - val_acc: 0.9844
Epoch 30/30
60000/60000 [=====] - 24s 405us/step - loss: 0.0546 - acc: 0.9848 - val_1
oss: 0.0567 - val_acc: 0.9855

```

In [63]:

```

score_drop_5 = model_drop_5.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score_drop_5[0])
print('Test accuracy:', score_drop_5[1])


fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.056678825084562415
Test accuracy: 0.9855



In [64]:

```

w_after = model_drop_5.get_weights()
h1 w = w_after[0].flatten().reshape(-1,1)

```

```


h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



Prettytable

In [74]:

```

number= [1,2,3,4,5,6,7,8,9]
name= ["Relu layer 2","Relu layer 3","Relu layer 5"," Batch layer 2"," Batch layer 3"," Batch layer 5","Drop layer 2","Drop layer 3","Drop layer 5"]
score=
[score_relu_2[0],score_relu_3[0],score_relu_5[0],score_batch_2[0],score_batch_3[0],score_batch_5[0],score_drop_2[0],score_drop_3[0],score_drop_5[0]]
acc=
[score_relu_2[1],score_relu_3[1],score_relu_5[1],score_batch_2[1],score_batch_3[1],score_batch_5[1],score_drop_2[1],score_drop_3[1],score_drop_5[1]]

#Initialize Prettytable
ptable = PrettyTable()
ptable.add_column("Index", number)
ptable.add_column("Architecture", name)
ptable.add_column("Test Score", score)
ptable.add_column("Test Accuracy", acc)

print(ptable)

```

Index	Architecture	Test Score	Test Accuracy
1	Relu layer 2	0.09564156731741441	0.9817
2	Relu layer 3	0.09514001322890968	0.9816
3	Relu layer 5	0.09693762213855957	0.9801
4	Batch layer 2	0.09071999986333348	0.9804
5	Batch layer 3	0.08639855843480118	0.9817
6	Batch layer 5	0.06694545334176073	0.9851
7	Drop layer 2	0.052775146638258594	0.9848
8	Drop layer 3	0.05677158319847949	0.9858
9	Drop layer 5	0.056678825084562415	0.9855

Conclusions

1. We implement Keras on MNIST dataset.
2. Different layered architectures like 2, 3, 5 hidden layers are used to build Neural networks.
3. Activation method RELU is used and for optimization ADAM is used.
4. In each layered network, Batch normalization and dropout layer is also added to check the performance of model.
5. Test accuracy is same for all architectures.
6. Test score is same in case of Relu, in case of dropout test score is low than Relu.

In []: