

CNN_MNIST

August 9, 2019

```
In [1]: import warnings
        warnings.filterwarnings("ignore")

        from __future__ import print_function
        import keras
        from keras.datasets import mnist
        from keras.models import Sequential
        from keras.layers import Dense, Dropout, Flatten
        from keras.layers import Conv2D, MaxPooling2D
        from keras import backend as K
        from keras.layers import Activation, BatchNormalization, regularizers
```

Using TensorFlow backend.

```
In [9]: batch_size = 128
        num_classes = 10
        epochs = 12

        # input image dimensions
        img_rows, img_cols = 28, 28

        # the data, split between train and test sets
        (x_train, y_train), (x_test, y_test) = mnist.load_data()

        if K.image_data_format() == 'channels_first':
            x_train = x_train.reshape(x_train.shape[0], 1, img_rows, img_cols)
            x_test = x_test.reshape(x_test.shape[0], 1, img_rows, img_cols)
            input_shape = (1, img_rows, img_cols)
        else:
            x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 1)
            x_test = x_test.reshape(x_test.shape[0], img_rows, img_cols, 1)
            input_shape = (img_rows, img_cols, 1)

        x_train = x_train.astype('float32')
        x_test = x_test.astype('float32')
        x_train /= 255
        x_test /= 255
```

```

print('x_train shape:', x_train.shape)
print(x_train.shape[0], 'train samples')
print(x_test.shape[0], 'test samples')

# convert class vectors to binary class matrices
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)

```

```

x_train shape: (60000, 28, 28, 1)
60000 train samples
10000 test samples

```

In [3]: # Credits: https://github.com/keras-team/keras/blob/master/examples/mnist_cnn.py

```

model = Sequential()
model.add(Conv2D(32, kernel_size=(3, 3),
                 activation='relu',
                 input_shape=input_shape))
model.add(Conv2D(64, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))
model.add(Flatten())
model.add(Dense(128, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax'))

model.compile(loss=keras.losses.categorical_crossentropy,
              optimizer=keras.optimizers.Adadelta(),
              metrics=['accuracy'])

model.fit(x_train, y_train,
          batch_size=batch_size,
          epochs=epochs,
          verbose=1,
          validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])

```

WARNING: Logging before flag parsing goes to stderr.

W0809 01:32:51.226855 12008 deprecation_wrapper.py:119] From C:\Users\ACER\Anaconda3\lib\site-p

W0809 01:32:51.320584 12008 deprecation_wrapper.py:119] From C:\Users\ACER\Anaconda3\lib\site-p

W0809 01:32:51.320584 12008 deprecation_wrapper.py:119] From C:\Users\ACER\Anaconda3\lib\site-p

W0809 01:32:51.383068 12008 deprecation_wrapper.py:119] From C:\Users\ACER\Anaconda3\lib\site-p

```

W0809 01:32:51.383068 12008 deprecation_wrapper.py:119] From C:\Users\ACER\Anaconda3\lib\site-packages\tensorflow\python\ops\stack.py:113: tf.nn.conv2d (from tensorflow.python.ops.nn_ops) is deprecated and will be removed in a future version.
Instructions for updating:
Please use `rate` instead of `keep_prob`. Rate should be set to `rate = 1 - keep_prob`.
W0809 01:32:51.523696 12008 deprecation_wrapper.py:119] From C:\Users\ACER\Anaconda3\lib\site-packages\tensorflow\python\ops\stack.py:113: tf.nn.conv2d (from tensorflow.python.ops.nn_ops) is deprecated and will be removed in a future version.
Instructions for updating:
Please use `rate` instead of `keep_prob`. Rate should be set to `rate = 1 - keep_prob`.
W0809 01:32:51.539281 12008 deprecation_wrapper.py:119] From C:\Users\ACER\Anaconda3\lib\site-packages\tensorflow\python\ops\stack.py:113: tf.nn.conv2d (from tensorflow.python.ops.nn_ops) is deprecated and will be removed in a future version.
Instructions for updating:
Please use `rate` instead of `keep_prob`. Rate should be set to `rate = 1 - keep_prob`.
W0809 01:32:51.711116 12008 deprecation.py:323] From C:\Users\ACER\Anaconda3\lib\site-packages\tensorflow\python\ops\stack.py:113: tf.nn.conv2d (from tensorflow.python.ops.nn_ops) is deprecated and will be removed in a future version.
Instructions for updating:
Please use `rate` instead of `keep_prob`. Rate should be set to `rate = 1 - keep_prob`.
Use tf.where in 2.0, which has the same broadcast rule as np.where

```

Train on 60000 samples, validate on 10000 samples

```

Epoch 1/12
60000/60000 [=====] - 104s 2ms/step - loss: 0.2648 - acc: 0.9184 - val_loss: 0.024807610545822537
Epoch 2/12
60000/60000 [=====] - 107s 2ms/step - loss: 0.0873 - acc: 0.9739 - val_loss: 0.024807610545822537
Epoch 3/12
60000/60000 [=====] - 106s 2ms/step - loss: 0.0671 - acc: 0.9799 - val_loss: 0.024807610545822537
Epoch 4/12
60000/60000 [=====] - 105s 2ms/step - loss: 0.0552 - acc: 0.9830 - val_loss: 0.024807610545822537
Epoch 5/12
60000/60000 [=====] - 106s 2ms/step - loss: 0.0464 - acc: 0.9856 - val_loss: 0.024807610545822537
Epoch 6/12
60000/60000 [=====] - 106s 2ms/step - loss: 0.0428 - acc: 0.9872 - val_loss: 0.024807610545822537
Epoch 7/12
60000/60000 [=====] - 105s 2ms/step - loss: 0.0381 - acc: 0.9888 - val_loss: 0.024807610545822537
Epoch 8/12
60000/60000 [=====] - 73s 1ms/step - loss: 0.0345 - acc: 0.9893 - val_loss: 0.024807610545822537
Epoch 9/12
60000/60000 [=====] - 74s 1ms/step - loss: 0.0299 - acc: 0.9902 - val_loss: 0.024807610545822537
Epoch 10/12
60000/60000 [=====] - 76s 1ms/step - loss: 0.0306 - acc: 0.9908 - val_loss: 0.024807610545822537
Epoch 11/12
60000/60000 [=====] - 75s 1ms/step - loss: 0.0277 - acc: 0.9915 - val_loss: 0.024807610545822537
Epoch 12/12
60000/60000 [=====] - 73s 1ms/step - loss: 0.0274 - acc: 0.9917 - val_loss: 0.024807610545822537
Test loss: 0.024807610545822537
Test accuracy: 0.9925

```

0.0.1 Example 1

```

In [10]: model_1 = Sequential()
          model_1.add(Conv2D(32, kernel_size=(2, 2),

```

```

        activation='relu',
        input_shape=input_shape))
model_1.add(Conv2D(32, (2, 2), activation='relu'))
model_1.add(MaxPooling2D(pool_size=(3, 3)))
model_1.add(BatchNormalization())
model_1.add(Dropout(0.25))
model_1.add(Flatten())
model_1.add(Dense(111, activation='relu'))
model_1.add(Dense(num_classes, activation='softmax'))

model_1.compile(loss=keras.losses.categorical_crossentropy,
                optimizer=keras.optimizers.Adadelta(),
                metrics=['accuracy'])

model_1.fit(x_train, y_train,
            batch_size=batch_size,
            epochs=epochs,
            verbose=1,
            validation_data=(x_test, y_test))
score_1 = model_1.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score_1[0])
print('Test accuracy:', score_1[1])

```

W0809 16:04:03.157577 2872 deprecation_wrapper.py:119] From C:\Users\ACER\Anaconda3\lib\site-p

Train on 60000 samples, validate on 10000 samples

```

Epoch 1/12
60000/60000 [=====] - 37s 618us/step - loss: 0.1461 - acc: 0.9548 - va
Epoch 2/12
60000/60000 [=====] - 40s 663us/step - loss: 0.0490 - acc: 0.9848 - va
Epoch 3/12
60000/60000 [=====] - 39s 656us/step - loss: 0.0345 - acc: 0.9890 - va
Epoch 4/12
60000/60000 [=====] - 40s 663us/step - loss: 0.0246 - acc: 0.9921 - va
Epoch 5/12
60000/60000 [=====] - 40s 666us/step - loss: 0.0191 - acc: 0.9940 - va
Epoch 6/12
60000/60000 [=====] - 39s 649us/step - loss: 0.0162 - acc: 0.9944 - va
Epoch 7/12
60000/60000 [=====] - 40s 670us/step - loss: 0.0131 - acc: 0.9960 - va
Epoch 8/12
60000/60000 [=====] - 41s 678us/step - loss: 0.0112 - acc: 0.9964 - va
Epoch 9/12
60000/60000 [=====] - 39s 654us/step - loss: 0.0096 - acc: 0.9968 - va
Epoch 10/12
60000/60000 [=====] - 41s 675us/step - loss: 0.0081 - acc: 0.9974 - va

```

```

Epoch 11/12
60000/60000 [=====] - 40s 672us/step - loss: 0.0070 - acc: 0.9977 - va
Epoch 12/12
60000/60000 [=====] - 39s 657us/step - loss: 0.0059 - acc: 0.9983 - va
Test loss: 0.03634111389952359
Test accuracy: 0.9902

```

0.0.2 Example 2

```
In [11]: epochs = 10
```

```

model_2 = Sequential()
model_2.add(Conv2D(32, kernel_size=(5, 5),
                  activation='relu',
                  input_shape=input_shape))
model_2.add(Conv2D(64, (5, 5), activation='relu'))
model_2.add(Conv2D(32, (5, 5), activation='relu'))
model_2.add(MaxPooling2D(pool_size=(3, 3)))
model_2.add(Flatten())
model_2.add(Dense(128, activation='relu'))
model_2.add(Dropout(0.5))
model_2.add(BatchNormalization())
model_2.add(Dropout(0.30))
model_2.add(Dense(num_classes, activation='softmax'))

model_2.compile(loss=keras.losses.categorical_crossentropy,
                optimizer=keras.optimizers.Adadelta(),
                metrics=['accuracy'])

model_2.fit(x_train, y_train,
            batch_size=batch_size,
            epochs=epochs,
            verbose=1,
            validation_data=(x_test, y_test))
score_2 = model_2.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score_2[0])
print('Test accuracy:', score_2[1])

```

Train on 60000 samples, validate on 10000 samples

```

Epoch 1/10
60000/60000 [=====] - 153s 3ms/step - loss: 0.2875 - acc: 0.9204 - va
Epoch 2/10
60000/60000 [=====] - 185s 3ms/step - loss: 0.0909 - acc: 0.9750 - va
Epoch 3/10
60000/60000 [=====] - 189s 3ms/step - loss: 0.0649 - acc: 0.9825 - va
Epoch 4/10
60000/60000 [=====] - 188s 3ms/step - loss: 0.0532 - acc: 0.9850 - va

```

```

Epoch 5/10
60000/60000 [=====] - 161s 3ms/step - loss: 0.0447 - acc: 0.9874 - va
Epoch 6/10
60000/60000 [=====] - 143s 2ms/step - loss: 0.0383 - acc: 0.9892 - va
Epoch 7/10
60000/60000 [=====] - 145s 2ms/step - loss: 0.0339 - acc: 0.9906 - va
Epoch 8/10
60000/60000 [=====] - 143s 2ms/step - loss: 0.0323 - acc: 0.9906 - va
Epoch 9/10
60000/60000 [=====] - 144s 2ms/step - loss: 0.0277 - acc: 0.9921 - va
Epoch 10/10
60000/60000 [=====] - 145s 2ms/step - loss: 0.0257 - acc: 0.9930 - va
Test loss: 0.015346395371133803
Test accuracy: 0.9957

```

0.0.3 Example 3

In [12]: epochs = 11

```

model_3 = Sequential()
model_3.add(Conv2D(64, kernel_size=(6, 6),
                  activation='relu',
                  input_shape=input_shape))
model_3.add(Conv2D(32, (6, 6), activation='relu'))
model_3.add(MaxPooling2D(pool_size=(2, 2)))
model_3.add(Dropout(0.25))
model_3.add(Flatten())
model_3.add(Dense(128, activation='relu'))
model_3.add(Dropout(0.5))
model_3.add(BatchNormalization())
model_3.add(Dense(num_classes, activation='softmax'))

model_3.compile(loss=keras.losses.categorical_crossentropy,
               optimizer=keras.optimizers.Adam(),
               metrics=['accuracy'])

model_3.fit(x_train, y_train,
          batch_size=batch_size,
          epochs=epochs,
          verbose=1,
          validation_data=(x_test, y_test))
score_3 = model_3.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score_3[0])
print('Test accuracy:', score_3[1])

```

Train on 60000 samples, validate on 10000 samples
Epoch 1/11

```

60000/60000 [=====] - 133s 2ms/step - loss: 0.2035 - acc: 0.9440 - va
Epoch 2/11
60000/60000 [=====] - 135s 2ms/step - loss: 0.0680 - acc: 0.9806 - va
Epoch 3/11
60000/60000 [=====] - 129s 2ms/step - loss: 0.0490 - acc: 0.9860 - va
Epoch 4/11
60000/60000 [=====] - 128s 2ms/step - loss: 0.0405 - acc: 0.9882 - va
Epoch 5/11
60000/60000 [=====] - 128s 2ms/step - loss: 0.0361 - acc: 0.9893 - va
Epoch 6/11
60000/60000 [=====] - 127s 2ms/step - loss: 0.0291 - acc: 0.9915 - va
Epoch 7/11
60000/60000 [=====] - 127s 2ms/step - loss: 0.0271 - acc: 0.9916 - va
Epoch 8/11
60000/60000 [=====] - 130s 2ms/step - loss: 0.0233 - acc: 0.9927 - va
Epoch 9/11
60000/60000 [=====] - 127s 2ms/step - loss: 0.0213 - acc: 0.9938 - va
Epoch 10/11
60000/60000 [=====] - 129s 2ms/step - loss: 0.0210 - acc: 0.9937 - va
Epoch 11/11
60000/60000 [=====] - 146s 2ms/step - loss: 0.0179 - acc: 0.9944 - va
Test loss: 0.018701013443106058
Test accuracy: 0.9949

```

0.0.4 Example 4

In [3]: *# With activation = softmax, batch normalization and dropout layer*

```

epochs = 10

model_4 = Sequential()
model_4.add(Conv2D(32, kernel_size=(4, 4),
                  activation='sigmoid',
                  input_shape=input_shape))
model_4.add(Conv2D(64, (4, 4), activation='sigmoid'))
model_4.add(MaxPooling2D(pool_size=(3, 3)))
model_4.add(Dropout(0.70))
model_4.add(Flatten())
model_4.add(Dense(128, activation='sigmoid'))
model_4.add(BatchNormalization(epsilon=0.001))
model_4.add(Dropout(0.30))
model_4.add(Dense(num_classes, activation='softmax'))

model_4.compile(loss=keras.losses.categorical_crossentropy,
               optimizer=keras.optimizers.Adadelta(),
               metrics=['accuracy'])

```

```

model_4.fit(x_train, y_train,
            batch_size=batch_size,
            epochs=epochs,
            verbose=1,
            validation_data=(x_test, y_test))
score_4 = model_4.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score_4[0])
print('Test accuracy:', score_4[1])

```

WARNING: Logging before flag parsing goes to stderr.

W0809 14:12:37.660238 2872 deprecation_wrapper.py:119] From C:\Users\ACER\Anaconda3\lib\site-packages\tensorflow\python\ops\nn_ops.py:4224: Large dropout rate: 0.7 (>0.5). In TensorFlow 2.x,

W0809 14:12:37.691510 2872 deprecation_wrapper.py:119] From C:\Users\ACER\Anaconda3\lib\site-packages\tensorflow\python\ops\nn_ops.py:4224: Large dropout rate: 0.7 (>0.5). In TensorFlow 2.x,

W0809 14:12:37.707100 2872 deprecation_wrapper.py:119] From C:\Users\ACER\Anaconda3\lib\site-packages\tensorflow\python\ops\nn_ops.py:4224: Large dropout rate: 0.7 (>0.5). In TensorFlow 2.x,

W0809 14:12:37.753999 2872 deprecation_wrapper.py:119] From C:\Users\ACER\Anaconda3\lib\site-packages\tensorflow\python\ops\nn_ops.py:4224: Large dropout rate: 0.7 (>0.5). In TensorFlow 2.x,

W0809 14:12:37.769587 2872 deprecation_wrapper.py:119] From C:\Users\ACER\Anaconda3\lib\site-packages\tensorflow\python\ops\nn_ops.py:4224: Large dropout rate: 0.7 (>0.5). In TensorFlow 2.x,

W0809 14:12:37.785243 2872 deprecation.py:506] From C:\Users\ACER\Anaconda3\lib\site-packages\tensorflow\python\ops\nn_ops.py:4224: Large dropout rate: 0.7 (>0.5). In TensorFlow 2.x,

Instructions for updating:

Please use `rate` instead of `keep_prob`. Rate should be set to `rate = 1 - keep_prob`.

W0809 14:12:37.785243 2872 nn_ops.py:4224] Large dropout rate: 0.7 (>0.5). In TensorFlow 2.x,

W0809 14:12:38.035182 2872 deprecation_wrapper.py:119] From C:\Users\ACER\Anaconda3\lib\site-packages\tensorflow\python\ops\nn_ops.py:4224: Large dropout rate: 0.7 (>0.5). In TensorFlow 2.x,

W0809 14:12:38.050771 2872 deprecation_wrapper.py:119] From C:\Users\ACER\Anaconda3\lib\site-packages\tensorflow\python\ops\nn_ops.py:4224: Large dropout rate: 0.7 (>0.5). In TensorFlow 2.x,

W0809 14:12:38.222638 2872 deprecation.py:323] From C:\Users\ACER\Anaconda3\lib\site-packages\tensorflow\python\ops\nn_ops.py:4224: Large dropout rate: 0.7 (>0.5). In TensorFlow 2.x,

Instructions for updating:

Use tf.where in 2.0, which has the same broadcast rule as np.where

Train on 60000 samples, validate on 10000 samples

Epoch 1/10

60000/60000 [=====] - 118s 2ms/step - loss: 1.8476 - acc: 0.3396 - val_loss: 1.8476 - val_acc: 0.3396

Epoch 2/10

60000/60000 [=====] - 118s 2ms/step - loss: 0.2793 - acc: 0.9129 - val_loss: 0.2793 - val_acc: 0.9129

Epoch 3/10

60000/60000 [=====] - 118s 2ms/step - loss: 0.1542 - acc: 0.9513 - val_loss: 0.1542 - val_acc: 0.9513

Epoch 4/10

60000/60000 [=====] - 116s 2ms/step - loss: 0.1222 - acc: 0.9622 - val_loss: 0.1222 - val_acc: 0.9622

Epoch 5/10

60000/60000 [=====] - 119s 2ms/step - loss: 0.1033 - acc: 0.9673 - val_loss: 0.1033 - val_acc: 0.9673

Epoch 6/10

60000/60000 [=====] - 123s 2ms/step - loss: 0.0941 - acc: 0.9703 - val_loss: 0.0941 - val_acc: 0.9703

Epoch 7/10

60000/60000 [=====] - 123s 2ms/step - loss: 0.0834 - acc: 0.9738 - val_loss: 0.0834 - val_acc: 0.9738


```

Epoch 8/10
60000/60000 [=====] - 118s 2ms/step - loss: 0.0784 - acc: 0.9759 - va
Epoch 9/10
60000/60000 [=====] - 120s 2ms/step - loss: 0.0756 - acc: 0.9768 - va
Epoch 10/10
60000/60000 [=====] - 119s 2ms/step - loss: 0.0703 - acc: 0.9778 - va
Test loss: 0.03086712787185097
Test accuracy: 0.9897

```

0.0.5 Example 5

In [4]: # With activation = tanh, without batch normalization, with dropout, with adam optimizer

```

epochs = 10

model_5 = Sequential()
model_5.add(Conv2D(32, kernel_size=(3, 3),
                  activation='tanh',
                  input_shape=input_shape))
model_5.add(Conv2D(32, (3, 3), activation='tanh'))
model_5.add(MaxPooling2D(pool_size=(3, 3)))
model_5.add(Dropout(0.44))
model_5.add(Flatten())
model_5.add(Dense(128, activation='sigmoid'))
model_5.add(Dropout(0.30))
model_5.add(Dense(num_classes, activation='softmax'))

model_5.compile(loss=keras.losses.categorical_crossentropy,
               optimizer=keras.optimizers.Adam(),
               metrics=['accuracy'])

model_5.fit(x_train, y_train,
          batch_size=batch_size,
          epochs=epochs,
          verbose=1,
          validation_data=(x_test, y_test))
score_5 = model_5.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score_5[0])
print('Test accuracy:', score_5[1])

```

Train on 60000 samples, validate on 10000 samples

```

Epoch 1/10
60000/60000 [=====] - 63s 1ms/step - loss: 0.3902 - acc: 0.8877 - val
Epoch 2/10
60000/60000 [=====] - 64s 1ms/step - loss: 0.1223 - acc: 0.9660 - val
Epoch 3/10
60000/60000 [=====] - 63s 1ms/step - loss: 0.0921 - acc: 0.9726 - val

```

```

Epoch 4/10
60000/60000 [=====] - 64s 1ms/step - loss: 0.0792 - acc: 0.9763 - val.
Epoch 5/10
60000/60000 [=====] - 63s 1ms/step - loss: 0.0707 - acc: 0.9786 - val.
Epoch 6/10
60000/60000 [=====] - 63s 1ms/step - loss: 0.0606 - acc: 0.9817 - val.
Epoch 7/10
60000/60000 [=====] - 63s 1ms/step - loss: 0.0577 - acc: 0.9824 - val.
Epoch 8/10
60000/60000 [=====] - 63s 1ms/step - loss: 0.0516 - acc: 0.9839 - val.
Epoch 9/10
60000/60000 [=====] - 62s 1ms/step - loss: 0.0484 - acc: 0.9849 - val.
Epoch 10/10
60000/60000 [=====] - 62s 1ms/step - loss: 0.0462 - acc: 0.9855 - val.
Test loss: 0.03044679508442059
Test accuracy: 0.9902

```

0.0.6 Example 6

In [5]: *# Activation = selu & softmax , without Batch normalization and dropout, Initializer =
bias_initializer, Optimizer = Adagrad*

```

epochs = 11

model_6 = Sequential()
model_6.add(Conv2D(64, kernel_size=(5, 5),
                  activation='relu',
                  input_shape=input_shape))
model_6.add(Dense(64,
                  kernel_initializer='random_uniform',
                  bias_initializer='zeros'))
model_6.add(Conv2D(32, (4, 4), activation='softmax'))
model_6.add(MaxPooling2D(pool_size=(3, 3)))
model_6.add(Flatten())
model_6.add(Dense(128, activation='relu'))
model_6.add(Dense(num_classes, activation='softmax'))

model_6.compile(loss=keras.losses.categorical_crossentropy,
                optimizer=keras.optimizers.Adagrad(),
                metrics=['accuracy'])

model_6.fit(x_train, y_train,
            batch_size=batch_size,
            epochs=epochs,
            verbose=1,
            validation_data=(x_test, y_test))
score_6 = model_6.evaluate(x_test, y_test, verbose=0)

```

```

print('Test loss:', score_6[0])
print('Test accuracy:', score_6[1])

```

Train on 60000 samples, validate on 10000 samples

```

Epoch 1/11
60000/60000 [=====] - 194s 3ms/step - loss: 0.2787 - acc: 0.9100 - va
Epoch 2/11
60000/60000 [=====] - 192s 3ms/step - loss: 0.0592 - acc: 0.9824 - va
Epoch 3/11
60000/60000 [=====] - 195s 3ms/step - loss: 0.0469 - acc: 0.9859 - va
Epoch 4/11
60000/60000 [=====] - 191s 3ms/step - loss: 0.0396 - acc: 0.9885 - va
Epoch 5/11
60000/60000 [=====] - 189s 3ms/step - loss: 0.0345 - acc: 0.9905 - va
Epoch 6/11
60000/60000 [=====] - 189s 3ms/step - loss: 0.0310 - acc: 0.9913 - va
Epoch 7/11
60000/60000 [=====] - 189s 3ms/step - loss: 0.0282 - acc: 0.9920 - va
Epoch 8/11
60000/60000 [=====] - 191s 3ms/step - loss: 0.0259 - acc: 0.9928 - va
Epoch 9/11
60000/60000 [=====] - 192s 3ms/step - loss: 0.0240 - acc: 0.9935 - va
Epoch 10/11
60000/60000 [=====] - 208s 3ms/step - loss: 0.0223 - acc: 0.9942 - va
Epoch 11/11
60000/60000 [=====] - 203s 3ms/step - loss: 0.0209 - acc: 0.9944 - va
Test loss: 0.03338987519220682
Test accuracy: 0.988

```

0.0.7 Example 7

```

In [7]: #Activation = softplus, Initializer = random_normal, with dropout and batch normalizat
from keras import initializers

```

```

epochs = 10

model_7 = Sequential()
model_7.add(Conv2D(32, kernel_size=(3, 3),
                  activation='softplus',
                  input_shape=input_shape))
model_7.add(Conv2D(32, (4, 4), activation='softplus'))
model_7.add(Dense(32,
                  kernel_initializer=initializers.random_normal(stddev=0.01)))
model_7.add(MaxPooling2D(pool_size=(3, 3)))
model_7.add(Dropout(0.50))
model_7.add(Flatten())
model_7.add(Dense(128, activation='softplus'))

```

```

model_7.add(BatchNormalization(epsilon=0.001))
model_7.add(Dropout(0.30))
model_7.add(Dense(num_classes, activation='softmax'))

model_7.compile(loss=keras.losses.categorical_crossentropy,
                 optimizer=keras.optimizers.Adamax(),
                 metrics=['accuracy'])

model_7.fit(x_train, y_train,
            batch_size=batch_size,
            epochs=epochs,
            verbose=1,
            validation_data=(x_test, y_test))
score_7 = model_7.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score_7[0])
print('Test accuracy:', score_7[1])

```

Train on 60000 samples, validate on 10000 samples

```

Epoch 1/10
60000/60000 [=====] - 114s 2ms/step - loss: 0.6562 - acc: 0.7877 - va
Epoch 2/10
60000/60000 [=====] - 109s 2ms/step - loss: 0.2479 - acc: 0.9251 - va
Epoch 3/10
60000/60000 [=====] - 108s 2ms/step - loss: 0.1668 - acc: 0.9492 - va
Epoch 4/10
60000/60000 [=====] - 105s 2ms/step - loss: 0.1280 - acc: 0.9607 - va
Epoch 5/10
60000/60000 [=====] - 107s 2ms/step - loss: 0.1124 - acc: 0.9657 - va
Epoch 6/10
60000/60000 [=====] - 108s 2ms/step - loss: 0.1011 - acc: 0.9692 - va
Epoch 7/10
60000/60000 [=====] - 109s 2ms/step - loss: 0.0914 - acc: 0.9715 - va
Epoch 8/10
60000/60000 [=====] - 109s 2ms/step - loss: 0.0839 - acc: 0.9738 - va
Epoch 9/10
60000/60000 [=====] - 107s 2ms/step - loss: 0.0793 - acc: 0.9754 - va
Epoch 10/10
60000/60000 [=====] - 109s 2ms/step - loss: 0.0748 - acc: 0.9765 - va
Test loss: 0.04983391810744069
Test accuracy: 0.9841

```

0.0.8 Prettytable

```
In [13]: from prettytable import PrettyTable
```

```

number= [1,2,3,4,5,6,7]
name= ["Model 1","Model 2","Model 3","Model 4","Model 5","Model 6","Model 7"]

```

```

loss= [score_1[0],score_2[0],score_3[0],score_4[0],score_5[0],score_6[0],score_7[0]]
acc= [score_1[1],score_2[1],score_3[1],score_4[1],score_5[1],score_6[1],score_7[1]]

#Initialize Prettytable
ptable = PrettyTable()
ptable.add_column("Index", number)
ptable.add_column("Model", name)
ptable.add_column("Test Loss", loss)
ptable.add_column("Test Accuracy", acc)
print(ptable)

```

Index	Model	Test Loss	Test Accuracy
1	Model 1	0.03634111389952359	0.9902
2	Model 2	0.015346395371133803	0.9957
3	Model 3	0.018701013443106058	0.9949
4	Model 4	0.03086712787185097	0.9897
5	Model 5	0.03044679508442059	0.9902
6	Model 6	0.03338987519220682	0.988
7	Model 7	0.04983391810744069	0.9841

0.0.9 Conclusions

1. For MNIST dataset, we build Convolutional Neural Networks using Conv2D.
2. We have used different kernel sizes like (3,3),(4,4),(5,5) etc.
3. Different layers are used in all the examples.
4. Some models with or without Batch normalization and Dropout layers to check the performance.
5. Activation methods like softplus, relu, selu, sigmoid etc. are used
6. Optimizers like Adagrad, Adam, Adamax, Adadelata are used.
7. Test accuracy is almost same in all cases.

In []: