

# Taxi demand prediction in New York City

In [1]:

```
#Importing Libraries
# pip3 install graphviz
#pip3 install dask
#pip3 install toolz
#pip3 install cloudpickle
# https://www.youtube.com/watch?v=ieW3G7ZzRZ0
# https://github.com/dask/dask-tutorial
# please do go through this python notebook: https://github.com/dask/dask-tutorial/blob/master/07_dataframe.ipynb
import dask.dataframe as dd #similar to pandas

import pandas as pd #pandas to create small dataframes

# pip3 install folium
# if this doesnt work refere install_folium.JPG in drive
import folium #open street map

# unix time: https://www.unixtimestamp.com/
import datetime #Convert to unix time

import time #Convert to unix time

# if numpy is not installed already : pip3 install numpy
import numpy as np #Do aritmetic operations on arrays

# matplotlib: used to plot graphs
import matplotlib
# matplotlib.use('nbagg') : matplotlib uses this protocall which makes plots more user intractive
like zoom in and zoom out
matplotlib.use('nbagg')
import matplotlib.pyplot as plt
import seaborn as sns #Plots
from matplotlib import rcParams #Size of plots

# this lib is used while we calculate the straight line distance between two (lat,lon) pairs in miles
import gpxpy.geo #Get the haversine distance

from sklearn.cluster import MiniBatchKMeans, KMeans #Clustering
import math
import pickle
import os
import matplotlib.pyplot as plt

# download mingw: https://mingw-w64.org/doku.php/download/mingw-builds
# install it in your system and keep the path, mingw_path ='installed path'
mingw_path = 'C:\\\\Program Files\\\\mingw-w64\\\\x86_64-5.3.0-posix-seh-rt_v4-rev0\\\\mingw64\\\\bin'
os.environ['PATH'] = mingw_path + ';' + os.environ['PATH']

# to install xgboost: pip3 install xgboost
# if it didnt happen check install_xgboost.JPG
import xgboost as xgb

# to install sklearn: pip install -U scikit-learn
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error
from sklearn.metrics import mean_absolute_error
import warnings
warnings.filterwarnings("ignore")
```

## Data Information

Get the data from : [http://www.nyc.gov/html/tlc/html/about/trip\\_record\\_data.shtml](http://www.nyc.gov/html/tlc/html/about/trip_record_data.shtml) (2016 data) The data used in the attached datasets

were collected and provided to the NYC Taxi and Limousine Commission (TLC)

## Information on taxis:

### **Yellow Taxi: Yellow Medallion Taxicabs**

These are the famous NYC yellow taxis that provide transportation exclusively through street-hails. The number of taxicabs is limited by a finite number of medallions issued by the TLC. You access this mode of transportation by standing in the street and hailing an available taxi with your hand. The pickups are not pre-arranged.

### **For Hire Vehicles (FHV)s**

FHV transportation is accessed by a pre-arrangement with a dispatcher or limo company. These FHVs are not permitted to pick up passengers via street hails, as those rides are not considered pre-arranged.

### **Green Taxi: Street Hail Livery (SHL)**

The SHL program will allow livery vehicle owners to license and outfit their vehicles with green borough taxi branding, meters, credit card machines, and ultimately the right to accept street hails in addition to pre-arranged rides.

Credits: Quora

#### **Footnote:**

In the given notebook we are considering only the yellow taxis for the time period between Jan - Mar 2015 & Jan - Mar 2016

## Data Collection

We Have collected all yellow taxi trips data from jan-2015 to dec-2016(Will be using only 2015 data)

file name	file name size	number of records	number of features
yellow_tripdata_2016-01	1. 59G	10906858	19
yellow_tripdata_2016-02	1. 66G	11382049	19
yellow_tripdata_2016-03	1. 78G	12210952	19
yellow_tripdata_2016-04	1. 74G	11934338	19
yellow_tripdata_2016-05	1. 73G	11836853	19
yellow_tripdata_2016-06	1. 62G	11135470	19
yellow_tripdata_2016-07	884Mb	10294080	17
yellow_tripdata_2016-08	854Mb	9942263	17
yellow_tripdata_2016-09	870Mb	10116018	17
yellow_tripdata_2016-10	933Mb	10854626	17
yellow_tripdata_2016-11	868Mb	10102128	17
yellow_tripdata_2016-12	897Mb	10449408	17
yellow_tripdata_2015-01	1.84Gb	12748986	19
yellow_tripdata_2015-02	1.81Gb	12450521	19
yellow_tripdata_2015-03	1.94Gb	13351609	19
yellow_tripdata_2015-04	1.90Gb	13071789	19
yellow_tripdata_2015-05	1.91Gb	13158262	19
yellow_tripdata_2015-06	1.79Gb	12324935	19
yellow_tripdata_2015-07	1.68Gb	11562783	19
yellow_tripdata_2015-08	1.62Gb	11130304	19
yellow_tripdata_2015-09	1.63Gb	11225063	19
yellow_tripdata_2015-10	1.79Gb	12315488	19
yellow_tripdata_2015-11	1.65Gb	11312676	19
yellow_tripdata_2015-12	1.67Gb	11460573	19

... [2] \*

```
#Looking at the features
# dask dataframe : # https://github.com/dask/dask-tutorial/blob/master/07_dataframe.ipynb
month = dd.read_csv('yellow_tripdata_2015-01.csv')
print(month.columns)

Index(['VendorID', 'tpep_pickup_datetime', 'tpep_dropoff_datetime',
       'passenger_count', 'trip_distance', 'pickup_longitude',
       'pickup_latitude', 'RateCodeID', 'store_and_fwd_flag',
       'dropoff_longitude', 'dropoff_latitude', 'payment_type', 'fare_amount',
       'extra', 'mta_tax', 'tip_amount', 'tolls_amount',
       'improvement_surcharge', 'total_amount'],
      dtype='object')
```

In [3]:

```
type(month)
```

Out[3]:


```
dask.dataframe.core.DataFrame
```

In [4]:

```
# However unlike Pandas, operations on dask.dataframes don't trigger immediate computation,
# instead they add key-value pairs to an underlying Dask graph. Recall that in the diagram below,
# circles are operations and rectangles are results.

# to see the visualization you need to install graphviz
# pip3 install graphviz if this doesnt work please check the install_graphviz.jpg in the drive
month.visualize()
```


Out[4]:



In [5]:

```
month.fare_amount.sum().visualize()
```

Out[5]:



## Features in the dataset:

Field Name	Description
VendorID	A code indicating the TPEP provider that provided the record. 1. Creative Mobile Technologies 2. VeriFone Inc.
tpep_pickup_datetime	The date and time when the meter was engaged.
tpep_dropoff_datetime	The date and time when the meter was disengaged.

Passenger_count	The number of passengers in the vehicle. This is a driver-entered value.	
Trip_distance	The elapsed trip distance in miles reported by the taximeter.	
Pickup_longitude	Longitude where the meter was engaged.	
Pickup_latitude	Latitude where the meter was engaged.	
RateCodeID	1. 2. 3. 4. 5. 6.	
	The final rate code in effect at the end of the trip. Standard rate JFK Newark Nassau or Westchester Negotiated fare Group ride	
Store_and_fwd_flag	This flag indicates whether the trip record was held in vehicle memory before sending to the vendor aka "store and forward," because the vehicle did not have a connection to the server. Y= store and forward trip N= not a store and forward trip	
Dropoff_longitude	Longitude where the meter was disengaged.	
Dropoff_latitude	Latitude where the meter was disengaged.	
Payment_type	1. 2. 3. 4. 5. 6.	
	A numeric code signifying how the passenger paid for the trip. Credit card Cash No charge Dispute Unknown Voided trip	
Fare_amount	The time-and-distance fare calculated by the meter.	
Extra	Miscellaneous extras and surcharges. Currently, this only includes the 0.50 and 1 rush hour and overnight charges.	
MTA_tax	0.50 MTA tax that is automatically triggered based on the metered rate in use.	
Improvement_surcharge	0.30 improvement surcharge assessed trips at the flag drop. The improvement surcharge began being levied in 2015.	
Tip_amount	Tip amount – This field is automatically populated for credit card tips. Cash tips are not included.	
Tolls_amount	Total amount of all tolls paid in trip.	
Total_amount	The total amount charged to passengers. Does not include cash tips.	

## ML Problem Formulation

### Time-series forecasting and Regression

- To find number of pickups, given location coordinates(latitude and longitude) and time, in the query region and surrounding regions.

To solve the above we would be using data collected in Jan - Mar 2015 to predict the pickups in Jan - Mar 2016.

## Performance metrics

1. Mean Absolute percentage error.
2. Mean Squared error.

## Data Cleaning

In this section we will be doing univariate analysis and removing outlier/illegitimate values which may be caused due to some error

In [4]:

```
#table below shows few datapoints along with all our features
month.head(5)
```

Out [4]:

VendorID	tpep_pickup_datetime	tpep_dropoff_datetime	passenger_count	trip_distance	pickup_longitude	pickup_latitude	RateCode
0	2 2015-01-15 19:05:39	2015-01-15 19:23:42	1	1.59	-73.993896	40.750111	
1	1 2015-01-10 20:33:38	2015-01-10 20:53:28	1	3.30	-74.001648	40.724243	
2	1 2015-01-10 20:33:38	2015-01-10 20:43:41	1	1.80	-73.963341	40.802788	
3	1 2015-01-10 20:33:39	2015-01-10 20:35:31	1	0.50	-74.009087	40.713818	
4	1 2015-01-10 20:33:39	2015-01-10 20:52:58	1	3.00	-73.971176	40.762428	

VendorID	tpep_pickup_datetime	tpep_dropoff_datetime	passenger_count	trip_distance	pickup_longitude	pickup_latitude	RateCode

## 1. Pickup Latitude and Pickup Longitude

It is inferred from the source <https://www.flickr.com/places/info/2459115> that New York is bounded by the location coordinates(lat,long) - (40.5774, -74.15) & (40.9176,-73.7004) so hence any coordinates not within these coordinates are not considered by us as we are only concerned with pickups which originate within New York.

In [5]:

```
# Plotting pickup coordinates which are outside the bounding box of New-York
# we will collect all the points outside the bounding box of newyork city to outlier_locations
# dask dataframe
outlier_locations = month[((month.pickup_longitude <= -74.15) | (month.pickup_latitude <= 40.5774) |
\ 
    (month.pickup_longitude >= -73.7004) | (month.pickup_latitude >= 40.9176))]

# creating a map with the a base location
# read more about the folium here: http://folium.readthedocs.io/en/latest/quickstart.html

# note: you dont need to remember any of these, you dont need indeepth knowledge on these maps and
```



In [8]:

```
type(sample_locations)
```

Out [8]:

```
pandas.core.frame.DataFrame
```

**Observation:-** As you can see above that there are some points just outside the boundary but there are a few that are in either South america, Mexico or Canada


## 2. Dropoff Latitude & Dropoff Longitude

It is inferred from the source <https://www.flickr.com/places/info/2459115> that New York is bounded by the location coordinates(lat,long) - (40.5774, -74.15) & (40.9176,-73.7004) so hence any cordinates not within these cordinates are not considered by us as we are only concerned with dropoffs which are within New York.

In [6]:

```
# Plotting dropoff cordinates which are outside the bounding box of New-York
# we will collect all the points outside the bounding box of newyork city to outlier_locations
outlier_locations = month[((month.dropoff_longitude <= -74.15) | (month.dropoff_latitude <= 40.5774)
) | \
(month.dropoff_longitude >= -73.7004) | (month.dropoff_latitude >= 40.9176))]

# creating a map with the a base location
# read more about the folium here: http://folium.readthedocs.io/en/latest/quickstart.html
```



**Observation:-** The observations here are similar to those obtained while analysing pickup latitude and longitude

## 3. Trip Durations:

According to NYC Taxi & Limousine Commission Regulations **the maximum allowed trip duration in a 24 hour interval is 12 hours.**

In [3]:

```
#The timestamps are converted to unix so as to get duration(trip-time) & speed also pickup-times in unix are used while binning

# in our data we have time in the format "YYYY-MM-DD HH:MM:SS" we convert this string to python time format and then into unix timestamp
# https://stackoverflow.com/a/27914405
def convert_to_unix(s):
    return time.mktime(datetime.datetime.strptime(s, "%Y-%m-%d %H:%M:%S").timetuple())

# we return a data frame which contains the columns
# 1.'passenger_count' : self explanatory
# 2.'trip_distance' : self explanatory
# 3.'pickup_longitude' : self explanatory
# 4.'pickup_latitude' : self explanatory
# 5.'dropoff_longitude' : self explanatory
# 6.'dropoff_latitude' : self explanatory
# 7.'total_amount' : total fare that was paid
# 8.'trip_times' : duration of each trip
# 9.'pickup_time' : pickup time converted into unix time
# 10.'Speed' : velocity of each trip
def return_with_trip_times(month):
    duration = month[['tpep_pickup_datetime', 'tpep_dropoff_datetime']].compute()
    #pickups and dropoffs to unix time
    duration_pickup = [convert_to_unix(x) for x in duration['tpep_pickup_datetime'].values]
    duration_drop = [convert_to_unix(x) for x in duration['tpep_dropoff_datetime'].values]
    #calculate duration of trips
    durations = (np.array(duration_drop) - np.array(duration_pickup))/float(60)

    #append durations of trips and speed in miles/hr to a new dataframe
    new_frame =
month[['passenger_count', 'trip_distance', 'pickup_longitude', 'pickup_latitude', 'dropoff_longitude', 'dropoff_latitude', 'total_amount']].compute()

    new_frame['trip_times'] = durations
    new_frame['pickup_time'] = duration_pickup
    new_frame['Speed'] = 60*(new_frame['trip_distance'])/new_frame['trip_times'])


    return new_frame

# print(frame_with_durations.head())
# passenger_count trip_distance pickup_longitude pickup_latitude dropoff_longitude
dropoff_latitude total_amount trip_times pickup_time Speed
# 1 1.59 -73.993896 40.750111 -73.974785 40.750618
17.05 18.050000 1.421329e+09 5.285319
# 1 3.30 -74.001648 40.724243 -73.994415 40.759109
.80 19.833333 1.420902e+09 9.983193
# 1 1.80 -73.963341 40.802788 -73.951820 40.824413
10.80 10.050000 1.420902e+09 10.746269
# 1 0.50 -74.009087 40.713818 -74.004326 40.719986
4.80 1.866667 1.420902e+09 16.071429
# 1 3.00 -73.971176 40.762428 -74.004181 40.742653
6.30 19.316667 1.420902e+09 9.318378
frame_with_durations = return_with_trip_times(month)
```

In [11]:

```
# the skewed box plot shows us the presence of outliers
sns.boxplot(y="trip_time", data=frame_with_durations)
plt.show()
```





In [12]:

```
#calculating 0-100th percentile to find a the correct percentile value for removal of outliers
for i in range(0,100,10):
    var = frame_with_durations["trip_times"].values
    var = np.sort(var, axis = None)
    print("{} percentile value is {}".format(i,var[int(len(var)*(float(i)/100))]))
print ("100 percentile value is ",var[-1])
```

```
0 percentile value is -1211.0166666666667
10 percentile value is 3.833333333333335
20 percentile value is 5.38333333333334
30 percentile value is 6.816666666666666
40 percentile value is 8.3
50 percentile value is 9.95
60 percentile value is 11.866666666666667
70 percentile value is 14.28333333333333
80 percentile value is 17.63333333333333
90 percentile value is 23.45
100 percentile value is 548555.633333333
```

In [13]:

```
print(var)
```

```
[ -1211.01666667 -1200.           -985.5          ... 548492.93333333
 548530.88333333 548555.63333333]
```

In [14]:

```
#looking further from the 99th percecntile
for i in range(90,100):
    var = frame_with_durations["trip_times"].values
    var = np.sort(var, axis = None)
    print("{} percentile value is {}".format(i,var[int(len(var)*(float(i)/100))]))
print ("100 percentile value is ",var[-1])
```

```
90 percentile value is 23.45
91 percentile value is 24.35
92 percentile value is 25.38333333333333
93 percentile value is 26.55
94 percentile value is 27.93333333333334
95 percentile value is 29.58333333333332
96 percentile value is 31.68333333333334
97 percentile value is 34.46666666666667
98 percentile value is 38.71666666666667
99 percentile value is 46.75
100 percentile value is 548555.633333333
```


In [4]:

```
#removing data based on our analysis and mtc visualizations
```

```
#removing data based on our analysis and the regulations
frame_with_durations_modified=frame_with_durations[(frame_with_durations.trip_times>1) &
(frame_with_durations.trip_times<720)]
```


In [14]:


```
#box-plot after removal of outliers
sns.boxplot(y="trip_times", data =frame_with_durations_modified)
plt.show()
```



In [16]:

```
#pdf of trip-times after removing the outliers
sns.FacetGrid(frame_with_durations_modified,size=6) \
.map(sns.kdeplot,"trip_times") \
.add_legend();
plt.show();
```






In [16]:

```
#converting the values to log-values to check for log-normal
import math
frame_with_durations_modified['log_times']=[math.log(i) for i in frame_with_durations_modified['trip_times'].values]
```

In [17]:


```
#pdf of log-values
sns.FacetGrid(frame_with_durations_modified,size=6) \
    .map(sns.kdeplot,"log_times") \
    .add_legend();
plt.show();
```



In [30]:

```
#Q-Q plot for checking if trip-times is log-normal
import scipy
scipy.stats.probplot(frame_with_durations_modified['log_times'].values, plot=plt)
plt.show()
```


## Probability Plot



## 4. Speed

In [9]:

```
# check for any outliers in the data after trip duration outliers removed
# box-plot for speeds with outliers
frame_with_durations_modified['Speed'] =
60*(frame_with_durations_modified['trip_distance']/frame_with_durations_modified['trip_times'])
sns.boxplot(y="Speed", data=frame_with_durations_modified)
plt.show()
```



In [17]:

```
#calculating speed values at each percentile 0,10,20,30,40,50,60,70,80,90,100
for i in range(0,100,10):
    var =frame_with_durations_modified["Speed"].values
    var = np.sort(var,axis = None)
```

```
    print("{} percentile value is {}".format(i,var[int(len(var)*(float(i)/100))]))  
print("100 percentile value is ",var[-1])
```

```
0 percentile value is 0.0  
10 percentile value is 6.409495548961425  
20 percentile value is 7.80952380952381  
30 percentile value is 8.929133858267717  
40 percentile value is 9.98019801980198  
50 percentile value is 11.06865671641791  
60 percentile value is 12.286689419795222  
70 percentile value is 13.796407185628745  
80 percentile value is 15.963224893917962  
90 percentile value is 20.186915887850468  
100 percentile value is 192857142.85714284
```

In [18]:

```
#calculating speed values at each percntile 90,91,92,93,94,95,96,97,98,99,100  
for i in range(90,100):  
    var =frame_with_durations_modified["Speed"].values  
    var = np.sort(var, axis = None)  
    print("{} percentile value is {}".format(i,var[int(len(var)*(float(i)/100))]))  
print("100 percentile value is ",var[-1])
```

```
90 percentile value is 20.186915887850468  
91 percentile value is 20.91645569620253  
92 percentile value is 21.752988047808763  
93 percentile value is 22.721893491124263  
94 percentile value is 23.844155844155843  
95 percentile value is 25.182552504038775  
96 percentile value is 26.80851063829787  
97 percentile value is 28.84304932735426  
98 percentile value is 31.591128254580514  
99 percentile value is 35.7513566847558  
100 percentile value is 192857142.85714284
```

In [18]:

```
#calculating speed values at each percntile 99.0,99.1,99.2,99.3,99.4,99.5,99.6,99.7,99.8,99.9,100  
for i in np.arange(0.0, 1.0, 0.1):  
    var =frame_with_durations_modified["Speed"].values  
    var = np.sort(var, axis = None)  
    print("{} percentile value is {}".format(99+i,var[int(len(var)*(float(99+i)/100))]))  
print("100 percentile value is ",var[-1])
```

```
99.0 percentile value is 35.7513566847558  
99.1 percentile value is 36.31084727468969  
99.2 percentile value is 36.91470054446461  
99.3 percentile value is 37.588235294117645  
99.4 percentile value is 38.33035714285714  
99.5 percentile value is 39.17580340264651  
99.6 percentile value is 40.15384615384615  
99.7 percentile value is 41.338301043219076  
99.8 percentile value is 42.86631016042781  
99.9 percentile value is 45.3107822410148  
100 percentile value is 192857142.85714284
```

In [5]:

```
#removing further outliers based on the 99.9th percentile value  
frame_with_durations_modified=frame_with_durations[(frame_with_durations.Speed>0) &  
(frame_with_durations.Speed<45.31)]
```

In [6]:

```
#avg.speed of cabs in New-York  
sum(frame_with_durations_modified["Speed"]) / float(len(frame_with_durations_modified["Speed"]))
```

Out[6]:


```
12.450173996027528
```

The avg speed in Newyork speed is 12.45miles/hr, so a cab driver can travel 2 miles per 10min on avg.

#### 4. Trip Distance

In [21]:

```
# up to now we have removed the outliers based on trip durations and cab speeds
# lets try if there are any outliers in trip distances
# box-plot showing outliers in trip-distance values
sns.boxplot(y="trip_distance", data =frame_with_durations_modified)
plt.show()
```



In [22]:

```
#calculating trip distance values at each percentile 0,10,20,30,40,50,60,70,80,90,100
for i in range(0,100,10):
    var =frame_with_durations_modified["trip_distance"].values
    var = np.sort(var,axis = None)
    print("{} percentile value is {}".format(i,var[int(len(var)*(float(i)/100))]))
print("100 percentile value is ",var[-1])
```

```
0 percentile value is 0.01
10 percentile value is 0.66
20 percentile value is 0.9
30 percentile value is 1.1
40 percentile value is 1.39
50 percentile value is 1.69
60 percentile value is 2.07
70 percentile value is 2.6
80 percentile value is 3.6
90 percentile value is 5.97
100 percentile value is 258.9
```

In [23]:

```
#calculating trip distance values at each percentile 90,91,92,93,94,95,96,97,98,99,100
for i in range(90,100):
    var =frame_with_durations_modified["trip_distance"].values
    var = np.sort(var,axis = None)
    print("{} percentile value is {}".format(i,var[int(len(var)*(float(i)/100))]))
```

```
print("100 percentile value is ",var[-1])
```

```
90 percentile value is 5.97
91 percentile value is 6.45
92 percentile value is 7.07
93 percentile value is 7.85
94 percentile value is 8.72
95 percentile value is 9.6
96 percentile value is 10.6
97 percentile value is 12.1
98 percentile value is 16.03
99 percentile value is 18.17
100 percentile value is 258.9
```

In [23]:

```
#calculating trip distance values at each percentile
99.0,99.1,99.2,99.3,99.4,99.5,99.6,99.7,99.8,99.9,100
for i in np.arange(0.0, 1.0, 0.1):
    var =frame_with_durations_modified["trip_distance"].values
    var = np.sort(var, axis = None)
    print("{} percentile value is {}".format(99+i,var[int(len(var)*(float(99+i)/100))]))
print("100 percentile value is ",var[-1])
```


```
99.0 percentile value is 18.17
99.1 percentile value is 18.37
99.2 percentile value is 18.6
99.3 percentile value is 18.83
99.4 percentile value is 19.13
99.5 percentile value is 19.5
99.6 percentile value is 19.96
99.7 percentile value is 20.5
99.8 percentile value is 21.22
99.9 percentile value is 22.57
100 percentile value is 258.9
```

In [7]:

```
#removing further outliers based on the 99.9th percentile value
frame_with_durations_modified=frame_with_durations[(frame_with_durations.trip_distance>0) &
(frame_with_durations.trip_distance<23)]
```

In [25]:

```
#box-plot after removal of outliers
sns.boxplot(y="trip_distance", data = frame_with_durations_modified)
plt.show()
```






## 5. Total Fare

In [25]:

```
# up to now we have removed the outliers based on trip durations, cab speeds, and trip distances
# lets try if there are any outliers in based on the total_amount
# box-plot showing outliers in fare
sns.boxplot(y="total_amount", data =frame_with_durations_modified)
plt.show()
```



In [26]:

```
#calculating total fare amount values at each percentile 0,10,20,30,40,50,60,70,80,90,100
for i in range(0,100,10):
    var = frame_with_durations_modified["total_amount"].values
    var = np.sort(var, axis = None)
    print("{} percentile value is {}".format(i,var[int(len(var)*(float(i)/100))]))
print("100 percentile value is ",var[-1])
```

```
0 percentile value is -242.55
10 percentile value is 6.3
20 percentile value is 7.8
30 percentile value is 8.8
40 percentile value is 9.8
50 percentile value is 11.16
60 percentile value is 12.8
70 percentile value is 14.8
80 percentile value is 18.3
90 percentile value is 25.8
100 percentile value is 3950611.6
```

In [45]:

```
#calculating total fare amount values at each percentile 90,91,92,93,94,95,96,97,98,99,100
for i in range(90,100):
    var = frame_with_durations_modified["total_amount"].values
```

```
var = np.sort(var, axis = None)
print("{} percentile value is {}".format(i,var[int(len(var)*(float(i)/100))]))
print("100 percentile value is ",var[-1])
```

```
90 percentile value is 25.8
91 percentile value is 27.3
92 percentile value is 29.3
93 percentile value is 31.8
94 percentile value is 34.8
95 percentile value is 38.53
96 percentile value is 42.6
97 percentile value is 48.13
98 percentile value is 58.13
99 percentile value is 66.13
100 percentile value is 3950611.6
```

In [27]:

```
#calculating total fare amount values at each percentile
99.0,99.1,99.2,99.3,99.4,99.5,99.6,99.7,99.8,99.9,100
for i in np.arange(0.0, 1.0, 0.1):
    var = frame_with_durations_modified["total_amount"].values
    var = np.sort(var, axis = None)
    print("{} percentile value is {}".format(99+i,var[int(len(var)*(float(99+i)/100))]))
print("100 percentile value is ",var[-1])
```


```
99.0 percentile value is 66.13
99.1 percentile value is 68.13
99.2 percentile value is 69.6
99.3 percentile value is 69.6
99.4 percentile value is 69.73
99.5 percentile value is 69.75
99.6 percentile value is 69.76
99.7 percentile value is 72.58
99.8 percentile value is 75.35
99.9 percentile value is 88.28
100 percentile value is 3950611.6
```

**Observation:-** As even the 99.9th percentile value doesn't look like an outlier, as there is not much difference between the 99.8th percentile and 99.9th percentile, we move on to do graphical analysis

In [28]:


```
#below plot shows us the fare values(sorted) to find a sharp increase to remove those values as outliers
# plot the fare amount excluding last two values in sorted data
plt.plot(var[:-2])
plt.show()
```






In [48]:


```
# a very sharp increase in fare values can be seen  
# plotting last three total fare values, and we can observe there is share increase in the values  
plt.plot(var[-3:])  
plt.show()
```



In [49]:

```
#now looking at values not including the last two points we again find a drastic increase at around 1000 fare value  
# we plot last 50 values excluding last two values  
plt.plot(var[-50:-2])  
plt.show()
```





## Remove all outliers/errorous points.

In [8]:

```
#removing all outliers based on our univariate analysis above
def remove_outliers(new_frame):

    a = new_frame.shape[0]
    print ("Number of pickup records = ",a)
    temp_frame = new_frame[((new_frame.dropoff_longitude >= -74.15) & (new_frame.dropoff_longitude <= -73.7004) & \
                           (new_frame.dropoff_latitude >= 40.5774) & (new_frame.dropoff_latitude <= 40.9176)) & \
                           ((new_frame.pickup_longitude >= -74.15) & (new_frame.pickup_latitude >= 40.5774) & \
                           (new_frame.pickup_longitude <= -73.7004) & (new_frame.pickup_latitude <= 40.9176))]
    b = temp_frame.shape[0]
    print ("Number of outlier coordinates lying outside NY boundaries:",(a-b))

    temp_frame = new_frame[(new_frame.trip_times > 0) & (new_frame.trip_times < 720)]
    c = temp_frame.shape[0]
    print ("Number of outliers from trip times analysis:",(a-c))

    temp_frame = new_frame[(new_frame.trip_distance > 0) & (new_frame.trip_distance < 23)]
    d = temp_frame.shape[0]
    print ("Number of outliers from trip distance analysis:",(a-d))

    temp_frame = new_frame[(new_frame.Speed <= 65) & (new_frame.Speed >= 0)]
    e = temp_frame.shape[0]
    print ("Number of outliers from speed analysis:",(a-e))

    temp_frame = new_frame[(new_frame.total_amount <1000) & (new_frame.total_amount >0)]
    f = temp_frame.shape[0]
    print ("Number of outliers from fare analysis:",(a-f))

    new_frame = new_frame[((new_frame.dropoff_longitude >= -74.15) & (new_frame.dropoff_longitude <= -73.7004) & \
                           (new_frame.dropoff_latitude >= 40.5774) & (new_frame.dropoff_latitude <= 40.9176)) & \
                           ((new_frame.pickup_longitude >= -74.15) & (new_frame.pickup_latitude >= 40.5774) & \
                           (new_frame.pickup_longitude <= -73.7004) & (new_frame.pickup_latitude <= 40.9176))]

    new_frame = new_frame[(new_frame.trip_times > 0) & (new_frame.trip_times < 720)]
    new_frame = new_frame[(new_frame.trip_distance > 0) & (new_frame.trip_distance < 23)]
    new_frame = new_frame[(new_frame.Speed < 45.31) & (new_frame.Speed > 0)]
    new_frame = new_frame[(new_frame.total_amount <1000) & (new_frame.total_amount >0)]

    print ("Total outliers removed",a - new_frame.shape[0])
    print ("---")
    return new_frame
```

In [9]:

```
print ("Removing outliers in the month of Jan-2015")
print ("---")
frame_with_durations_outliers_removed = remove_outliers(frame_with_durations)
print("fraction of data points that remain after removing outliers",
      float(len(frame_with_durations_outliers_removed))/len(frame_with_durations))
```

```

Removing outliers in the month of Jan-2015
----
Number of pickup records = 12748986
Number of outlier coordinates lying outside NY boundaries: 293919
Number of outliers from trip times analysis: 23889
Number of outliers from trip distance analysis: 92597
Number of outliers from speed analysis: 24473
Number of outliers from fare analysis: 5275
Total outliers removed 377910
---
fraction of data points that remain after removing outliers 0.9703576425607495

```

## Data-preperation

### Clustering/Segmentation

In [10]:

```

#trying different cluster sizes to choose the right K in K-means
coords = frame_with_durations_outliers_removed[['pickup_latitude', 'pickup_longitude']].values
neighbours=[]

def find_min_distance(cluster_centers, cluster_len):
    nice_points = 0
    wrong_points = 0
    less2 = []
    more2 = []
    min_dist=1000
    for i in range(0, cluster_len):
        nice_points = 0
        wrong_points = 0
        for j in range(0, cluster_len):
            if j!=i:
                distance = gpxpy.geo.haversine_distance(cluster_centers[i][0], cluster_centers[i][1],
                cluster_centers[j][0], cluster_centers[j][1])
                min_dist = min(min_dist,distance/(1.60934*1000))
                if (distance/(1.60934*1000)) <= 2:
                    nice_points +=1
                else:
                    wrong_points += 1
            less2.append(nice_points)
            more2.append(wrong_points)
    neighbours.append(less2)
    print ("On choosing a cluster size of ",cluster_len,"\\nAvg. Number of Clusters within the vicinity (i.e. intercluster-distance < 2):", np.ceil(sum(less2)/len(less2)), "\\nAvg. Number of Clusters outside the vicinity (i.e. intercluster-distance > 2):", np.ceil(sum(more2)/len(more2)), "\\nMin inter-cluster distance = ",min_dist,"\\n---")

def find_clusters(increment):
    kmeans = MiniBatchKMeans(n_clusters=increment, batch_size=10000, random_state=42).fit(coords)
    frame_with_durations_outliers_removed['pickup_cluster'] =
    kmeans.predict(frame_with_durations_outliers_removed[['pickup_latitude', 'pickup_longitude']])
    cluster_centers = kmeans.cluster_centers_
    cluster_len = len(cluster_centers)
    return cluster_centers, cluster_len

# we need to choose number of clusters so that, there are more number of cluster regions
#that are close to any cluster center
# and make sure that the minimum inter cluster should not be very less
for increment in range(10, 100, 10):
    cluster_centers, cluster_len = find_clusters(increment)
    find_min_distance(cluster_centers, cluster_len)

```

```

On choosing a cluster size of 10
Avg. Number of Clusters within the vicinity (i.e. intercluster-distance < 2): 2.0
Avg. Number of Clusters outside the vicinity (i.e. intercluster-distance > 2): 8.0
Min inter-cluster distance = 1.0945442325142543
---
On choosing a cluster size of 20
Avg. Number of Clusters within the vicinity (i.e. intercluster-distance < 2): 4.0
Avg. Number of Clusters outside the vicinity (i.e. intercluster-distance > 2): 16.0
Min inter-cluster distance = 0.7131298007387813

```

```

---
On choosing a cluster size of 30
Avg. Number of Clusters within the vicinity (i.e. intercluster-distance < 2): 8.0
Avg. Number of Clusters outside the vicinity (i.e. intercluster-distance > 2): 22.0
Min inter-cluster distance = 0.5185088176172206
---
On choosing a cluster size of 40
Avg. Number of Clusters within the vicinity (i.e. intercluster-distance < 2): 8.0
Avg. Number of Clusters outside the vicinity (i.e. intercluster-distance > 2): 32.0
Min inter-cluster distance = 0.5069768450363973
---
On choosing a cluster size of 50
Avg. Number of Clusters within the vicinity (i.e. intercluster-distance < 2): 12.0
Avg. Number of Clusters outside the vicinity (i.e. intercluster-distance > 2): 38.0
Min inter-cluster distance = 0.365363025983595
---
On choosing a cluster size of 60
Avg. Number of Clusters within the vicinity (i.e. intercluster-distance < 2): 14.0
Avg. Number of Clusters outside the vicinity (i.e. intercluster-distance > 2): 46.0
Min inter-cluster distance = 0.34704283494187155
---
On choosing a cluster size of 70
Avg. Number of Clusters within the vicinity (i.e. intercluster-distance < 2): 16.0
Avg. Number of Clusters outside the vicinity (i.e. intercluster-distance > 2): 54.0
Min inter-cluster distance = 0.30502203163244707
---
On choosing a cluster size of 80
Avg. Number of Clusters within the vicinity (i.e. intercluster-distance < 2): 18.0
Avg. Number of Clusters outside the vicinity (i.e. intercluster-distance > 2): 62.0
Min inter-cluster distance = 0.29220324531738534
---
On choosing a cluster size of 90
Avg. Number of Clusters within the vicinity (i.e. intercluster-distance < 2): 21.0
Avg. Number of Clusters outside the vicinity (i.e. intercluster-distance > 2): 69.0
Min inter-cluster distance = 0.18257992857034985
---

```

In [16]:

```

print(cluster_centers[40][0])
print(cluster_centers[40][1])

```

```

40.76295756776805
-73.95990035721279

```

In [35]:

```

print(neighbours)

```

```

[[0, 4, 0, 2, 1, 3, 0, 3, 2, 3], [6, 0, 5, 2, 0, 7, 5, 0, 6, 3, 6, 4, 5, 0, 6, 0, 7, 0, 0, 2], [5,
12, 10, 0, 0, 12, 11, 11, 12, 5, 0, 1, 12, 1, 7, 9, 11, 1, 13, 2, 7, 3, 7, 12, 10, 12, 13, 2, 0, 1
3], [13, 0, 10, 0, 14, 5, 11, 14, 13, 1, 15, 14, 11, 9, 3, 13, 11, 1, 15, 0, 13, 11, 8, 0, 14, 12,
1, 4, 1, 0, 14, 0, 3, 0, 15, 11, 11, 7, 7, 13], [8, 17, 17, 0, 2, 22, 16, 19, 21, 8, 1, 2, 20, 2,
15, 14, 21, 2, 18, 5, 11, 5, 14, 17, 15, 21, 20, 4, 0, 20, 19, 5, 18, 0, 20, 16, 1, 14, 3, 22, 2,
7, 0, 3, 20, 11, 10, 17, 21, 16], [8, 24, 1, 17, 2, 24, 15, 14, 25, 6, 22, 14, 20, 3, 0, 2, 2, 20,
21, 24, 22, 11, 19, 24, 16, 8, 20, 0, 25, 22, 25, 13, 17, 8, 1, 2, 17, 21, 25, 23, 5, 22, 0, 24,
1, 0, 5, 2, 18, 8, 22, 22, 18, 2, 23, 0, 3, 17, 24, 1], [20, 27, 1, 16, 1, 28, 13, 19, 24, 23, 2,
28, 6, 1, 27, 2, 22, 25, 2, 23, 21, 24, 27, 9, 3, 16, 27, 7, 18, 0, 1, 20, 0, 29, 24, 0, 29, 0, 19
, 21, 6, 28, 3, 10, 2, 20, 1, 28, 1, 26, 24, 12, 28, 9, 7, 0, 23, 24, 27, 28, 21, 19, 27, 27, 16,
10, 0, 13, 27, 18], [27, 1, 32, 25, 2, 20, 33, 15, 6, 31, 27, 3, 3, 12, 7, 26, 22, 25, 29, 15, 18,
30, 29, 26, 28, 26, 9, 30, 33, 23, 9, 31, 2, 29, 0, 0, 28, 20, 2, 20, 22, 14, 32, 0, 27, 25, 22,
31, 30, 28, 3, 2, 0, 28, 21, 1, 28, 0, 29, 8, 0, 0, 32, 25, 14, 15, 6, 29, 4, 3, 30, 28, 5, 34, 3,
12, 22, 5, 1, 15], [27, 37, 1, 1, 36, 18, 26, 33, 28, 31, 6, 7, 4, 38, 7, 18, 25, 38, 2, 31, 13, 1
, 1, 29, 36, 34, 35, 11, 32, 0, 36, 14, 26, 0, 37, 2, 32, 38, 31, 26, 1, 32, 9, 37, 0, 22, 16, 0,
39, 30, 2, 3, 30, 1, 32, 31, 0, 26, 38, 0, 0, 36, 24, 34, 22, 13, 24, 1, 6, 12, 36, 32, 16, 32, 37
, 39, 32, 21, 29, 5, 0, 36, 21, 27, 4, 0, 35, 37, 0, 4, 1]]

```

## Inference:

- The main objective was to find a optimal min. distance(Which roughly estimates to the radius of a cluster) between the clusters which we got was 30

## For Cluster size = 30

In [11]:

```
# if check for the 50 clusters you can observe that there are two clusters with only 0.3 miles apart from each other
# so we choose 30 clusters for solve the further problem

# Getting 30 clusters using the kmeans
kmeans = MiniBatchKMeans(n_clusters=30, batch_size=10000, random_state=0).fit(coords)
frame_with_durations_outliers_removed['pickup_cluster'] =
kmeans.predict(frame_with_durations_outliers_removed[['pickup_latitude', 'pickup_longitude']])
```

## Plotting the cluster centers:

In [12]:

```
# Plotting the cluster centers on OSM
cluster_centers = kmeans.cluster_centers_
cluster_len = len(cluster_centers)
map_osm = folium.Map(location=[40.734695, -73.990372], tiles='Stamen Toner')
for i in range(cluster_len):
    folium.Marker(list((cluster_centers[i][0],cluster_centers[i][1])), popup=(str(cluster_centers[i][0])+str(cluster_centers[i][1]))).add_to(map_osm)
map_osm
```

Out[12]:

## Plotting the clusters:

In [19]:


```
#Visualising the clusters on a map
def plot_clusters(frame):
    city_long_border = (-74.03, -73.75)
    city_lat_border = (40.63, 40.85)
    fig, ax = plt.subplots(ncols=1, nrows=1)
    ax.scatter(frame.pickup_longitude.values[:100000], frame.pickup_latitude.values[:100000], s=10,
lw=0,
               c=frame.pickup_cluster.values[:100000], cmap='tab20', alpha=0.2)
```

```

    ax.set_xlim(city_long_border)
    ax.set_ylim(city_lat_border)
    ax.set_xlabel('Longitude')
    ax.set_ylabel('Latitude')
    plt.show()

plot_clusters(frame_with_durations_outliers_removed)

```



## Time-binning

In [13]:

```

#Refer:https://www.unixtimestamp.com/
# 1420070400 : 2015-01-01 00:00:00
# 1422748800 : 2015-02-01 00:00:00
# 1425168000 : 2015-03-01 00:00:00
# 1427846400 : 2015-04-01 00:00:00
# 1430438400 : 2015-05-01 00:00:00
# 1433116800 : 2015-06-01 00:00:00

# 1451606400 : 2016-01-01 00:00:00
# 1454284800 : 2016-02-01 00:00:00
# 1456790400 : 2016-03-01 00:00:00
# 1459468800 : 2016-04-01 00:00:00
# 1462060800 : 2016-05-01 00:00:00
# 1464739200 : 2016-06-01 00:00:00

def add_pickup_bins(frame,month,year):
    unix_pickup_times=[i for i in frame['pickup_times'].values]
    unix_times = [[1420070400,1422748800,1425168000,1427846400,1430438400,1433116800],\
                  [1451606400,1454284800,1456790400,1459468800,1462060800,1464739200]]

    start_pickup_unix=unix_times[year-2015][month-1]
    # https://www.timeanddate.com/time/zones/est
    # (int((i-start_pickup_unix)/600)+33) : our unix time is in gmt so we are converting it to est
    tenminutewise_binned_unix_pickup_times=[(int((i-start_pickup_unix)/600)+33) for i in unix_pickup_times]
    frame['pickup_bins'] = np.array(tenminutewise_binned_unix_pickup_times)
    return frame

```

In [14]:

```
# clustering, making pickup bins and grouping by pickup cluster and pickup bins
```

```

frame_with_durations_outliers_removed['pickup_cluster'] =
kmeans.predict(frame_with_durations_outliers_removed[['pickup_latitude', 'pickup_longitude']])
jan_2015_frame = add_pickup_bins(frame_with_durations_outliers_removed,1,2015)
jan_2015_groupby =
jan_2015_frame[['pickup_cluster','pickup_bins','trip_distance']].groupby(['pickup_cluster','pickup_bins']).count()

```

In [15]:

```

# we add two more columns 'pickup_cluster' (to which cluster it belongs to)
# and 'pickup_bins' (to which 10min intravel the trip belongs to)
jan_2015_frame.head()

```

Out[15]:

	passenger_count	trip_distance	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	total_amount	trip_times	pickup_cluster
0	1	1.59	-73.993896	40.750111	-73.974785	40.750618	17.05	18.050000	1.42
1	1	3.30	-74.001648	40.724243	-73.994415	40.759109	17.80	19.833333	1.42
2	1	1.80	-73.963341	40.802788	-73.951820	40.824413	10.80	10.050000	1.42
3	1	0.50	-74.009087	40.713818	-74.004326	40.719986	4.80	1.866667	1.42
4	1	3.00	-73.971176	40.762428	-74.004181	40.742653	16.30	19.316667	1.42

In [23]:

```

# hear the trip_distance represents the number of pickups that are happened in that particular 10min intravel
# this data frame has two indices
# primary index: pickup_cluster (cluster number)
# secondary index : pickup_bins (we devide whole months time into 10min intravels 24*31*60/10 =4464 bins)
jan_2015_groupby.head()

```

Out[23]:

trip_distance		
pickup_cluster	pickup_bins	
0	1	105
	2	199
	3	208
	4	141
	5	155

In [16]:

```

# upto now we cleaned data and prepared data for the month 2015,
# now do the same operations for months Jan, Feb, March of 2016
# 1. get the dataframe which includes only required columns
# 2. adding trip times, speed, unix time stamp of pickup_time
# 4. remove the outliers based on trip_times, speed, trip_duration, total_amount
# 5. add pickup_cluster to each data point
# 6. add pickup_bin (index of 10min intravel to which that trip belongs to)
# 7. group by data, based on 'pickup_cluster' and 'pickup_bin'

# Data Preparation for the months of Jan, Feb and March 2016
def datapreparation(month,kmeans,month_no,year_no):

    print ("Return with trip times..")

    frame_with_durations = return_with_trip_times(month)

    print ("Remove outliers..")
    frame_with_durations_outliers_removed = remove_outliers(frame_with_durations)

    print ("Estimating clusters..")
    frame_with_durations_outliers_removed['pickup_cluster'] =

```

```

frame_with_durations_outliers_removed['pickup_cluster'],
kmeans.predict(frame_with_durations_outliers_removed[['pickup_latitude', 'pickup_longitude']])
#frame_with_durations_outliers_removed_2016['pickup_cluster'] =
kmeans.predict(frame_with_durations_outliers_removed_2016[['pickup_latitude',
'pickup_longitude']])

print ("Final groupbying..")
final_updated_frame = add_pickup_bins(frame_with_durations_outliers_removed,month_no,year_no)
final_groupby_frame = final_updated_frame[['pickup_cluster','pickup_bins','trip_distance']].groupby(['pickup_cluster','pickup_bins']).count()

return final_updated_frame,final_groupby_frame

month_jan_2016 = dd.read_csv('yellow_tripdata_2016-01.csv')
month_feb_2016 = dd.read_csv('yellow_tripdata_2016-02.csv')
month_mar_2016 = dd.read_csv('yellow_tripdata_2016-03.csv')

jan_2016_frame,jan_2016_groupby = datapreparation(month_jan_2016,kmeans,1,2016)
feb_2016_frame,feb_2016_groupby = datapreparation(month_feb_2016,kmeans,2,2016)
mar_2016_frame,mar_2016_groupby = datapreparation(month_mar_2016,kmeans,3,2016)

```

Return with trip times..

Remove outliers..

```

Number of pickup records = 10906858
Number of outlier coordinates lying outside NY boundaries: 214677
Number of outliers from trip times analysis: 27190
Number of outliers from trip distance analysis: 79742
Number of outliers from speed analysis: 21047
Number of outliers from fare analysis: 4991
Total outliers removed 297784
---
```

Estimating clusters..

Final groupbying..

Return with trip times..

Remove outliers..

```

Number of pickup records = 11382049
Number of outlier coordinates lying outside NY boundaries: 223161
Number of outliers from trip times analysis: 27670
Number of outliers from trip distance analysis: 81902
Number of outliers from speed analysis: 22437
Number of outliers from fare analysis: 5476
Total outliers removed 308177
---
```

Estimating clusters..

Final groupbying..

Return with trip times..

Remove outliers..

```

Number of pickup records = 12210952
Number of outlier coordinates lying outside NY boundaries: 232444
Number of outliers from trip times analysis: 30868
Number of outliers from trip distance analysis: 87318
Number of outliers from speed analysis: 23889
Number of outliers from fare analysis: 5859
Total outliers removed 324635
---
```

Estimating clusters..

Final groupbying..

## Smoothing

In [18]:

```

# Gets the unique bins where pickup values are present for each each reigion

# for each cluster region we will collect all the indices of 10min intravels in which the pickups
# are happened
# we got an observation that there are some pickpbins that doesnt have any pickups
def return_unq_pickup_bins(frame):
    values = []
    for i in range(0,30):
        new = frame[frame['pickup_cluster'] == i]
        list_unq = list(set(new['pickup_bins']))
        list_unq.sort()
        values.append(list_unq)

```

```
return values
```

In [19]:

```
# for every month we get all indices of 10min intravels in which atleast one pickup got happened

#jan
jan_2015_unique = return_unq_pickup_bins(jan_2015_frame)
jan_2016_unique = return_unq_pickup_bins(jan_2016_frame)

#feb
feb_2016_unique = return_unq_pickup_bins(feb_2016_frame)

#march
mar_2016_unique = return_unq_pickup_bins(mar_2016_frame)
```

In [20]:

```
# for each cluster number of 10min intavels with 0 pickups
for i in range(30):
    print("for the ",i,"th cluster number of 10min intavels with zero pickups: ",4464 -
len(set(jan_2015_unique[i])))
    print('-'*60)
```

```
for the 0 th cluster number of 10min intavels with zero pickups: 26
-----
for the 1 th cluster number of 10min intavels with zero pickups: 30
-----
for the 2 th cluster number of 10min intavels with zero pickups: 150
-----
for the 3 th cluster number of 10min intavels with zero pickups: 35
-----
for the 4 th cluster number of 10min intavels with zero pickups: 170
-----
for the 5 th cluster number of 10min intavels with zero pickups: 40
-----
for the 6 th cluster number of 10min intavels with zero pickups: 320
-----
for the 7 th cluster number of 10min intavels with zero pickups: 35
-----
for the 8 th cluster number of 10min intavels with zero pickups: 39
-----
for the 9 th cluster number of 10min intavels with zero pickups: 46
-----
for the 10 th cluster number of 10min intavels with zero pickups: 98
-----
for the 11 th cluster number of 10min intavels with zero pickups: 32
-----
for the 12 th cluster number of 10min intavels with zero pickups: 37
-----
for the 13 th cluster number of 10min intavels with zero pickups: 326
-----
for the 14 th cluster number of 10min intavels with zero pickups: 35
-----
for the 15 th cluster number of 10min intavels with zero pickups: 29
-----
for the 16 th cluster number of 10min intavels with zero pickups: 25
-----
for the 17 th cluster number of 10min intavels with zero pickups: 40
-----
for the 18 th cluster number of 10min intavels with zero pickups: 30
-----
for the 19 th cluster number of 10min intavels with zero pickups: 35
-----
for the 20 th cluster number of 10min intavels with zero pickups: 40
-----
for the 21 th cluster number of 10min intavels with zero pickups: 38
-----
for the 22 th cluster number of 10min intavels with zero pickups: 34
-----
for the 23 th cluster number of 10min intavels with zero pickups: 49
-----
for the 24 th cluster number of 10min intavels with zero pickups: 49
-----
for the 25 th cluster number of 10min intavels with zero pickups: 27
```

```

for the 25 th cluster number of 10min intavels with zero pickups: 25
-----
for the 26 th cluster number of 10min intavels with zero pickups: 26
-----
for the 27 th cluster number of 10min intavels with zero pickups: 720
-----
for the 28 th cluster number of 10min intavels with zero pickups: 34
-----
for the 29 th cluster number of 10min intavels with zero pickups: 29
-----
```

there are two ways to fill up these values

- Fill the missing value with 0's
- Fill the missing values with the avg values
  - Case 1:(values missing at the start)
   
Ex1:  $\lfloor \frac{x}{4} \rfloor, \lceil \frac{x}{4} \rceil, \lceil \frac{x}{4} \rceil, \lceil \frac{x}{4} \rceil$ 
  
Ex2:  $\lfloor \frac{x}{3} \rfloor, \lceil \frac{x}{3} \rceil, \lceil \frac{x}{3} \rceil$
  - Case 2:(values missing in middle)
   
Ex1:  $\lfloor \frac{x+y}{4} \rfloor, \lceil \frac{x+y}{4} \rceil, \lceil \frac{x+y}{4} \rceil, \lceil \frac{x+y}{4} \rceil$ 
  
Ex2:  $\lfloor \frac{x+y}{5} \rfloor, \lceil \frac{x+y}{5} \rceil, \lceil \frac{x+y}{5} \rceil, \lceil \frac{x+y}{5} \rceil, \lceil \frac{x+y}{5} \rceil$
  - Case 3:(values missing at the end)
   
Ex1:  $\lfloor \frac{x}{4} \rfloor, \lceil \frac{x}{4} \rceil, \lceil \frac{x}{4} \rceil, \lceil \frac{x}{4} \rceil$ 
  
Ex2:  $\lfloor \frac{x}{2} \rfloor, \lceil \frac{x}{2} \rceil$

In [21]:

```

# Fills a value of zero for every bin where no pickup data is present
# the count_values: number pickps that are happened in each region for each 10min intravel
# there wont be any value if there are no pickups.
# values: number of unique bins

# for every 10min intravel(pickup_bin) we will check it is there in our unique bin,
# if it is there we will add the count_values[index] to smoothed data
# if not we add 0 to the smoothed data
# we finally return smoothed data
def fill_missing(count_values,values):
    smoothed_regions=[]
    ind=0
    for r in range(0,30):
        smoothed_bins=[]
        for i in range(4464):
            if i in values[r]:
                smoothed_bins.append(count_values[ind])
                ind+=1
            else:
                smoothed_bins.append(0)
        smoothed_regions.extend(smoothed_bins)
    return smoothed_regions
```

In [22]:

```

# Fills a value of zero for every bin where no pickup data is present
# the count_values: number pickps that are happened in each region for each 10min intravel
# there wont be any value if there are no pickups.
# values: number of unique bins

# for every 10min intravel(pickup_bin) we will check it is there in our unique bin,
# if it is there we will add the count_values[index] to smoothed data
# if not we add smoothed data (which is calculated based on the methods that are discussed in the
# above markdown cell)
# we finally return smoothed data
def smoothing(count_values,values):
    smoothed_regions=[] # stores list of final smoothed values of each reigion
    ind=0
    repeat=0
    smoothed_value=0
    for r in range(0,30):
        smoothed_bins=[] #stores the final smoothed values
        repeat=0
        for i in range(4464):
            if repeat!=0: # prevents iteration for a value which is already visited/resolved
                repeat-=1
            else:
```

```

repeat =+
continue
if i in values[r]: #checks if the pickup-bin exists
    smoothed_bins.append(count_values[ind]) # appends the value of the pickup bin if it
exists
else:
    if i!=0:
        right_hand_limit=0
        for j in range(i,4464):
            if j not in values[r]: #searches for the left-limit or the pickup-bin
value which has a pickup value
                continue
            else:
                right_hand_limit=j
                break
        if right_hand_limit==0:
            #Case 1: When we have the last/last few values are found to be missing,hence we
have no right-limit here
            smoothed_value=count_values[ind-1]*1.0/((4463-i)+2)*1.0

            for j in range(i,4464):
                smoothed_bins.append(math.ceil(smoothed_value))
                smoothed_bins[i-1] = math.ceil(smoothed_value)
                repeat=(4463-i)
                ind-=1
            else:
                #Case 2: When we have the missing values between two known values
                smoothed_value=(count_values[ind-1]+count_values[ind])*1.0/((right_hand_lim
t-i)+2)*1.0
                for j in range(i,right_hand_limit+1):
                    smoothed_bins.append(math.ceil(smoothed_value))
                    smoothed_bins[i-1] = math.ceil(smoothed_value)
                    repeat=(right_hand_limit-i)
            else:
                #Case 3: When we have the first/first few values are found to be missing,hence we
have no left-limit here
                right_hand_limit=0
                for j in range(i,4464):
                    if j not in values[r]:
                        continue
                    else:
                        right_hand_limit=j
                        break
                smoothed_value=count_values[ind]*1.0/((right_hand_limit-i)+1)*1.0
                for j in range(i,right_hand_limit+1):
                    smoothed_bins.append(math.ceil(smoothed_value))
                    repeat=(right_hand_limit-i)
                ind+=1
            smoothed_regions.extend(smoothed_bins)
return smoothed_regions

```

In [23]:

```

#Filling Missing values of Jan-2015 with 0
# here in jan_2015_groupby dataframe the trip_distance represents the number of pickups that are h
appened
jan_2015_fill = fill_missing(jan_2015_groupby['trip_distance'].values,jan_2015_unique)

#Smoothing Missing values of Jan-2015
jan_2015_smooth = smoothing(jan_2015_groupby['trip_distance'].values,jan_2015_unique)

```

In [24]:

```


# number of 10min indices for jan 2015= 24*31*60/10 = 4464
# number of 10min indices for jan 2016 = 24*31*60/10 = 4464
# number of 10min indices for feb 2016 = 24*29*60/10 = 4176
# number of 10min indices for march 2016 = 24*30*60/10 = 4320
# for each cluster we will have 4464 values, therefore 40*4464 = 178560 (length of the
jan_2015_fill)
print("number of 10min intravels among all the clusters ",len(jan_2015_fill))

```

number of 10min intravels among all the clusters 133920

```
In [0]:
```

```
# Smoothing vs Filling
# sample plot that shows two variations of filling missing values
# we have taken the number of pickups for cluster region 2
plt.figure(figsize=(10,5))
plt.plot(jan_2015_fill[4464:8920], label="zero filled values")
plt.plot(jan_2015_smooth[4464:8920], label="filled with avg values")
plt.legend()
plt.show()
```



```
In [0]:
```

```
# why we choose, these methods and which method is used for which data?

# Ans: consider we have data of some month in 2015 jan 1st, 10 _ _ _ 20, i.e there are 10 pickups
# that are happened in 1st
# 10min intravel, 0 pickups happened in 2nd 10mins intravel, 0 pickups happened in 3rd 10min intravel
# and 20 pickups happened in 4th 10min intravel.
# in fill_missing method we replace these values like 10, 0, 0, 20
# where as in smoothing method we replace these values as 6,6,6,6,6, if you can check the number of pickups
# that are happened in the first 40min are same in both cases, but if you can observe that we looking at the future values
# wheen you are using smoothing we are looking at the future number of pickups which might cause a data leakage.

# so we use smoothing for jan 2015th data since it acts as our training data
# and we use simple fill_misssing method for 2016th data.
```

```
In [25]:
```

```
# Jan-2015 data is smoothed, Jan, Feb & March 2016 data missing values are filled with zero
jan_2015_smooth = smoothing(jan_2015.groupby(['trip_distance']).values,jan_2015.unique)
jan_2016_smooth = fill_missing(jan_2016.groupby(['trip_distance']).values,jan_2016.unique)
feb_2016_smooth = fill_missing(feb_2016.groupby(['trip_distance']).values,feb_2016.unique)
mar_2016_smooth = fill_missing(mar_2016.groupby(['trip_distance']).values,mar_2016.unique)

# Making list of all the values of pickup data in every bin for a period of 3 months and storing them region-wise
regions_cum = []

# a =[1,2,3]
# b = [2,3,4]
# a+b = [1, 2, 3, 2, 3, 4]

# number of 10min indices for jan 2015= 24*31*60/10 = 4464
# number of 10min indices for jan 2016 = 24*31*60/10 = 4464
# number of 10min indices for feb 2016 = 24*30*60/10 = 4320
# number of 10min indices for mar 2016 = 24*29*60/10 = 4176
```

```

# number of 1min indices for rep 2016 = 24*29*60/10 = 416
# number of 10min indices for march 2016 = 24*31*60/10 = 4464
# regions_cum: it will contain 40 lists, each list will contain 4464+4176+4464 values which represents the number of pickups
# that are happened for three months in 2016 data

for i in range(0,30):
    regions_cum.append(jan_2016_smooth[4464*i:4464*(i+1)]+feb_2016_smooth[4176*i:4176*(i+1)]+mar_2016_smooth[4464*i:4464*(i+1)])

# print(len(regions_cum))
# 40
# print(len(regions_cum[0]))
# 13104

```



## Time series and Fourier Transforms

In [26]:





```





def uniqueish_color():
    """There're better ways to generate unique colors, but this isn't awful."""
    return plt.cm.gist_ncar(np.random.random())
first_x = list(range(0,4464))
second_x = list(range(4464,8640))
third_x = list(range(8640,13104))
for i in range(30):
    plt.figure(figsize=(10,4))
    plt.plot(first_x,regions_cum[i][:4464], color=uniqueish_color(), label='2016 Jan month data')
    plt.plot(second_x,regions_cum[i][4464:8640], color=uniqueish_color(), label='2016 feb month data')
    plt.plot(third_x,regions_cum[i][8640:], color=uniqueish_color(), label='2016 march month data')
    plt.legend()
    plt.show()





```











0 2000 4000 6000 8000 10000 12000






























In [27]:

```
# getting peaks: https://blog.ytotech.com/2015/11/01/findpeaks-in-python/
# read more about fft function :
https://docs.scipy.org/doc/numpy/reference/generated/numpy.fft.fft.html
Y      = np.fft.fft(np.array(jan_2016_smooth)[0:4460])
# read more about the fftfreq:
https://docs.scipy.org/doc/numexpr/reference/generated/numexpr.fft.freq.html
```

```

from os import path
import numpy as np
import matplotlib.pyplot as plt
freq = np.fft.fftfreq(4460, 1)
n = len(freq)
plt.figure()
plt.plot(freq[:int(n/2)], np.abs(Y)[:int(n/2)])
plt.xlabel("Frequency")
plt.ylabel("Amplitude")
plt.show()

```



In [28]:

```

#Preparing the Dataframe only with x(i) values as jan-2015 data and y(i) values as jan-2016
ratios_jan = pd.DataFrame()
ratios_jan['Given']=jan_2015_smooth
ratios_jan['Prediction']=jan_2016_smooth
ratios_jan['Ratios']=ratios_jan['Prediction']*1.0/ratios_jan['Given']*1.0

```

## Modelling: Baseline Models

Now we get into modelling in order to forecast the pickup densities for the months of Jan, Feb and March of 2016 for which we are using multiple models with two variations

1. Using Ratios of the 2016 data to the 2015 data i.e  $R_t = P_t^{2016}/P_t^{2015}$
2. Using Previous known values of the 2016 data itself to predict the future values

### Simple Moving Averages

The First Model used is the Moving Averages Model which uses the previous n values in order to predict the next value

Using Ratio Values -  $R_t = (R_{t-1} + R_{t-2} + R_{t-3} \dots R_{t-n})/n$

In [29]:

```

def MA_R_Predictions(ratios,month):
    predicted_ratio=(ratios['Ratios'].values)[0]
    error=[]
    predicted_values=[]
    window_size=3
    predicted_ratio_values=[]
    for i in range(0,4464*30):
        predicted_ratio_values.append(predicted_ratio)

```

```

if i%4464==0:
    predicted_ratio_values.append(0)
    predicted_values.append(0)
    error.append(0)
    continue
predicted_ratio_values.append(predicted_ratio)
predicted_values.append(int(((ratios['Given'].values)[i])*predicted_ratio))
error.append(abs((math.pow(int((ratios['Given'].values)[i])*predicted_ratio)-(ratios['Prediction'].values)[i],1))))
if i+1>=window_size:
    predicted_ratio=sum((ratios['Ratios'].values)[(i+1)-window_size:(i+1)]) / window_size
else:
    predicted_ratio=sum((ratios['Ratios'].values)[0:(i+1)]) / (i+1)

ratios['MA_R_Predicted'] = predicted_values
ratios['MA_R_Error'] = error
mape_err = (sum(error)/len(error)) / (sum(ratios['Prediction'].values)/len(ratios['Prediction'].values))
mse_err = sum([e**2 for e in error]) / len(error)
return ratios, mape_err, mse_err

```

For the above the Hyperparameter is the window-size (n) which is tuned manually and it is found that the window-size of 3 is optimal for getting the best results using Moving Averages using previous Ratio values therefore we get  $R_t = (R_{t-1} + R_{t-2} + R_{t-3})/3$

Next we use the Moving averages of the 2016 values itself to predict the future value using  $P_t = (P_{t-1} + P_{t-2} + P_{t-3} \dots P_{t-n})/n$

In [30]:

```

def MA_P_Predictions(ratios,month):
    predicted_value=(ratios['Prediction'].values)[0]
    error=[]
    predicted_values=[]
    window_size=1
    predicted_ratio_values=[]
    for i in range(0,4464*30):
        predicted_values.append(predicted_value)
        error.append(abs((predicted_value-(ratios['Prediction'].values)[i],1)))
        if i+1>=window_size:
            predicted_value=int(sum((ratios['Prediction'].values)[(i+1)-window_size:(i+1)]) / window_size)
        else:
            predicted_value=int(sum((ratios['Prediction'].values)[0:(i+1)]) / (i+1))

    ratios['MA_P_Predicted'] = predicted_values
    ratios['MA_P_Error'] = error
    mape_err = (sum(error)/len(error)) / (sum(ratios['Prediction'].values)/len(ratios['Prediction'].values))
    mse_err = sum([e**2 for e in error]) / len(error)
    return ratios, mape_err, mse_err

```

For the above the Hyperparameter is the window-size (n) which is tuned manually and it is found that the window-size of 1 is optimal for getting the best results using Moving Averages using previous 2016 values therefore we get  $P_t = P_{t-1}$

## Weighted Moving Averages

The Moving Average Model used gave equal importance to all the values in the window used, but we know intuitively that the future is more likely to be similar to the latest values and less similar to the older values. Weighted Averages converts this analogy into a mathematical relationship giving the highest weight while computing the averages to the latest previous value and decreasing weights to the subsequent older ones

Weighted Moving Averages using Ratio Values -  $R_t = (N * R_{t-1} + (N-1) * R_{t-2} + (N-2) * R_{t-3} \dots 1 * R_{t-n}) / (N * (N+1)/2)$

In [31]:

```

def WA_R_Predictions(ratios,month):
    predicted_ratio=(ratios['Ratios'].values)[0]
    alpha=0.5
    error=[]
    predicted_values=[]

```

```

window_size=5
predicted_ratio_values=[]
for i in range(0,4464*30):
    if i%4464==0:
        predicted_ratio_values.append(0)
        predicted_values.append(0)
        error.append(0)
        continue
    predicted_ratio_values.append(predicted_ratio)
    predicted_values.append(int(((ratios['Given'].values)[i])*predicted_ratio))
    error.append(abs((math.pow(int(((ratios['Given'].values)[i])*predicted_ratio)-(ratios['Prediction'].values)[i],1))))
    if i+1>=window_size:
        sum_values=0
        sum_of_coeff=0
        for j in range(window_size,0,-1):
            sum_values += j*(ratios['Ratios'].values)[i-window_size+j]
            sum_of_coeff+=j
        predicted_ratio=sum_values/sum_of_coeff
    else:
        sum_values=0
        sum_of_coeff=0
        for j in range(i+1,0,-1):
            sum_values += j*(ratios['Ratios'].values)[j-1]
            sum_of_coeff+=j
        predicted_ratio=sum_values/sum_of_coeff

ratios['WA_R_Predicted'] = predicted_values
ratios['WA_R_Error'] = error
mape_err = (sum(error)/len(error))/(sum(ratios['Prediction'].values)/len(ratios['Prediction'].values))
mse_err = sum([e**2 for e in error])/len(error)
return ratios,mape_err,mse_err

```

For the above the Hyperparameter is the window-size (n) which is tuned manually and it is found that the window-size of 5 is optimal for getting the best results using Weighted Moving Averages using previous Ratio values therefore we get

$$R_t = (5 * R_{t-1} + 4 * R_{t-2} + 3 * R_{t-3} + 2 * R_{t-4} + R_{t-5}) / 15$$

Weighted Moving Averages using Previous 2016 Values -  $P_t = (N * P_{t-1} + (N-1) * P_{t-2} + (N-2) * P_{t-3} + \dots + 1 * P_{t-n}) / (N * (N+1)/2)$

In [32]:

```

def WA_P_Predictions(ratios,month):
    predicted_value=(ratios['Prediction'].values)[0]
    error=[]
    predicted_values=[]
    window_size=2
    for i in range(0,4464*30):
        predicted_values.append(predicted_value)
        error.append(abs((math.pow(predicted_value-(ratios['Prediction'].values)[i],1))))
        if i+1>=window_size:
            sum_values=0
            sum_of_coeff=0
            for j in range(window_size,0,-1):
                sum_values += j*(ratios['Prediction'].values)[i-window_size+j]
                sum_of_coeff+=j
            predicted_value=int(sum_values/sum_of_coeff)

        else:
            sum_values=0
            sum_of_coeff=0
            for j in range(i+1,0,-1):
                sum_values += j*(ratios['Prediction'].values)[j-1]
                sum_of_coeff+=j
            predicted_value=int(sum_values/sum_of_coeff)

    ratios['WA_P_Predicted'] = predicted_values
    ratios['WA_P_Error'] = error
    mape_err = (sum(error)/len(error))/(sum(ratios['Prediction'].values)/len(ratios['Prediction'].values))
    mse_err = sum([e**2 for e in error])/len(error)
    return ratios,mape_err,mse_err

```

For the above the Hyperparameter is the window-size (n) which is tuned manually and it is found that the window-size of 2 is optimal for getting the best results using Weighted Moving Averages using previous 2016 values therefore we get  $P_t = (2 * P_{t-1} + P_{t-2})/3$

## Exponential Weighted Moving Averages

[https://en.wikipedia.org/wiki/Moving\\_average#Exponential\\_moving\\_average](https://en.wikipedia.org/wiki/Moving_average#Exponential_moving_average) Through weighted averaged we have satisfied the analogy of giving higher weights to the latest value and decreasing weights to the subsequent ones but we still do not know which is the correct weighting scheme as there are infinitely many possibilities in which we can assign weights in a non-increasing order and tune the the hyperparameter window-size. To simplify this process we use Exponential Moving Averages which is a more logical way towards assigning weights and at the same time also using an optimal window-size.

In exponential moving averages we use a single hyperparameter alpha ( $\alpha$ ) which is a value between 0 & 1 and based on the value of the hyperparameter alpha the weights and the window sizes are configured.

For eg. If  $\alpha = 0.9$  then the number of days on which the value of the current iteration is based is  $\sim 1/(1 - \alpha) = 10$  i.e. we consider values 10 days prior before we predict the value for the current iteration. Also the weights are assigned using  $2/(N+1) = 0.18$ , where N = number of prior values being considered, hence from this it is implied that the first or latest value is assigned a weight of 0.18 which keeps exponentially decreasing for the subsequent values.

$$R_t' = \alpha * R_{t-1}' + (1 - \alpha) * R_{t-1}'$$

In [33]:

```
def EA_R1_Predictions(ratios,month):
    predicted_ratio=(ratios['Ratios'].values)[0]
    alpha=0.6
    error=[]
    predicted_values=[]
    predicted_ratio_values=[]
    for i in range(0,4464*30):
        if i%4464==0:
            predicted_ratio_values.append(0)
            predicted_values.append(0)
            error.append(0)
            continue
        predicted_ratio_values.append(predicted_ratio)
        predicted_values.append(int(((ratios['Given'].values)[i])*predicted_ratio))
        error.append(abs((math.pow(int(((ratios['Given'].values)[i])*predicted_ratio))-(ratios['Prediction'].values)[i],1))))
        predicted_ratio = (alpha*predicted_ratio) + (1-alpha)*((ratios['Ratios'].values)[i])

    ratios['EA_R1_Predicted'] = predicted_values
    ratios['EA_R1_Error'] = error
    mape_err = (sum(error)/len(error))/(sum(ratios['Prediction'].values)/len(ratios['Prediction'].values))
    mse_err = sum([e**2 for e in error])/len(error)
    return ratios,mape_err,mse_err
```

$$P_t' = \alpha * P_{t-1}' + (1 - \alpha) * P_{t-1}'$$

In [34]:

```
def EA_P1_Predictions(ratios,month):
    predicted_value= (ratios['Prediction'].values)[0]
    alpha=0.3
    error=[]
    predicted_values=[]
    for i in range(0,4464*30):
        if i%4464==0:
            predicted_values.append(0)
            error.append(0)
            continue
        predicted_values.append(predicted_value)
        error.append(abs((math.pow(predicted_value-(ratios['Prediction'].values)[i],1))))
        predicted_value =int((alpha*predicted_value) + (1-alpha)*((ratios['Prediction'].values)[i])))

    ratios['EA_P1_Predicted'] = predicted_values
    ratios['EA_P1_Error'] = error
    mape_err = (sum(error)/len(error))/(sum(ratios['Prediction'].values)/len(ratios['Prediction'].values))
    mse_err = sum([e**2 for e in error])/len(error)
    return ratios,mape_err,mse_err
```

```

        ratios,
        mse_err = sum([e**2 for e in error])/len(error)
    return ratios,mape_err,mse_err

```

In [35]:

```

mean_err=[0]*10
median_err=[0]*10
ratios_jan,mean_err[0],median_err[0]=MA_R_Predictions(ratios_jan,'jan')
ratios_jan,mean_err[1],median_err[1]=MA_P_Predictions(ratios_jan,'jan')
ratios_jan,mean_err[2],median_err[2]=WA_R_Predictions(ratios_jan,'jan')
ratios_jan,mean_err[3],median_err[3]=WA_P_Predictions(ratios_jan,'jan')
ratios_jan,mean_err[4],median_err[4]=EA_R1_Predictions(ratios_jan,'jan')
ratios_jan,mean_err[5],median_err[5]=EA_P1_Predictions(ratios_jan,'jan')

```

## Comparison between baseline models

We have chosen our error metric for comparison between models as **MAPE (Mean Absolute Percentage Error)** so that we can know that on an average how good is our model with predictions and **MSE (Mean Squared Error)** is also used so that we have a clearer understanding as to how well our forecasting model performs with outliers so that we make sure that there is not much of a error margin between our prediction and the actual value

In [36]:

```

print ("Error Metric Matrix (Forecasting Methods) - MAPE & MSE")
print ("-----")
print ("Moving Averages (Ratios) - " , mean_err[0], " MAPE: ",mean_err[0], " MSE: ",me
ian_err[0])
print ("Moving Averages (2016 Values) - " , mean_err[1], " MAPE: ",mean_err[1], " MSE: ",m
edian_err[1])
print ("-----")
print ("Weighted Moving Averages (Ratios) - " , mean_err[2], " MAPE: ",mean_err[2], " MSE: ",me
dian_err[2])
print ("Weighted Moving Averages (2016 Values) - " , mean_err[3], " MAPE: ",mean_err[3], " MSE: ",me
dian_err[3])
print ("-----")
print ("Exponential Moving Averages (Ratios) - " , mean_err[4], " MAPE: ",mean_err[4], " MSE: ",media
n_err[4])
print ("Exponential Moving Averages (2016 Values) - " , mean_err[5], " MAPE: ",mean_err[5], " MSE: ",media
n_err[5])

```

Error Metric Matrix (Forecasting Methods) - MAPE & MSE

```

-----
-
Moving Averages (Ratios) - 7977897252 MAPE: 0.1628685972027342 MSE: 561.04
Moving Averages (2016 Values) - 4901433691756 MAPE: 0.12651735674574424 MSE: 241.
-----
-
Weighted Moving Averages (Ratios) - 548.5285170250896 MAPE: 0.1599661761243253 MSE:
Weighted Moving Averages (2016 Values) - 229.33734318996414 MAPE: 0.12121086157001072 MSE:
-----
-
Exponential Moving Averages (Ratios) - 546.5861260454002 MAPE: 0.1593403910652334 MSE:
Exponential Moving Averages (2016 Values) - 226.0377688172043 MAPE: 0.1209639974233378 MSE:

```

**Please Note:-** The above comparisons are made using Jan 2015 and Jan 2016 only

From the above matrix it is inferred that the best forecasting model for our prediction would be:-  $P_t' = \alpha * P_{t-1} + (1 - \alpha) * P_{t-1}'$  i.e Exponential Moving Averages using 2016 Values

# Regression Models

## Train-Test Split

Before we start predictions using the tree based regression models we take 3 months of 2016 pickup data and split it such that for every region we have 70% data in train and 30% in test, ordered date-wise for every region

In [37]:

```
# Preparing data to be split into train and test, The below prepares data in cumulative form which
# will be later split into test and train
# number of 10min indices for jan 2015= 24*31*60/10 = 4464
# number of 10min indices for jan 2016 = 24*31*60/10 = 4464
# number of 10min indices for feb 2016 = 24*29*60/10 = 4176
# number of 10min indices for march 2016 = 24*31*60/10 = 4464
# regions_cum: it will contain 40 lists, each list will contain 4464+4176+4464 values which represents
# the number of pickups
# that are happened for three months in 2016 data

# print(len(regions_cum))
# 40
# print(len(regions_cum[0]))
# 12960

# we take number of pickups that are happened in last 5 10min intravels
number_of_time_stamps = 5

# output varaiable
# it is list of lists
# it will contain number of pickups 13099 for each cluster
output = []

# tsne_lat will contain 13104-5=13099 times lattitude of cluster center for every cluster
# Ex: [[cent_lat 13099times],[cent_lat 13099times], [cent_lat 13099times].... 40 lists]
# it is list of lists
tsne_lat = []

# tsne_lon will contain 13104-5=13099 times logitude of cluster center for every cluster
# Ex: [[cent_long 13099times],[cent_long 13099times], [cent_long 13099times].... 40 lists]
# it is list of lists
tsne_lon = []

# we will code each day
# sunday = 0, monday=1, tue = 2, wed=3, thur=4, fri=5,sat=6
# for every cluster we will be adding 13099 values, each value represent to which day of the week
# that pickup bin belongs to
# it is list of lists
tsne_weekday = []

# its an numpy array, of shape (523960, 5)
# each row corresponds to an entry in out data
# for the first row we will have [f0,f1,f2,f3,f4] fi=number of pickups happened in i+1th 10min intravel(bin)
# the second row will have [f1,f2,f3,f4,f5]
# the third row will have [f2,f3,f4,f5,f6]
# and so on...
tsne_feature = []

tsne_feature = [0]*number_of_time_stamps
for i in range(0,30):
    tsne_lat.append([kmeans.cluster_centers_[i][0]]*13099)
    tsne_lon.append([kmeans.cluster_centers_[i][1]]*13099)
    # jan 1st 2016 is thursday, so we start our day from 4: "(int(k/144))%7+4"
    # our prediction start from 5th 10min intravel since we need to have number of pickups that are
    # happened in last 5 pickup bins
    tsne_weekday.append([int(((int(k/144))%7+4)%7) for k in range(5,4464+4176+4464)])
    # regions_cum is a list of lists [[x1,x2,x3..x13104], [x1,x2,x3..x13104], [x1,x2,x3..x13104],
    # [x1,x2,x3..x13104], [x1,x2,x3..x13104], .. 40 lsits]
    tsne_feature = np.vstack((tsne_feature, [regions_cum[i][r:r+number_of_time_stamps] for r in range(0,len(regions_cum[i])-number_of_time_stamps)]))
    output.append(regions_cum[i][5:])
    tsne_feature = tsne_feature[1:]
```

```
tsne_feature = tsne_feature[1:]
```

In [38]:

```
len(tsne_lat[0])*len(tsne_lat) == tsne_feature.shape[0] == len(tsne_weekday)*len(tsne_weekday[0]) == 30*13099 == len(output)*len(output[0])
```

Out [38]:

```
True
```

In [39]:

```
# Getting the predictions of exponential moving averages to be used as a feature in cumulative form

# upto now we computed 8 features for every data point that starts from 50th min of the day
# 1. cluster center latitude
# 2. cluster center longitude
# 3. day of the week
# 4. f_t_1: number of pickups that are happened previous t-1th 10min intravel
# 5. f_t_2: number of pickups that are happened previous t-2th 10min intravel
# 6. f_t_3: number of pickups that are happened previous t-3th 10min intravel
# 7. f_t_4: number of pickups that are happened previous t-4th 10min intravel
# 8. f_t_5: number of pickups that are happened previous t-5th 10min intravel

# from the baseline models we said the exponential weighted moving average gives us the best error
# we will try to add the same exponential weighted moving average at t as a feature to our data
# exponential weighted moving average => p'(t) = alpha*p'(t-1) + (1-alpha)*P(t-1)
alpha=0.3

# it is a temporary array that store exponential weighted moving average for each 10min intravel,
# for each cluster it will get reset
# for every cluster it contains 13104 values
predicted_values=[]

# it is similar like tsne_lat
# it is list of lists
# predict_list is a list of lists [[x5,x6,x7..x13104], [x5,x6,x7..x13104], [x5,x6,x7..x13104], [x5,x6,x7..x13104], [x5,x6,x7..x13104], .. 40 lsits]
predict_list = []
tsne_flat_exp_avg = []
for r in range(0,30):
    for i in range(0,13104):
        if i==0:
            predicted_value= regions_cum[r][0]
            predicted_values.append(0)
            continue
        predicted_values.append(predicted_value)
        predicted_value =int((alpha*predicted_value) + (1-alpha)*(regions_cum[r][i]))
    predict_list.append(predicted_values[5:])
    predicted_values=[]
```

## Fourier Features

In [40]:

```
amplitude = []
frequency = []
for i in range(30):
    amp = np.abs(np.fft.fft(regions_cum[i][0:13104])) #amplitude calculation
    freq = np.abs(np.fft.fft(freq(13104, 1))) #frequencies calculation
    amp_indices = np.argsort(-amp)[1:] #sorting amplitude
    amp_values = []
    freq_values = []
    for j in range(0, 9, 2): #taking top 5 amplitudes and frequencies
        amp_values.append(amp[amp_indices[j]])
        freq_values.append(freq[amp_indices[j]])
    for k in range(13104): #those top 5 frequencies and amplitudes are same for all the points in one cluster
        amplitude.append(amp_values)
        frequency.append(freq_values)
```

```
In [41]:
```

```
# train, test split : 70% 30% split
# Before we start predictions using the tree based regression models we take 3 months of 2016 pick
up data
# and split it such that for every region we have 70% data in train and 30% in test,
# ordered date-wise for every region
print("size of train data :", int(13099*0.7))
print("size of test data :", int(13099*0.3))
```

```
size of train data : 9169
size of test data : 3929
```

```
In [42]:
```

```
#for train and test data fourier frequencies
train_frequencies = [frequency[i*13099:(13099*i+9169)] for i in range(0,30)]
test_frequencies = [frequency[(13099*(i))+9169:13099*(i+1)] for i in range(0,30)]
```

```
In [43]:
```

```
#for train and test data fourier amplitudes
train_amplitudes = [amplitude[i*13099:(13099*i+9169)] for i in range(0,30)]
test_amplitudes = [amplitude[(13099*(i))+9169:13099*(i+1)] for i in range(0,30)]
```

```
In [44]:
```

```
# extracting first 9169 timestamp values i.e 70% of 13099 (total timestamps) for our training data
train_features = [tsne_feature[i*13099:(13099*i+9169)] for i in range(0,30)]
# temp = [0]*(12955 - 9068)
test_features = [tsne_feature[(13099*(i))+9169:13099*(i+1)] for i in range(0,30)]
```

```
In [45]:
```

```
print("Number of data clusters", len(train_features), "Number of data points in trian data",
len(train_features[0]), "Each data point contains", len(train_features[0][0]), "features")
print("Number of data clusters", len(train_features), "Number of data points in test data",
len(test_features[0]), "Each data point contains", len(test_features[0][0]), "features")
```

```
Number of data clusters 30 Number of data points in trian data 9169 Each data point contains 5 fea
tures
```

```
Number of data clusters 30 Number of data points in test data 3930 Each data point contains 5 feat
ures
```

```
In [46]:
```

```
# extracting first 9169 timestamp values i.e 70% of 13099 (total timestamps) for our training data
tsne_train_flat_lat = [i[:9169] for i in tsne_lat]
tsne_train_flat_lon = [i[:9169] for i in tsne_lon]
tsne_train_flat_weekday = [i[:9169] for i in tsne_weekday]
tsne_train_flat_output = [i[:9169] for i in output]
tsne_train_flat_exp_avg = [i[:9169] for i in predict_list]
```

```
In [47]:
```

```
# extracting the rest of the timestamp values i.e 30% of 12956 (total timestamps) for our test dat
a
tsne_test_flat_lat = [i[9169:] for i in tsne_lat]
tsne_test_flat_lon = [i[9169:] for i in tsne_lon]
tsne_test_flat_weekday = [i[9169:] for i in tsne_weekday]
tsne_test_flat_output = [i[9169:] for i in output]
tsne_test_flat_exp_avg = [i[9169:] for i in predict_list]
```

```
In [48]:
```

```
# the above contains values in the form of list of lists (i.e. list of values of each region), her
e we make all of them in one list
train_new_features = [
```

```
train_new_features = []
for i in range(0,30):
    train_new_features.extend(train_features[i])
test_new_features = []
for i in range(0,30):
    test_new_features.extend(test_features[i])
```

In [49]:

```
train_freq=[]
test_freq=[]
train_amp=[]
test_amp=[]
for i in range(0,30):
    train_freq.extend(train_frequencies[i])
    test_freq.extend(test_frequencies[i])
    train_amp.extend(train_amplitudes[i])
    test_amp.extend(test_amplitudes[i])
```

In [50]:

```
#stacking new features,frequencies and amplitudes horizontally using hstack.
train_brand_new_features=np.hstack((train_new_features,train_freq,train_amp))
test_brand_new_features=np.hstack((test_new_features,test_freq,test_amp))
```

In [51]:

```
# converting lists of lists into sinle list i.e flatten
# a = [[1,2,3,4],[4,6,7,8]]
# print(sum(a,[]))
# [1, 2, 3, 4, 4, 6, 7, 8]

tsne_train_lat = sum(tsne_train_flat_lat, [])
tsne_train_lon = sum(tsne_train_flat_lon, [])
tsne_train_weekday = sum(tsne_train_flat_weekday, [])
tsne_train_output = sum(tsne_train_flat_output, [])
tsne_train_exp_avg = sum(tsne_train_flat_exp_avg, [])
```

In [52]:

```
# converting lists of lists into sinle list i.e flatten
# a = [[1,2,3,4],[4,6,7,8]]
# print(sum(a,[]))
# [1, 2, 3, 4, 4, 6, 7, 8]

tsne_test_lat = sum(tsne_test_flat_lat, [])
tsne_test_lon = sum(tsne_test_flat_lon, [])
tsne_test_weekday = sum(tsne_test_flat_weekday, [])
tsne_test_output = sum(tsne_test_flat_output, [])
tsne_test_exp_avg = sum(tsne_test_flat_exp_avg, [])
```

In [53]:

```
# Preparing the data frame for our train data
columns = ['ft_5','ft_4','ft_3','ft_2','ft_1']
df_train = pd.DataFrame(data=train_new_features, columns=columns)
df_train['lat'] = tsne_train_lat
df_train['lon'] = tsne_train_lon
df_train['weekday'] = tsne_train_weekday
df_train['exp_avg'] = tsne_train_exp_avg

print(df_train.shape)
```

(275070, 9)

In [54]:

```
df_train.head()
```

Out[54]:

	ft_5	ft_4	ft_3	ft_2	ft_1	lat	lon	weekday	exp_avg
0	0	106	243	299	328	40.777809	-73.954054	4	309
1	106	243	299	328	340	40.777809	-73.954054	4	330
2	243	299	328	340	316	40.777809	-73.954054	4	320
3	299	328	340	316	327	40.777809	-73.954054	4	324
4	328	340	316	327	323	40.777809	-73.954054	4	323

In [55]:

```
# Preparing the data frame for our test data
df_test = pd.DataFrame(data=test_new_features, columns=columns)
df_test['lat'] = tsne_test_lat
df_test['lon'] = tsne_test_lon
df_test['weekday'] = tsne_test_weekday
df_test['exp_avg'] = tsne_test_exp_avg
print(df_test.shape)
```

(117900, 9)

In [56]:

```
df_test.head()
```

Out[56]:

	ft_5	ft_4	ft_3	ft_2	ft_1	lat	lon	weekday	exp_avg
0	240	213	243	222	234	40.777809	-73.954054	4	231
1	213	243	222	234	291	40.777809	-73.954054	4	273
2	243	222	234	291	256	40.777809	-73.954054	4	261
3	222	234	291	256	266	40.777809	-73.954054	4	264
4	234	291	256	266	268	40.777809	-73.954054	4	266

## Merging Fourier Features

In [57]:

```
# Preparing the data frame for our train data
columns = ['ft_5', 'ft_4', 'ft_3', 'ft_2', 'ft_1', 'fre1', 'fre2', 'fre3', 'fre4', 'fre5', 'amp1', 'amp2', 'amp3', 'amp4', 'amp5']
df_train_fr = pd.DataFrame(data=train_brand_new_features, columns=columns)
df_train_fr['lat'] = tsne_train_lat
df_train_fr['lon'] = tsne_train_lon
df_train_fr['weekday'] = tsne_train_weekday
df_train_fr['exp_avg'] = tsne_train_exp_avg

print(df_train_fr.shape)
```

(275070, 19)

In [58]:

```
# Preparing the data frame for our test data
df_test_fr = pd.DataFrame(data=test_brand_new_features, columns=columns)
df_test_fr['lat'] = tsne_test_lat
df_test_fr['lon'] = tsne_test_lon
df_test_fr['weekday'] = tsne_test_weekday
df_test_fr['exp_avg'] = tsne_test_exp_avg
print(df_test_fr.shape)
```

(117900, 19)

In [59]:

```
df_train_fr.head()
```

Out[59]:

	ft_5	ft_4	ft_3	ft_2	ft_1	fre1	fre2	fre3	fre4	fre5	amp1	amp2	amp3
0	0.0	106.0	243.0	299.0	328.0	0.006944	0.013889	0.012897	0.020833	0.007937	760430.104617	298303.16954	157723.832133
1	106.0	243.0	299.0	328.0	340.0	0.006944	0.013889	0.012897	0.020833	0.007937	760430.104617	298303.16954	157723.832133
2	243.0	299.0	328.0	340.0	316.0	0.006944	0.013889	0.012897	0.020833	0.007937	760430.104617	298303.16954	157723.832133
3	299.0	328.0	340.0	316.0	327.0	0.006944	0.013889	0.012897	0.020833	0.007937	760430.104617	298303.16954	157723.832133
4	328.0	340.0	316.0	327.0	323.0	0.006944	0.013889	0.012897	0.020833	0.007937	760430.104617	298303.16954	157723.832133

In [60]:

```
df_test_fr.head()
```

Out[60]:

	ft_5	ft_4	ft_3	ft_2	ft_1	fre1	fre2	fre3	fre4	fre5	amp1	amp2	amp3
0	240.0	213.0	243.0	222.0	234.0	0.006944	0.013889	0.012897	0.020833	0.007937	760430.104617	298303.16954	157723.832133
1	213.0	243.0	222.0	234.0	291.0	0.006944	0.013889	0.012897	0.020833	0.007937	760430.104617	298303.16954	157723.832133
2	243.0	222.0	234.0	291.0	256.0	0.006944	0.013889	0.012897	0.020833	0.007937	760430.104617	298303.16954	157723.832133
3	222.0	234.0	291.0	256.0	266.0	0.006944	0.013889	0.012897	0.020833	0.007937	760430.104617	298303.16954	157723.832133
4	234.0	291.0	256.0	266.0	268.0	0.006944	0.013889	0.012897	0.020833	0.007937	760430.104617	298303.16954	157723.832133

## Using Linear Regression

In [61]:

```
# find more about LinearRegression function here http://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LinearRegression.html
# -----
# default parameters
# sklearn.linear_model.LinearRegression(fit_intercept=True, normalize=False, copy_X=True, n_jobs=1)

# some of methods of LinearRegression()
# fit(X, y[, sample_weight]) Fit linear model.
# get_params([deep]) Get parameters for this estimator.
# predict(X) Predict using the linear model
# score(X, y[, sample_weight]) Returns the coefficient of determination R^2 of the prediction.
# set_params(**params) Set the parameters of this estimator.
# -----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/geometric-intuition-1-2-copy-8/
# -----


from sklearn.linear_model import LinearRegression
from sklearn.model_selection import GridSearchCV

parameters = {'fit_intercept':[True, False], 'normalize':[True, False]}

model = LinearRegression(n_jobs = -1)
lr = GridSearchCV(model, parameters, scoring = 'neg_mean_absolute_error', cv = 3)
lr.fit(df_train, tsne_train_output)
```

```

print("Best Parameters",lr.best_params_)

lr_pred_test = lr.predict(df_test)
lr_test_predictions = [round(value) for value in lr_pred_test]
lr_pred_train = lr.predict(df_train)
lr_train_predictions = [round(value) for value in lr_pred_train]

```

Best Parameters {'fit\_intercept': False, 'normalize': True}

In [62]:

```

train_MAPE_lr = mean_absolute_error(tsne_train_output, lr_pred_train)/ (sum(tsne_train_output)/len(tsne_train_output))
train_MSE_lr = mean_squared_error(tsne_train_output, lr_pred_train)

test_MAPE_lr = mean_absolute_error(tsne_test_output, lr_pred_test)/ (sum(tsne_test_output)/len(tsne_test_output))
test_MSE_lr = mean_squared_error(tsne_test_output, lr_pred_test)

```

In [63]:

```

print('Mean absolute percentage train error for Linear Regressor =',train_MAPE_lr)
print('Mean absolute percentage test error for Linear Regressor =',test_MAPE_lr)
print('-'*50)
print('Mean Squared train error for Linear Regressor =',train_MSE_lr)
print('Mean Squared test error for Linear Regressor =',test_MSE_lr)

```

```

Mean absolute percentage train error for Linear Regressor = 0.11979574640286851
Mean absolute percentage test error for Linear Regressor = 0.11589840442892027
-----
Mean Squared train error for Linear Regressor = 233.4872499180883
Mean Squared test error for Linear Regressor = 251.67848517852556

```

## Using Random Forest Regressor

In [64]:

```

# Training a hyper-parameter tuned random forest regressor on our train data
# find more about LinearRegression function here http://scikit-learn.org/stable/modules/generated/sklearn.ensemble.RandomForestRegressor.html
# -----
# default parameters
# sklearn.ensemble.RandomForestRegressor(n_estimators=10, criterion='mse', max_depth=None, min_samples_split=2,
#                                         min_samples_leaf=1, min_weight_fraction_leaf=0.0, max_features='auto', max_leaf_nodes=None, min_impurity_decrease=0.0,
#                                         min_impurity_split=None, bootstrap=True, oob_score=False, n_jobs=1, random_state=None, verbose=0, warm_start=False)

# some of methods of RandomForestRegressor()
# apply(X) Apply trees in the forest to X, return leaf indices.
# decision_path(X) Return the decision path in the forest
# fit(X, y[, sample_weight]) Build a forest of trees from the training set (X, y).
# get_params([deep]) Get parameters for this estimator.
# predict(X) Predict regression target for X.
# score(X, y[, sample_weight]) Returns the coefficient of determination R^2 of the prediction.
# -----
# video link1: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/regression-using-decision-trees-2/
# video link2: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/what-are-ensembles/
# -----


from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import RandomizedSearchCV
from scipy.stats import randint as sp_randint
from scipy.stats import uniform

parameters = {"n_estimators": [10, 20, 50, 100, 150, 200, 250, 300],
              "max_depth": [2, 5, 10, 20, 50, 100, 200, 250],
              "min_samples_split": sp_randint(120, 190),

```

```
        "min_samples_leaf": sp_randint(25,65)
    }

clf = RandomForestRegressor(random_state=25,n_jobs=-1)
model = RandomizedSearchCV(clf, param_distributions=parameters,
                           n_iter=5, cv=3, scoring='neg_mean_absolute_error', random_state=25)
model.fit(df_train_fr, tsne_train_output)
print(model.best_estimator_)
```

RandomForestRegressor(bootstrap=True, criterion='mse', max\_depth=50,  
 max\_features='auto', max\_leaf\_nodes=None,  
 min\_impurity\_decrease=0.0, min\_impurity\_split=None,  
 min\_samples\_leaf=51, min\_samples\_split=135,  
 min\_weight\_fraction\_leaf=0.0, n\_estimators=200, n\_jobs=-1,  
 oob\_score=False, random\_state=25, verbose=0,  
 warm\_start=False)

In [65]:

```
model = RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=50,  
                           max_features='auto', max_leaf_nodes=None,  
                           min_impurity_decrease=0.0, min_impurity_split=None,  
                           min_samples_leaf=51, min_samples_split=135,  
                           min_weight_fraction_leaf=0.0, n_estimators=200, n_jobs=-1,  
                           oob_score=False, random_state=25, verbose=0,  
                           warm_start=False)  
model.fit(df_train_fr, tsne_train_output)
```

Out[65]:

```
RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=50,  
                      max_features='auto', max_leaf_nodes=None,  
                      min_impurity_decrease=0.0, min_impurity_split=None,  
                      min_samples_leaf=51, min_samples_split=135,  
                      min_weight_fraction_leaf=0.0, n_estimators=200, n_jobs=-1,  
                      oob_score=False, random_state=25, verbose=0,  
                      warm_start=False)
```

In [66]:

```
# Predicting on test data using our trained random forest model  
  
# the models regr1 is already hyper parameter tuned  
# the parameters that we got above are found using grid search  
  
rf_pred_test = model.predict(df_test_fr)  
rndf_test_predictions = [round(value) for value in rf_pred_test]  
rf_pred_train = model.predict(df_train_fr)  
rndf_train_predictions = [round(value) for value in rf_pred_train]
```

In [67]:

```
train_MAPE_rf = mean_absolute_error(tsne_train_output, rf_pred_train) / (sum(tsne_train_output)/len(tsne_train_output))  
train_MSE_rf = mean_squared_error(tsne_train_output, rf_pred_train)  
  
test_MAPE_rf = mean_absolute_error(tsne_test_output, rf_pred_test) / (sum(tsne_test_output)/len(tsne_test_output))  
test_MSE_rf = mean_squared_error(tsne_test_output, rf_pred_test)
```

In [68]:

```
print('Mean absolute percentage train error for Random Forest Regressor =', train_MAPE_rf)  
print('Mean absolute percentage test error for Random Forest Regressor =', test_MAPE_rf)  
print('-'*50)  
print('Mean Squared train error for Random Forest Regressor =', train_MSE_rf)  
print('Mean Squared test error for Random Forest Regressor =', test_MSE_rf)
```

Mean absolute percentage train error for Random Forest Regressor = 0.1128182170644065  
Mean absolute percentage test error for Random Forest Regressor = 0.11307315190075964

```
Mean Squared train error for Random Forest Regressor = 209.43061630822763
Mean Squared test error for Random Forest Regressor= 242.6311733695796
```

## Using XgBoost Regressor

In [70]:

```
# Training a hyper-parameter tuned Xg-Boost regressor on our train data

# find more about XGBRegressor function here
http://xgboost.readthedocs.io/en/latest/python/python_api.html?#module-xgboost.sklearn
# -----
# default paramters
# xgboost.XGBRegressor(max_depth=3, learning_rate=0.1, n_estimators=100, silent=True,
objective='reg:linear',
# booster='gbtree', n_jobs=1, nthread=None, gamma=0, min_child_weight=1, max_delta_step=0, subsample=1, colsample_bytree=1,
# colsample_bylevel=1, reg_alpha=0, reg_lambda=1, scale_pos_weight=1, base_score=0.5,
random_state=0, seed=None,
# missing=None, **kwargs)

# some of methods of RandomForestRegressor()
# fit(X, y, sample_weight=None, eval_set=None, eval_metric=None, early_stopping_rounds=None, verbose=True, xgb_model=None)
# get_params([deep]) Get parameters for this estimator.
# predict(data, output_margin=False, ntree_limit=0) : Predict with data. NOTE: This function is not thread safe.
# get_score(importance_type='weight') -> get the feature importance
# -----
# video link1: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/regression-using-decision-trees-2/
# video link2: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/what-are-ensembles/
# -----
import warnings
warnings.filterwarnings("ignore")

import xgboost as xgb
clf = xgb.XGBRegressor()

parameters = {"n_estimators": [10, 20, 50, 100, 150, 200, 250, 300],
              "max_depth": [2, 5, 10, 20, 50, 100, 200, 250]
              }
model = RandomizedSearchCV(clf, param_distributions=parameters,
                           n_iter=5, cv=3, scoring='neg_mean_absolute_error', random_state=25)

model.fit(df_train_fr, tsne_train_output)
print(model.best_estimator_)
```

[23:06:57] WARNING: C:/Jenkins/workspace/xgboost-win64\_release\_0.90/src/objective/regression\_obj.cu:152: reg:linear is now deprecated in favor of reg:squarederror.  
[23:14:01] WARNING: C:/Jenkins/workspace/xgboost-win64\_release\_0.90/src/objective/regression\_obj.cu:152: reg:linear is now deprecated in favor of reg:squarederror.  
[23:21:01] WARNING: C:/Jenkins/workspace/xgboost-win64\_release\_0.90/src/objective/regression\_obj.cu:152: reg:linear is now deprecated in favor of reg:squarederror.  
[23:28:04] WARNING: C:/Jenkins/workspace/xgboost-win64\_release\_0.90/src/objective/regression\_obj.cu:152: reg:linear is now deprecated in favor of reg:squarederror.  
[23:28:07] WARNING: C:/Jenkins/workspace/xgboost-win64\_release\_0.90/src/objective/regression\_obj.cu:152: reg:linear is now deprecated in favor of reg:squarederror.  
[23:28:10] WARNING: C:/Jenkins/workspace/xgboost-win64\_release\_0.90/src/objective/regression\_obj.cu:152: reg:linear is now deprecated in favor of reg:squarederror.  
[23:28:14] WARNING: C:/Jenkins/workspace/xgboost-win64\_release\_0.90/src/objective/regression\_obj.cu:152: reg:linear is now deprecated in favor of reg:squarederror.  
[23:29:36] WARNING: C:/Jenkins/workspace/xgboost-win64\_release\_0.90/src/objective/regression\_obj.cu:152: reg:linear is now deprecated in favor of reg:squarederror.  
[23:31:00] WARNING: C:/Jenkins/workspace/xgboost-

```
win64_release_0.90/src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of
reg:squarederror.
[23:32:24] WARNING: C:/Jenkins/workspace/xgboost-
win64_release_0.90/src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of
reg:squarederror.
[23:33:30] WARNING: C:/Jenkins/workspace/xgboost-
win64_release_0.90/src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of
reg:squarederror.
[23:34:37] WARNING: C:/Jenkins/workspace/xgboost-
win64_release_0.90/src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of
reg:squarederror.
[23:35:44] WARNING: C:/Jenkins/workspace/xgboost-
win64_release_0.90/src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of
reg:squarederror.
[23:52:23] WARNING: C:/Jenkins/workspace/xgboost-
win64_release_0.90/src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of
reg:squarederror.
[00:08:51] WARNING: C:/Jenkins/workspace/xgboost-
win64_release_0.90/src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of
reg:squarederror.
[00:21:45] WARNING: C:/Jenkins/workspace/xgboost-
win64_release_0.90/src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of
reg:squarederror.
XGBRegressor(base_score=0.5, booster='gbtree', colsample_bylevel=1,
             colsample_bynode=1, colsample_bytree=1, gamma=0,
             importance_type='gain', learning_rate=0.1, max_delta_step=0,
             max_depth=2, min_child_weight=1, missing=None, n_estimators=300,
             n_jobs=1, nthread=None, objective='reg:linear', random_state=0,
             reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
             silent=None, subsample=1, verbosity=1)
```

In [71]:

```
model=xgb.XGBRegressor(base_score=0.5, booster='gbtree', colsample_bylevel=1,
                       colsample_bynode=1, colsample_bytree=1, gamma=0,
                       importance_type='gain', learning_rate=0.1, max_delta_step=0,
                       max_depth=2, min_child_weight=1, missing=None, n_estimators=300,
                       n_jobs=1, nthread=None, objective='reg:linear', random_state=0,
                       reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
                       silent=None, subsample=1, verbosity=1)
model.fit(df_train_fr,tsne_train_output)
```

```
[00:23:56] WARNING: C:/Jenkins/workspace/xgboost-
win64_release_0.90/src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of
reg:squarederror.
```

Out[71]:

```
XGBRegressor(base_score=0.5, booster='gbtree', colsample_bylevel=1,
             colsample_bynode=1, colsample_bytree=1, gamma=0,
             importance_type='gain', learning_rate=0.1, max_delta_step=0,
             max_depth=2, min_child_weight=1, missing=None, n_estimators=300,
             n_jobs=1, nthread=None, objective='reg:linear', random_state=0,
             reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
             silent=None, subsample=1, verbosity=1)
```

In [72]:

```
#predicting with our trained Xg-Boost regressor
# the models x_model is already hyper parameter tuned
# the parameters that we got above are found using grid search

xg_pred_test = model.predict(df_test_fr)
xgb_test_predictions = [round(value) for value in xg_pred_test]
xg_pred_train = model.predict(df_train_fr)
xgb_train_predictions = [round(value) for value in xg_pred_train]
```

In [73]:

```
train_MAPE_xgb = mean_absolute_error(tsne_train_output, xg_pred_train)/ (sum(tsne_train_output)/len
(tsne_train_output))
train_MSE_xgb = mean_squared_error(tsne_train_output, xg_pred_train)
```

```
test_MAPE_xgb = mean_absolute_error(tsne_test_output, xg_pred_test)/ (sum(tsne_test_output)/len(tsne_test_output))
test_MSE_xgb = mean_squared_error(tsne_test_output, xg_pred_test)
```

In [74]:

```
print('Mean absolute percentage train error for XGB Regressor =',train_MAPE_xgb)
print('Mean absolute percentage test error for XGB Regressor =',test_MAPE_xgb)
print('-'*50)
print('Mean Squared train error for XGB Regressor =',train_MSE_xgb)
print('Mean Squared test error for XGB Regressor=',test_MSE_xgb)
```

```
Mean absolute percentage train error for XGB Regressor = 0.11862981653127391
Mean absolute percentage test error for XGB Regressor = 0.11543768319330362
-----
Mean Squared train error for XGB Regressor = 226.81898352759362
Mean Squared test error for XGB Regressor= 248.9385963751369
```

## Calculating the error metric values for various models

In [77]:

```
train_mape=[]
test_mape=[]

train_mape.append((mean_absolute_error(tsne_train_output,df_train['ft_1'].values))/(sum(tsne_train_output)/len(tsne_train_output)))
train_mape.append((mean_absolute_error(tsne_train_output,df_train['exp_avg'].values))/(sum(tsne_train_output)/len(tsne_train_output)))
train_mape.append((mean_absolute_error(tsne_train_output,rndf_train_predictions))/(sum(tsne_train_output)/len(tsne_train_output)))
train_mape.append((mean_absolute_error(tsne_train_output,
xgb_train_predictions))/(sum(tsne_train_output)/len(tsne_train_output)))
train_mape.append((mean_absolute_error(tsne_train_output,
lr_train_predictions))/(sum(tsne_train_output)/len(tsne_train_output)))

test_mape.append((mean_absolute_error(tsne_test_output, df_test['ft_1'].values))/(sum(tsne_test_output)/len(tsne_test_output)))
test_mape.append((mean_absolute_error(tsne_test_output,
df_test['exp_avg'].values))/(sum(tsne_test_output)/len(tsne_test_output)))
test_mape.append((mean_absolute_error(tsne_test_output,
rndf_test_predictions))/(sum(tsne_test_output)/len(tsne_test_output)))
test_mape.append((mean_absolute_error(tsne_test_output,
xgb_test_predictions))/(sum(tsne_test_output)/len(tsne_test_output)))
test_mape.append((mean_absolute_error(tsne_test_output,
lr_test_predictions))/(sum(tsne_test_output)/len(tsne_test_output)))
```

In [78]:

```
print ("Error Metric Matrix (Tree Based Regression Methods) - MAPE")
print ("-----")
print ("Baseline Model - Train: ",train_mape[0]," Test: ",test_mape[0])
print ("Exponential Averages Forecasting - Train: ",train_mape[1]," Test: ",test_mape[1])
print ("Linear Regression - Train: ",train_mape[3]," Test: ",test_mape[3])
print ("Random Forest Regression - Train: ",train_mape[2]," Test: ",test_mape[2])
```

```
Error Metric Matrix (Tree Based Regression Methods) - MAPE
-----
-
Baseline Model - Train: 0.12477882091940766 Test:
0.12137217161272074
Exponential Averages Forecasting - Train: 0.11976904266333344 Test:
0.11613179453264473
Linear Regression - Train: 0.11856445520314589 Test:
0.11540288391331441
Random Forest Regression - Train: 0.11271998849737422 Test:
0.11300542071723199
```

## Error Metric Matrix

In [79]:

```
print ("Error Metric Matrix (Tree Based Regression Methods) - MAPE")
print ("-----")
print ("Baseline Model - Train: ",train_mape[0], " Test: ",test_map
[0])
print ("Exponential Averages Forecasting - Train: ",train_mape[1], " Test: ",test_map
e[1])
print ("Linear Regression - Train: ",train_mape[4], " Test: ",test_mape
4])
print ("Random Forest Regression - Train: ",train_mape[2], " Test: ",test_mape
[2])
print ("XgBoost Regression - Train: ",train_mape[3], " Test: ",test_map
[3])
print ("-----")
```

Error Metric Matrix (Tree Based Regression Methods) - MAPE

```
-----
-
Baseline Model - Train: 0.12477882091940766 Test:
0.12137217161272074
Exponential Averages Forecasting - Train: 0.11976904266333344 Test:
0.11613179453264473
Linear Regression - Train: 0.11972981297412652 Test:
0.11583027191510348
Random Forest Regression - Train: 0.11271998849737422 Test:
0.11300542071723199
XgBoost Regression - Train: 0.11856445520314589 Test: 0.11540288391
31441
-----
```

## Prettytable

In [81]:

```
from prettytable import PrettyTable
x = PrettyTable()
x = PrettyTable(["Models", "MAPE train", "MAPE test"])
x.add_row(['Linear Regression',train_mape[4],test_mape[4]])
x.add_row(['Random Forest',train_mape[2],test_mape[2]])
x.add_row(['XGBOOST',train_mape[3],test_mape[3]])
print(x)
```

Models	MAPE train	MAPE test
Linear Regression	0.11972981297412652	0.11583027191510348
Random Forest	0.11271998849737422	0.11300542071723199
XGBOOST	0.11856445520314589	0.11540288391331441

Our best MAPE test is 11.30 % for RandomForest model after adding fourier features amplitude and frequencies. We have used cluster size = 30

## For Cluster size = 40

In [11]:

```
# if check for the 50 clusters you can observe that there are two clusters with only 0.3 miles apart from each other
# so we choose 40 clusters for solve the further problem
```

```
# Getting 40 clusters using the kmeans
kmeans = MiniBatchKMeans(n_clusters=40, batch_size=10000, random_state=0).fit(coords)
frame_with_durations_outliers_removed['pickup_cluster'] =
kmeans.predict(frame_with_durations_outliers_removed[['pickup_latitude', 'pickup_longitude']])
```

## Time-Binning

In [12]:

```
#Refer:https://www.unixtimestamp.com/
# 1420070400 : 2015-01-01 00:00:00
# 1422748800 : 2015-02-01 00:00:00
# 1425168000 : 2015-03-01 00:00:00
# 1427846400 : 2015-04-01 00:00:00
# 1430438400 : 2015-05-01 00:00:00
# 1433116800 : 2015-06-01 00:00:00

# 1451606400 : 2016-01-01 00:00:00
# 1454284800 : 2016-02-01 00:00:00
# 1456790400 : 2016-03-01 00:00:00
# 1459468800 : 2016-04-01 00:00:00
# 1462060800 : 2016-05-01 00:00:00
# 1464739200 : 2016-06-01 00:00:00

def add_pickup_bins(frame,month,year):
    unix_pickup_times=[i for i in frame['pickup_times'].values]
    unix_times = [[1420070400,1422748800,1425168000,1427846400,1430438400,1433116800],\
                   [1451606400,1454284800,1456790400,1459468800,1462060800,1464739200]]

    start_pickup_unix=unix_times[year-2015][month-1]
    # https://www.timeanddate.com/time/zones/est
    # (int((i-start_pickup_unix)/600)+33) : our unix time is in gmt to we are converting it to est
    tenminutewise_binned_unix_pickup_times=[(int((i-start_pickup_unix)/600)+33) for i in unix_pickup_times]
    frame['pickup_bins'] = np.array(tenminutewise_binned_unix_pickup_times)
    return frame
```

In [13]:

```
# clustering, making pickup bins and grouping by pickup cluster and pickup bins
frame_with_durations_outliers_removed['pickup_cluster'] =
kmeans.predict(frame_with_durations_outliers_removed[['pickup_latitude', 'pickup_longitude']])
jan_2015_frame = add_pickup_bins(frame_with_durations_outliers_removed,1,2015)
jan_2015_groupby =
jan_2015_frame[['pickup_cluster','pickup_bins','trip_distance']].groupby(['pickup_cluster','pickup_bins']).count()
```

In [14]:

```
# we add two more columns 'pickup_cluster' (to which cluster it belongs to)
# and 'pickup_bins' (to which 10min intravel the trip belongs to)
jan_2015_frame.head()
```

Out[14]:

	passenger_count	trip_distance	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	total_amount	trip_time	pickup_cluster
0	1	1.59	-73.993896	40.750111	-73.974785	40.750618	17.05	18.050000	1.420
1	1	3.30	-74.001648	40.724243	-73.994415	40.759109	17.80	19.833333	1.420
2	1	1.80	-73.963341	40.802788	-73.951820	40.824413	10.80	10.050000	1.420
3	1	0.50	-74.009087	40.713818	-74.004326	40.719986	4.80	1.866667	1.420
4	1	3.00	-73.971176	40.762428	-74.004181	40.742653	16.30	19.316667	1.420

In [15]:

```
# hear the trip_distance represents the number of pickups that are happend in that particular 10mi
n intravel
# this data frame has two indices
```

```
" Since data frame has one index
# primary index: pickup_cluster (cluster number)
# secondary index : pickup_bins (we devide whole months time into 10min intravels 24*31*60/10 =4464 bins)
jan_2015_groupby.head()
```

Out[15]:

trip_distance		
pickup_cluster	pickup_bins	
0	1	105
	2	199
	3	208
	4	141
	5	155

In [16]:

```
# upto now we cleaned data and prepared data for the month 2015,
# now do the same operations for months Jan, Feb, March of 2016
# 1. get the dataframe which includes only required columns
# 2. adding trip times, speed, unix time stamp of pickup_time
# 4. remove the outliers based on trip_times, speed, trip_duration, total_amount
# 5. add pickup_cluster to each data point
# 6. add pickup_bin (index of 10min intravel to which that trip belongs to)
# 7. group by data, based on 'pickup_cluster' and 'pickup_bin'

# Data Preparation for the months of Jan, Feb and March 2016
def datapreparation(month,kmeans,month_no,year_no):

    print ("Return with trip times..")

    frame_with_durations = return_with_trip_times(month)

    print ("Remove outliers..")
    frame_with_durations_outliers_removed = remove_outliers(frame_with_durations)

    print ("Estimating clusters..")
    frame_with_durations_outliers_removed['pickup_cluster'] =
    kmeans.predict(frame_with_durations_outliers_removed[['pickup_latitude', 'pickup_longitude']])
    #frame_with_durations_outliers_removed_2016['pickup_cluster'] =
    #kmeans.predict(frame_with_durations_outliers_removed_2016[['pickup_latitude',
    #'pickup_longitude']])

    print ("Final groupbying..")
    final_updated_frame = add_pickup_bins(frame_with_durations_outliers_removed,month_no,year_no)
    final_groupby_frame = final_updated_frame[['pickup_cluster','pickup_bins','trip_distance']].groupby(['pickup_cluster','pickup_bins']).count()

    return final_updated_frame,final_groupby_frame

month_jan_2016 = dd.read_csv('yellow_tripdata_2016-01.csv')
month_feb_2016 = dd.read_csv('yellow_tripdata_2016-02.csv')
month_mar_2016 = dd.read_csv('yellow_tripdata_2016-03.csv')

jan_2016_frame,jan_2016_groupby = datapreparation(month_jan_2016,kmeans,1,2016)
feb_2016_frame,feb_2016_groupby = datapreparation(month_feb_2016,kmeans,2,2016)
mar_2016_frame,mar_2016_groupby = datapreparation(month_mar_2016,kmeans,3,2016)
```

```
Return with trip times..
Remove outliers..
Number of pickup records = 10906858
Number of outlier coordinates lying outside NY boundaries: 214677
Number of outliers from trip times analysis: 27190
Number of outliers from trip distance analysis: 79742
Number of outliers from speed analysis: 21047
Number of outliers from fare analysis: 4991
Total outliers removed 297784
---
Estimating clusters..
```

```

Final groupbying..
Return with trip times..
Remove outliers..
Number of pickup records = 11382049
Number of outlier coordinates lying outside NY boundaries: 223161
Number of outliers from trip times analysis: 27670
Number of outliers from trip distance analysis: 81902
Number of outliers from speed analysis: 22437
Number of outliers from fare analysis: 5476
Total outliers removed 308177
---
Estimating clusters..
Final groupbying..
Return with trip times..
Remove outliers..
Number of pickup records = 12210952
Number of outlier coordinates lying outside NY boundaries: 232444
Number of outliers from trip times analysis: 30868
Number of outliers from trip distance analysis: 87318
Number of outliers from speed analysis: 23889
Number of outliers from fare analysis: 5859
Total outliers removed 324635
---
Estimating clusters..
Final groupbying..

```

## Smoothing

In [17]:

```

# Gets the unique bins where pickup values are present for each each reigion

# for each cluster region we will collect all the indices of 10min intravels in which the pickups
# are happened
# we got an observation that there are some pickpbins that doesnt have any pickups
def return_unq_pickup_bins(frame):
    values = []
    for i in range(0,40):
        new = frame[frame['pickup_cluster'] == i]
        list_unq = list(set(new['pickup_bins']))
        list_unq.sort()
        values.append(list_unq)
    return values

```

In [18]:

```

# for every month we get all indices of 10min intravels in which atleast one pickup got happened

#jan
jan_2015_unique = return_unq_pickup_bins(jan_2015_frame)
jan_2016_unique = return_unq_pickup_bins(jan_2016_frame)

#feb
feb_2016_unique = return_unq_pickup_bins(feb_2016_frame)

#march
mar_2016_unique = return_unq_pickup_bins(mar_2016_frame)

```

In [19]:

```

# for each cluster number of 10min intavels with 0 pickups
for i in range(40):
    print("for the ",i,"th cluster number of 10min intavels with zero pickups: ",4464 -
len(set(jan_2015_unique[i])))
    print('-'*60)

```

```

for the 0 th cluster number of 10min intavels with zero pickups: 41
-----
for the 1 th cluster number of 10min intavels with zero pickups: 1986
-----
for the 2 th cluster number of 10min intavels with zero pickups: 30
-----
```

```
for the 3 th cluster number of 10min intavels with zero pickups: 355
-----
for the 4 th cluster number of 10min intavels with zero pickups: 38
-----
for the 5 th cluster number of 10min intavels with zero pickups: 154
-----
for the 6 th cluster number of 10min intavels with zero pickups: 35
-----
for the 7 th cluster number of 10min intavels with zero pickups: 34
-----
for the 8 th cluster number of 10min intavels with zero pickups: 118
-----
for the 9 th cluster number of 10min intavels with zero pickups: 41
-----
for the 10 th cluster number of 10min intavels with zero pickups: 26
-----
for the 11 th cluster number of 10min intavels with zero pickups: 45
-----
for the 12 th cluster number of 10min intavels with zero pickups: 43
-----
for the 13 th cluster number of 10min intavels with zero pickups: 29
-----
for the 14 th cluster number of 10min intavels with zero pickups: 27
-----
for the 15 th cluster number of 10min intavels with zero pickups: 32
-----
for the 16 th cluster number of 10min intavels with zero pickups: 41
-----
for the 17 th cluster number of 10min intavels with zero pickups: 59
-----
for the 18 th cluster number of 10min intavels with zero pickups: 1191
-----
for the 19 th cluster number of 10min intavels with zero pickups: 1358
-----
for the 20 th cluster number of 10min intavels with zero pickups: 54
-----
for the 21 th cluster number of 10min intavels with zero pickups: 30
-----
for the 22 th cluster number of 10min intavels with zero pickups: 30
-----
for the 23 th cluster number of 10min intavels with zero pickups: 164
-----
for the 24 th cluster number of 10min intavels with zero pickups: 36
-----
for the 25 th cluster number of 10min intavels with zero pickups: 42
-----
for the 26 th cluster number of 10min intavels with zero pickups: 32
-----
for the 27 th cluster number of 10min intavels with zero pickups: 215
-----
for the 28 th cluster number of 10min intavels with zero pickups: 37
-----
for the 29 th cluster number of 10min intavels with zero pickups: 42
-----
for the 30 th cluster number of 10min intavels with zero pickups: 1181
-----
for the 31 th cluster number of 10min intavels with zero pickups: 43
-----
for the 32 th cluster number of 10min intavels with zero pickups: 45
-----
for the 33 th cluster number of 10min intavels with zero pickups: 44
-----
for the 34 th cluster number of 10min intavels with zero pickups: 40
-----
for the 35 th cluster number of 10min intavels with zero pickups: 43
-----
for the 36 th cluster number of 10min intavels with zero pickups: 37
-----
for the 37 th cluster number of 10min intavels with zero pickups: 322
-----
for the 38 th cluster number of 10min intavels with zero pickups: 37
-----
for the 39 th cluster number of 10min intavels with zero pickups: 44
-----
```

In [20]:

```

# Fills a value of zero for every bin where no pickup data is present
# the count_values: number pickups that are happened in each region for each 10min intravel
# there wont be any value if there are no pickups.
# values: number of unique bins

# for every 10min intravel(pickup_bin) we will check it is there in our unique bin,
# if it is there we will add the count_values[index] to smoothed data
# if not we add 0 to the smoothed data
# we finally return smoothed data
def fill_missing(count_values,values):
    smoothed_regions=[]
    ind=0
    for r in range(0,40):
        smoothed_bins=[]
        for i in range(4464):
            if i in values[r]:
                smoothed_bins.append(count_values[ind])
                ind+=1
            else:
                smoothed_bins.append(0)
        smoothed_regions.extend(smoothed_bins)
    return smoothed_regions

```

In [21]:

```

# Fills a value of zero for every bin where no pickup data is present
# the count_values: number pickups that are happened in each region for each 10min intravel
# there wont be any value if there are no pickups.
# values: number of unique bins

# for every 10min intravel(pickup_bin) we will check it is there in our unique bin,
# if it is there we will add the count_values[index] to smoothed data
# if not we add smoothed data (which is calculated based on the methods that are discussed in the
# above markdown cell)
# we finally return smoothed data
def smoothing(count_values,values):
    smoothed_regions=[] # stores list of final smoothed values of each region
    ind=0
    repeat=0
    smoothed_value=0
    for r in range(0,40):
        smoothed_bins=[] #stores the final smoothed values
        repeat=0
        for i in range(4464):
            if repeat!=0: # prevents iteration for a value which is already visited/resolved
                repeat-=1
                continue
            if i in values[r]: #checks if the pickup-bin exists
                smoothed_bins.append(count_values[ind]) # appends the value of the pickup bin if it
exists
            else:
                if i!=0:
                    right_hand_limit=0
                    for j in range(i,4464):
                        if j not in values[r]: #searches for the left-limit or the pickup-bin
value which has a pickup value
                            continue
                        else:
                            right_hand_limit=j
                            break
                    if right_hand_limit==0:
                        #Case 1: When we have the last/last few values are found to be missing,hence we
have no right-limit here
                        smoothed_value=count_values[ind-1]*1.0/((4463-i)+2)*1.0
                        for j in range(i,4464):
                            smoothed_bins.append(math.ceil(smoothed_value))
                            smoothed_bins[i-1] = math.ceil(smoothed_value)
                            repeat=(4463-i)
                            ind-=1
                        else:
                            #Case 2: When we have the missing values between two known values
                            smoothed_value=(count_values[ind-1]+count_values[ind])*1.0/((right_hand_lim
t-i)+2)*1.0
                            for j in range(i,right_hand_limit+1):

```

```

        smoothed_bins.append(math.ceil(smoothed_value))
        smoothed_bins[i-1] = math.ceil(smoothed_value)
        repeat=(right_hand_limit-i)
    else:
        #Case 3: When we have the first/first few values are found to be missing,hence
        we have no left-limit here
        right_hand_limit=0
        for j in range(i,4464):
            if j not in values[r]:
                continue
            else:
                right_hand_limit=j
                break
        smoothed_value=count_values[ind]*1.0/((right_hand_limit-i)+1)*1.0
        for j in range(i,right_hand_limit+1):
            smoothed_bins.append(math.ceil(smoothed_value))
        repeat=(right_hand_limit-i)
    ind+=1
    smoothed_regions.extend(smoothed_bins)
return smoothed_regions

```

In [22]:

```

#Filling Missing values of Jan-2015 with 0
# here in jan_2015_groupby dataframe the trip_distance represents the number of pickups that are h
appened
jan_2015_fill = fill_missing(jan_2015_groupby['trip_distance'].values,jan_2015_unique)

#Smoothing Missing values of Jan-2015
jan_2015_smooth = smoothing(jan_2015_groupby['trip_distance'].values,jan_2015_unique)

```

In [23]:

```

# number of 10min indices for jan 2015= 24*31*60/10 = 4464
# number of 10min indices for jan 2016 = 24*31*60/10 = 4464
# number of 10min indices for feb 2016 = 24*29*60/10 = 4176
# number of 10min indices for march 2016 = 24*30*60/10 = 4320
# for each cluster we will have 4464 values, therefore 40*4464 = 178560 (length of the
jan_2015_fill)
print("number of 10min intravels among all the clusters ",len(jan_2015_fill))

```


number of 10min intravels among all the clusters 178560


In [24]:

```

# Smoothing vs Filling
# sample plot that shows two variations of filling missing values
# we have taken the number of pickups for cluster region 2
plt.figure(figsize=(10,5))
plt.plot(jan_2015_fill[4464:8920], label="zero filled values")
plt.plot(jan_2015_smooth[4464:8920], label="filled with avg values")
plt.legend()
plt.show()

```





In [25]:

```
# Jan-2015 data is smoothed, Jan,Feb & March 2016 data missing values are filled with zero
jan_2015_smooth = smoothing(jan_2015_groupby['trip_distance'].values,jan_2015_unique)
jan_2016_smooth = fill_missing(jan_2016_groupby['trip_distance'].values,jan_2016_unique)
feb_2016_smooth = fill_missing(feb_2016_groupby['trip_distance'].values,feb_2016_unique)
mar_2016_smooth = fill_missing(mar_2016_groupby['trip_distance'].values,mar_2016_unique)

# Making list of all the values of pickup data in every bin for a period of 3 months and storing them region-wise
regions_cum = []

# a =[1,2,3]
# b = [2,3,4]
# a+b = [1, 2, 3, 2, 3, 4]

# number of 10min indices for jan 2015= 24*31*60/10 = 4464
# number of 10min indices for jan 2016 = 24*31*60/10 = 4464
# number of 10min indices for feb 2016 = 24*29*60/10 = 4176
# number of 10min indices for march 2016 = 24*31*60/10 = 4464
# regions_cum: it will contain 40 lists, each list will contain 4464+4176+4464 values which represents the number of pickups
# that are happened for three months in 2016 data


for i in range(0,40):
    regions_cum.append(jan_2016_smooth[4464*i:4464*(i+1)]+feb_2016_smooth[4176*i:4176*(i+1)]+mar_2016_smooth[4464*i:4464*(i+1)])





# print(len(regions_cum))
# 40
# print(len(regions_cum[0]))
# 13104
```





## Time Series and Fourier Transform






In [26]:





```
def uniqueish_color():
    """There're better ways to generate unique colors, but this isn't awful."""
    return plt.cm.gist_ncar(np.random.random())
first_x = list(range(0,4464))
second_x = list(range(4464,8640))
third_x = list(range(8640,13104))
for i in range(40):
    plt.figure(figsize=(10,4))
    plt.plot(first_x,regions_cum[i][:4464], color=uniqueish_color(), label='2016 Jan month data')
    plt.plot(second_x,regions_cum[i][4464:8640], color=uniqueish_color(), label='2016 feb month data')
    plt.plot(third_x,regions_cum[i][8640:], color=uniqueish_color(), label='2016 march month data')
    plt.legend()
    plt.show()
```






















































```
0      2000     4000     6000     8000    10000    12000
```

In [27]:

```
# getting peaks: https://blog.ytotech.com/2015/11/01/findpeaks-in-python/
# read more about fft function :
https://docs.scipy.org/doc/numpy/reference/generated/numpy.fft.fft.html
Y      = np.fft.fft(np.array(jan_2016_smooth)[0:4460])
# read more about the fftfreq:
https://docs.scipy.org/doc/numpy/reference/generated/numpy.fft.fftfreq.html
freq = np.fft.fftfreq(4460, 1)
n = len(freq)
plt.figure()
plt.plot( freq[:int(n/2)], np.abs(Y)[:int(n/2)] )
plt.xlabel("Frequency")
plt.ylabel("Amplitude")
plt.show()
```



In [28]:

```
#Preparing the Dataframe only with x(i) values as jan-2015 data and y(i) values as jan-2016
ratios_jan = pd.DataFrame()
ratios_jan['Given']=jan_2015_smooth
ratios_jan['Prediction']=jan_2016_smooth
ratios_jan['Ratios']=ratios_jan['Prediction']*1.0/ratios_jan['Given']*1.0
```

## Regression models

### Train-Test split

In [29]:

```
# Preparing data to be split into train and test, The below prepares data in cumulative form which
will be later split into test and train
# number of 10min indices for jan 2015= 24*31*60/10 = 4464
# number of 10min indices for jan 2016 = 24*31*60/10 = 4464
# number of 10min indices for feb 2016 = 24*29*60/10 = 4176
# number of 10min indices for march 2016 = 24*31*60/10 = 4464
# regions_cum: it will contain 40 lists, each list will contain 4464+4176+4464 values which represents the number of pickups
# that are happened for three months in 2016 data
```

```

# this are happened for three months in 2016 data

# print(len(regions_cum))
# 40
# print(len(regions_cum[0]))
# 12960

# we take number of pickups that are happened in last 5 10min intravels
number_of_time_stamps = 5

# output variable
# it is list of lists
# it will contain number of pickups 13099 for each cluster
output = []

# tsne_lat will contain 13104-5=13099 times lattitude of cluster center for every cluster
# Ex: [[cent_lat 13099times], [cent_lat 13099times], [cent_lat 13099times].... 40 lists]
# it is list of lists
tsne_lat = []

# tsne_lon will contain 13104-5=13099 times logitude of cluster center for every cluster
# Ex: [[cent_long 13099times], [cent_long 13099times], [cent_long 13099times].... 40 lists]
# it is list of lists
tsne_lon = []

# we will code each day
# sunday = 0, monday=1, tue = 2, wed=3, thur=4, fri=5,sat=6
# for every cluster we will be adding 13099 values, each value represent to which day of the week
# that pickup bin belongs to
# it is list of lists
tsne_weekday = []

# its an numpy array, of shape (523960, 5)
# each row corresponds to an entry in out data
# for the first row we will have [f0,f1,f2,f3,f4] fi=number of pickups happened in i+1th 10min int
ravel(bin)
# the second row will have [f1,f2,f3,f4,f5]
# the third row will have [f2,f3,f4,f5,f6]
# and so on...
tsne_feature = []

tsne_feature = [0]*number_of_time_stamps
for i in range(0,40):
    tsne_lat.append([kmeans.cluster_centers_[i][0]]*13099)
    tsne_lon.append([kmeans.cluster_centers_[i][1]]*13099)
    # jan 1st 2016 is thursday, so we start our day from 4: "(int(k/144))%7+4"
    # our prediction start from 5th 10min intravel since we need to have number of pickups that ar
e happened in last 5 pickup bins
    tsne_weekday.append([(int((int(k/144))%7+4)%7) for k in range(5,4464+4176+4464)])
    # regions_cum is a list of lists [[x1,x2,x3..x13104], [x1,x2,x3..x13104], [x1,x2,x3..x13104],
    [x1,x2,x3..x13104], [x1,x2,x3..x13104], .. 40 lsits]
    tsne_feature = np.vstack((tsne_feature, [regions_cum[i][r:r+number_of_time_stamps] for r in ran
ge(0,len(regions_cum[i])-number_of_time_stamps)]))
    output.append(regions_cum[i][5:])
tsne_feature = tsne_feature[1:]

```

In [30]:

```

len(tsne_lat[0])*len(tsne_lat) == tsne_feature.shape[0] == len(tsne_weekday)*len(tsne_weekday[0]) =
= 40*13099 == len(output)*len(output[0])

```

Out[30]:

True

In [31]:

```

# Getting the predictions of exponential moving averages to be used as a feature in cumulative for
m

# upto now we computed 8 features for every data point that starts from 50th min of the day
# 1. cluster center lattitude
# 2. cluster center longitude
# 3. day of the week
# 4. hour of the day
# 5. month of the year
# 6. day of the month
# 7. day of the year
# 8. week of the year

```

```

# 2. cluster center longitude
# 3. day of the week
# 4. f_t_1: number of pickups that are happened previous t-1th 10min intravel
# 5. f_t_2: number of pickups that are happened previous t-2th 10min intravel
# 6. f_t_3: number of pickups that are happened previous t-3th 10min intravel
# 7. f_t_4: number of pickups that are happened previous t-4th 10min intravel
# 8. f_t_5: number of pickups that are happened previous t-5th 10min intravel

# from the baseline models we said the exponential weighted moving avarage gives us the best error
# we will try to add the same exponential weighted moving avarage at t as a feature to our data
# exponential weighted moving avarage => p'(t) = alpha*p'(t-1) + (1-alpha)*P(t-1)
alpha=0.3

# it is a temporary array that store exponential weighted moving avarage for each 10min intravel,
# for each cluster it will get reset
# for every cluster it contains 13104 values
predicted_values=[]

predict_list = []
tsne_flat_exp_avg = []
for r in range(0,40):
    for i in range(0,13104):
        if i==0:
            predicted_value= regions_cum[r][0]
            predicted_values.append(0)
            continue
        predicted_values.append(predicted_value)
        predicted_value =int((alpha*predicted_value) + (1-alpha)*(regions_cum[r][i]))
    predict_list.append(predicted_values[5:])
    predicted_values=[]

```

## Fourier Features

In [32]:

```

amplitude = []
frequency = []
for i in range(40):
    amp = np.abs(np.fft.fft(regions_cum[i][0:13104])) #amplitude calculation
    freq = np.abs(np.fft.fftfreq(13104, 1)) #frequencies calculation
    amp_indices = np.argsort(-amp)[1:] #sorting amplitude
    amp_values = []
    freq_values = []
    for j in range(0, 9, 2): #taking top 5 amplitudes and frequencies
        amp_values.append(amp[amp_indices[j]])
        freq_values.append(freq[amp_indices[j]])
    for k in range(13104): #those top 5 frequencies and amplitudes are same for all the points
    in one cluster
        amplitude.append(amp_values)
        frequency.append(freq_values)

```

In [33]:

```

# train, test split : 70% 30% split
# Before we start predictions using the tree based regression models we take 3 months of 2016 pick
up data
# and split it such that for every region we have 70% data in train and 30% in test,
# ordered date-wise for every region
print("size of train data :", int(13099*0.7))
print("size of test data :", int(13099*0.3))

```

```

size of train data : 9169
size of test data : 3929

```

In [34]:

```

#for train and test data fourier frequencies
train_frequencies = [frequency[i*13099:(13099*i+9169)] for i in range(0,40)]
test_frequencies = [frequency[(13099*(i))+9169:13099*(i+1)] for i in range(0,40)]

```

In [35]:

```
#for train and test data fourier amplitudes
train_amplitudes = [amplitude[i*13099:(13099*i+9169)] for i in range(0,40)]
test_amplitudes = [amplitude[(13099*(i))+9169:13099*(i+1)] for i in range(0,40)]
```

In [36]:

```
# extracting first 9169 timestamp values i.e 70% of 13099 (total timestamps) for our training data
train_features = [tsne_feature[i*13099:(13099*i+9169)] for i in range(0,40)]
# temp = [0]*12955 - 9068
test_features = [tsne_feature[(13099*(i))+9169:13099*(i+1)] for i in range(0,40)]
```

In [37]:

```
print("Number of data clusters",len(train_features), "Number of data points in trian data",
      len(train_features[0]), "Each data point contains", len(train_features[0][0]),"features")
print("Number of data clusters",len(train_features), "Number of data points in test data",
      len(test_features[0]), "Each data point contains", len(test_features[0][0]),"features")
```

Number of data clusters 40 Number of data points in trian data 9169 Each data point contains 5 features

Number of data clusters 40 Number of data points in test data 3930 Each data point contains 5 features

In [38]:

```
# extracting first 9169 timestamp values i.e 70% of 13099 (total timestamps) for our training data
tsne_train_flat_lat = [i[:9169] for i in tsne_lat]
tsne_train_flat_lon = [i[:9169] for i in tsne_lon]
tsne_train_flat_weekday = [i[:9169] for i in tsne_weekday]
tsne_train_flat_output = [i[:9169] for i in output]
tsne_train_flat_exp_avg = [i[:9169] for i in predict_list]
```

In [39]:

```
# extracting the rest of the timestamp values i.e 30% of 12956 (total timestamps) for our test dat
a
tsne_test_flat_lat = [i[9169:] for i in tsne_lat]
tsne_test_flat_lon = [i[9169:] for i in tsne_lon]
tsne_test_flat_weekday = [i[9169:] for i in tsne_weekday]
tsne_test_flat_output = [i[9169:] for i in output]
tsne_test_flat_exp_avg = [i[9169:] for i in predict_list]
```

In [40]:

```
# the above contains values in the form of list of lists (i.e. list of values of each region), here we make all of them in one list
train_new_features = []
for i in range(0,40):
    train_new_features.extend(train_features[i])
test_new_features = []
for i in range(0,40):
    test_new_features.extend(test_features[i])
train_freq=[]
test_freq=[]
train_amp=[]
test_amp=[]
for i in range(0,40):
    train_freq.extend(train_frequencies[i])
    test_freq.extend(test_frequencies[i])
    train_amp.extend(train_amplitudes[i])
    test_amp.extend(test_amplitudes[i])
```

In [41]:

```
#stacking new features,frequencies and amplitudes horizontally using hstack.
train_brand_new_features=np.hstack((train_new_features,train_freq,train_amp))
test_brand_new_features=np.hstack((test_new_features,test_freq,test_amp))
```

....

In [42]:

```
# converting lists of lists into sinle list i.e flatten
# a = [[1,2,3,4],[4,6,7,8]]
# print(sum(a,[]))
# [1, 2, 3, 4, 4, 6, 7, 8]

tsne_train_lat = sum(tsne_train_flat_lat, [])
tsne_train_lon = sum(tsne_train_flat_lon, [])
tsne_train_weekday = sum(tsne_train_flat_weekday, [])
tsne_train_output = sum(tsne_train_flat_output, [])
tsne_train_exp_avg = sum(tsne_train_flat_exp_avg, [])
```

In [43]:

```
# converting lists of lists into sinle list i.e flatten
# a = [[1,2,3,4],[4,6,7,8]]
# print(sum(a,[]))
# [1, 2, 3, 4, 4, 6, 7, 8]

tsne_test_lat = sum(tsne_test_flat_lat, [])
tsne_test_lon = sum(tsne_test_flat_lon, [])
tsne_test_weekday = sum(tsne_test_flat_weekday, [])
tsne_test_output = sum(tsne_test_flat_output, [])
tsne_test_exp_avg = sum(tsne_test_flat_exp_avg, [])
```

In [44]:

```
# Preparing the data frame for our train data
columns = ['ft_5','ft_4','ft_3','ft_2','ft_1']
df_train = pd.DataFrame(data=train_new_features, columns=columns)
df_train['lat'] = tsne_train_lat
df_train['lon'] = tsne_train_lon
df_train['weekday'] = tsne_train_weekday
df_train['exp_avg'] = tsne_train_exp_avg

print(df_train.shape)
```

(366760, 9)

In [45]:

```
# Preparing the data frame for our test data
df_test = pd.DataFrame(data=test_new_features, columns=columns)
df_test['lat'] = tsne_test_lat
df_test['lon'] = tsne_test_lon
df_test['weekday'] = tsne_test_weekday
df_test['exp_avg'] = tsne_test_exp_avg
print(df_test.shape)
```

(157200, 9)

In [46]:

```
# Preparing the data frame for our train data
columns = ['ft_5','ft_4','ft_3','ft_2','ft_1','fre1','fre2','fre3','fre4','fre5','amp1','amp2','amp3','amp4','amp5']
df_train_fr = pd.DataFrame(data=train_brand_new_features, columns=columns)
df_train_fr['lat'] = tsne_train_lat
df_train_fr['lon'] = tsne_train_lon
df_train_fr['weekday'] = tsne_train_weekday
df_train_fr['exp_avg'] = tsne_train_exp_avg

print(df_train_fr.shape)
```

(366760, 19)

In [47]:

```
# Preparing the data frame for our test data
df_test_fr = pd.DataFrame(data=test_brand_new_features, columns=columns)
```

```

df_test_fr = pd.DataFrame(data=tsne_pca_new_features, columns=columns)
df_test_fr['lat'] = tsne_test_lat
df_test_fr['lon'] = tsne_test_lon
df_test_fr['weekday'] = tsne_test_weekday
df_test_fr['exp_avg'] = tsne_test_exp_avg
print(df_test_fr.shape)

```

(157200, 19)

In [65]:

```
df_train_fr.columns
```

Out[65]:

```

Index(['ft_5', 'ft_4', 'ft_3', 'ft_2', 'ft_1', 'fre1', 'fre2', 'fre3', 'fre4',
       'fre5', 'amp1', 'amp2', 'amp3', 'amp4', 'amp5', 'lat', 'lon', 'weekday',
       'exp_avg'],
      dtype='object')

```

## Triple Exponential Smoothing

**Exponential smoothing is a time series forecasting method for univariate data.**

In [48]:

```

#Reference is taken from this link :-- https://grisha.org/blog/2016/02/17/triple-exponential-smoothing-forecasting-part-iii/
#Triple Exponential Smoothing a.k.a Holt-Winters Method
#Exponential smoothing is a time series forecasting method for univariate data.
def initial_trend(series, slen): #computing the average of trend averages across seasons
    sum = 0.0
    for i in range(slen):
        sum += float(series[i+slen] - series[i]) / slen
    return sum / slen

#initial seasonal components
def initial_seasonal_components(series, slen):
    seasonals = {}
    season_averages = []
    n_seasons = int(len(series)/slen)
    # compute season averages
    for j in range(n_seasons):
        season_averages.append(sum(series[slen*j:slen*j+slen])/float(slen))
    # compute initial values
    for i in range(slen):
        sum_of_vals_over_avg = 0.0
        for j in range(n_seasons):
            sum_of_vals_over_avg += series[slen*j+i]-season_averages[j]
        seasonals[i] = sum_of_vals_over_avg/n_seasons
    return seasonals

#triple exponential smoothing
def triple_exponential_smoothing(series, slen, alpha, beta, gamma, n_preds):
    result = []
    seasonals = initial_seasonal_components(series, slen)
    for i in range(len(series)+n_preds):
        if i == 0: # initial values
            smooth = series[0]
            trend = initial_trend(series, slen)
            result.append(series[0])
            continue
        if i >= len(series): # we are forecasting
            m = i - len(series) + 1
            result.append((smooth + m*trend) + seasonals[i%slen])
        else:
            val = series[i]
            last_smooth, smooth = smooth, alpha*(val-seasonals[i%slen]) + (1-alpha)*(smooth+trend)
            trend = beta * (smooth-last_smooth) + (1-beta)*trend
            seasonals[i%slen] = gamma*(val-smooth) + (1-gamma)*seasonals[i%slen]
            result.append(smooth+trend+seasonals[i%slen])
    return result

```

```
In [49]:
```

```
alpha,beta,gama,slen = 0.2,.2,0.3,24
triple_exp_smoothing = []
for i in range(len(train_features)):
    triple_exp_smoothing.append(triple_exponential_smoothing(regions_cum[i][0:13104], slen, alpha,
beta, gama, 0))
```

```
In [50]:
```

```
tsne_train_triple_exp_smooth = [i[:9169] for i in triple_exp_smoothing]
tsne_test_triple_exp_smooth = [i[9169:] for i in triple_exp_smoothing]
```

```
In [51]:
```

```
#flattening
train_triple_exp_smooth=sum(tsne_train_triple_exp_smooth,[])
test_triple_exp_smooth=sum(tsne_test_triple_exp_smooth,[])
```

```
In [66]:
```

```
df_train_fr['triple_exp_smooth']=train_triple_exp_smooth
df_test_fr['triple_exp_smooth']=test_triple_exp_smooth[:157200]
```

```
In [67]:
```

```
df_train_fr.columns
```

```
Out[67]:
```

```
Index(['ft_5', 'ft_4', 'ft_3', 'ft_2', 'ft_1', 'fre1', 'fre2', 'fre3', 'fre4',
       'fre5', 'amp1', 'amp2', 'amp3', 'amp4', 'amp5', 'lat', 'lon', 'weekday',
       'exp_avg', 'triple_exp_smooth'],
      dtype='object')
```

```
In [68]:
```

```
df_test_fr.head()
```

```
Out[68]:
```

	ft_5	ft_4	ft_3	ft_2	ft_1	fre1	fre2	fre3	fre4	fre5	amp1	amp2	amp3	
0	118.0	106.0	104.0	93.0	102.0	0.006944	0.013889	0.012897	0.034722	0.007937	366274.780935	187964.6817	82956.377856	681
1	106.0	104.0	93.0	102.0	101.0	0.006944	0.013889	0.012897	0.034722	0.007937	366274.780935	187964.6817	82956.377856	681
2	104.0	93.0	102.0	101.0	120.0	0.006944	0.013889	0.012897	0.034722	0.007937	366274.780935	187964.6817	82956.377856	681
3	93.0	102.0	101.0	120.0	131.0	0.006944	0.013889	0.012897	0.034722	0.007937	366274.780935	187964.6817	82956.377856	681
4	102.0	101.0	120.0	131.0	164.0	0.006944	0.013889	0.012897	0.034722	0.007937	366274.780935	187964.6817	82956.377856	681

```
In [76]:
```

```
#Mean
df_train_fr['mean'] = df_train[['ft_5', 'ft_4', 'ft_3', 'ft_2', 'ft_1']].mean(axis=1)
df_test_fr['mean'] = df_test[['ft_5', 'ft_4', 'ft_3', 'ft_2', 'ft_1']].mean(axis=1)
```

```
In [77]:
```

```
#Variance
df_train_fr['variance'] = df_train[['ft_5', 'ft_4', 'ft_3', 'ft_2', 'ft_1']].var(axis=1)
df_test_fr['variance'] = df_test[['ft_5', 'ft_4', 'ft_3', 'ft_2', 'ft_1']].var(axis=1)
```

In [78]:

```
#skewness
df_train_fr['skewness'] = df_train[['ft_5', 'ft_4', 'ft_3', 'ft_2', 'ft_1']].skew(axis=1)
df_test_fr['skewness'] = df_test[['ft_5', 'ft_4', 'ft_3', 'ft_2', 'ft_1']].skew(axis=1)
```

In [79]:

```
#kurtosis
df_train_fr['kurtosis'] = df_train[['ft_5', 'ft_4', 'ft_3', 'ft_2', 'ft_1']].kurtosis(axis=1)
df_test_fr['kurtosis'] = df_test[['ft_5', 'ft_4', 'ft_3', 'ft_2', 'ft_1']].kurtosis(axis=1)
```

In [80]:

```
df_train.columns
```

Out[80]:

```
Index(['ft_5', 'ft_4', 'ft_3', 'ft_2', 'ft_1', 'lat', 'lon', 'weekday',
       'exp_avg', 'triple_exp_smooth', 'mean', 'variance', 'skewness',
       'kurtosis'],
      dtype='object')
```

In [81]:

```
df_train_fr.head()
```

Out[81]:

	ft_5	ft_4	ft_3	ft_2	ft_1	fre1	fre2	fre3	fre4	fre5	...	amp5	lat	lon	weekday
0	0.0	63.0	217.0	189.0	137.0	0.006944	0.013889	0.012897	0.034722	0.007937	...	62368.956611	40.776228	-	73.982119
1	63.0	217.0	189.0	137.0	135.0	0.006944	0.013889	0.012897	0.034722	0.007937	...	62368.956611	40.776228	-	73.982119
2	217.0	189.0	137.0	135.0	129.0	0.006944	0.013889	0.012897	0.034722	0.007937	...	62368.956611	40.776228	-	73.982119
3	189.0	137.0	135.0	129.0	150.0	0.006944	0.013889	0.012897	0.034722	0.007937	...	62368.956611	40.776228	-	73.982119
4	137.0	135.0	129.0	150.0	164.0	0.006944	0.013889	0.012897	0.034722	0.007937	...	62368.956611	40.776228	-	73.982119

5 rows × 24 columns

In [82]:

```
df_test_fr.head()
```

Out[82]:

	ft_5	ft_4	ft_3	ft_2	ft_1	fre1	fre2	fre3	fre4	fre5	...	amp5	lat	lon	weekday
0	118.0	106.0	104.0	93.0	102.0	0.006944	0.013889	0.012897	0.034722	0.007937	...	62368.956611	40.776228	-	73.982119
1	106.0	104.0	93.0	102.0	101.0	0.006944	0.013889	0.012897	0.034722	0.007937	...	62368.956611	40.776228	-	73.982119
2	104.0	93.0	102.0	101.0	120.0	0.006944	0.013889	0.012897	0.034722	0.007937	...	62368.956611	40.776228	-	73.982119
3	93.0	102.0	101.0	120.0	131.0	0.006944	0.013889	0.012897	0.034722	0.007937	...	62368.956611	40.776228	-	73.982119
4	102.0	101.0	120.0	131.0	164.0	0.006944	0.013889	0.012897	0.034722	0.007937	...	62368.956611	40.776228	-	73.982119

5 rows × 24 columns

## Linear Regression

In [84]:

```
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import GridSearchCV

parameters = {'fit_intercept':[True, False], 'normalize':[True, False]}

model = LinearRegression(n_jobs = -1)
lr = GridSearchCV(model, parameters, scoring = 'neg_mean_absolute_error', cv = 3)
lr.fit(df_train_fr, tsne_train_output)

print("Best Parameters", lr.best_params_)

lr_pred_test = lr.predict(df_test_fr)
lr_test_predictions = [round(value) for value in lr_pred_test]
lr_pred_train = lr.predict(df_train_fr)
lr_train_predictions = [round(value) for value in lr_pred_train]
```

Best Parameters {'fit\_intercept': True, 'normalize': True}

In [85]:

```
train_MAPE_lr = mean_absolute_error(tsne_train_output, lr_pred_train)/ (sum(tsne_train_output)/len(tsne_train_output))
train_MSE_lr = mean_squared_error(tsne_train_output, lr_pred_train)

test_MAPE_lr = mean_absolute_error(tsne_test_output, lr_pred_test)/ (sum(tsne_test_output)/len(tsne_test_output))
test_MSE_lr = mean_squared_error(tsne_test_output, lr_pred_test)
```

In [86]:

```
print('Mean absolute percentage train error for Linear Regressor =',train_MAPE_lr)
print('Mean absolute percentage test error for Linear Regressor =',test_MAPE_lr)
print('---'*50)
print('Mean Squared train error for Linear Regressor =',train_MSE_lr)
print('Mean Squared test error for Linear Regressor=',test_MSE_lr)
```

```
Mean absolute percentage train error for Linear Regressor = 0.1331421875328096
Mean absolute percentage test error for Linear Regressor = 0.14242584380608697
-----
Mean Squared train error for Linear Regressor = 164.23440561474337
Mean Squared test error for Linear Regressor= 207.24187276437382
```

## Random Forest Regressor

In [87]:

```
# Training a hyper-parameter tuned random forest regressor on our train data
# find more about LinearRegression function here http://scikit-learn.org/stable/modules/generated/sklearn.ensemble.RandomForestRegressor.html
# -----
# default parameters
# sklearn.ensemble.RandomForestRegressor(n_estimators=10, criterion='mse', max_depth=None, min_samples_split=2,
# min_samples_leaf=1, min_weight_fraction_leaf=0.0, max_features='auto', max_leaf_nodes=None, min_impurity_decrease=0.0,
# min_impurity_split=None, bootstrap=True, oob_score=False, n_jobs=1, random_state=None,
# verbose=0, warm_start=False)

# some of methods of RandomForestRegressor()
# apply(X) Apply trees in the forest to X, return leaf indices.
# decision_path(X) Return the decision path in the forest
# fit(X, y[, sample_weight]) Build a forest of trees from the training set (X, y).
# get_params([deep]) Get parameters for this estimator.
# predict(X) Predict regression target for X.
# score(X, y[, sample_weight]) Returns the coefficient of determination R^2 of the prediction.
# -----
# video link1: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/regression-
```

```
" video link1: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/decision-trees-2/
# video link2: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/what-are-ensembles/
# ----

from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import RandomizedSearchCV
from scipy.stats import randint as sp_randint
from scipy.stats import uniform

parameters = {"n_estimators": [50, 100, 150, 200, 300, 400, 450],
              "max_depth": [2, 5, 7, 10, 15, 18, 20, 25],
              "min_samples_split": sp_randint(120, 190),
              "min_samples_leaf": sp_randint(25, 65)
             }

clf = RandomForestRegressor(random_state=25, n_jobs=-1)
model = RandomizedSearchCV(clf, param_distributions=parameters,
                           n_iter=5, cv=3, scoring='neg_mean_absolute_error', random_state=25)
model.fit(df_train_fr, tsne_train_output)
print(model.best_estimator_)

RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=15,
                     max_features='auto', max_leaf_nodes=None,
                     min_impurity_decrease=0.0, min_impurity_split=None,
                     min_samples_leaf=51, min_samples_split=135,
                     min_weight_fraction_leaf=0.0, n_estimators=400, n_jobs=-1,
                     oob_score=False, random_state=25, verbose=0,
                     warm_start=False)
```

In [88]:

```
model = RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=15,
                             max_features='auto', max_leaf_nodes=None,
                             min_impurity_decrease=0.0, min_impurity_split=None,
                             min_samples_leaf=51, min_samples_split=135,
                             min_weight_fraction_leaf=0.0, n_estimators=400, n_jobs=-1,
                             oob_score=False, random_state=25, verbose=0,
                             warm_start=False)
model.fit(df_train_fr, tsne_train_output)
```

Out[88]:

```
RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=15,
                     max_features='auto', max_leaf_nodes=None,
                     min_impurity_decrease=0.0, min_impurity_split=None,
                     min_samples_leaf=51, min_samples_split=135,
                     min_weight_fraction_leaf=0.0, n_estimators=400, n_jobs=-1,
                     oob_score=False, random_state=25, verbose=0,
                     warm_start=False)
```

In [89]:

```
rf_pred_test = model.predict(df_test_fr)
rndf_test_predictions = [round(value) for value in rf_pred_test]
rf_pred_train = model.predict(df_train_fr)
rndf_train_predictions = [round(value) for value in rf_pred_train]
```

In [90]:

```
train_MAPE_rf = mean_absolute_error(tsne_train_output, rf_pred_train) / (sum(tsne_train_output)/len(tsne_train_output))
train_MSE_rf = mean_squared_error(tsne_train_output, rf_pred_train)

test_MAPE_rf = mean_absolute_error(tsne_test_output, rf_pred_test) / (sum(tsne_test_output)/len(tsne_test_output))
test_MSE_rf = mean_squared_error(tsne_test_output, rf_pred_test)
```

In [91]:

```
print('Mean absolute percentage train error for Random Forest Regressor =', train_MAPE_rf)
print('Mean absolute percentage test error for Random Forest Regressor =', test_MAPE_rf)
```

```
print('Mean absolute percentage train error for Random Forest Regressor =',train_MSE_rf)
print('Mean Squared train error for Random Forest Regressor =',train_MSE_rf)
print('Mean Squared test error for Random Forest Regressor =',test_MSE_rf)
```

```
Mean absolute percentage train error for Random Forest Regressor = 0.12447930452348953
Mean absolute percentage test error for Random Forest Regressor = 0.13236543170553486
-----
Mean Squared train error for Random Forest Regressor = 145.87698615266189
Mean Squared test error for Random Forest Regressor = 185.76486404494585
```

## XGBoost Regressor

In [143]:

```
import warnings
warnings.filterwarnings("ignore")

import xgboost as xgb
clf = xgb.XGBRegressor()

parameters = {"n_estimators": [10, 20, 50, 100, 150, 200, 250, 300],
              "max_depth": [2, 5, 10, 20, 50, 100, 200, 250]
              }
model = RandomizedSearchCV(clf, param_distributions=parameters,
                           n_iter=5, cv=3, scoring='neg_mean_absolute_error', random_state=25)

model.fit(df_train_fr, tsne_train_output)
print(model.best_estimator_)
```

```
[18:21:18] WARNING: C:/Jenkins/workspace/xgboost-
win64_release_0.90/src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of
reg:squarederror.
[18:27:45] WARNING: C:/Jenkins/workspace/xgboost-
win64_release_0.90/src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of
reg:squarederror.
[18:34:27] WARNING: C:/Jenkins/workspace/xgboost-
win64_release_0.90/src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of
reg:squarederror.
[18:40:59] WARNING: C:/Jenkins/workspace/xgboost-
win64_release_0.90/src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of
reg:squarederror.
[18:41:02] WARNING: C:/Jenkins/workspace/xgboost-
win64_release_0.90/src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of
reg:squarederror.
[18:41:05] WARNING: C:/Jenkins/workspace/xgboost-
win64_release_0.90/src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of
reg:squarederror.
[18:41:08] WARNING: C:/Jenkins/workspace/xgboost-
win64_release_0.90/src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of
reg:squarederror.
[18:42:28] WARNING: C:/Jenkins/workspace/xgboost-
win64_release_0.90/src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of
reg:squarederror.
[18:43:50] WARNING: C:/Jenkins/workspace/xgboost-
win64_release_0.90/src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of
reg:squarederror.
[18:45:12] WARNING: C:/Jenkins/workspace/xgboost-
win64_release_0.90/src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of
reg:squarederror.
[18:46:12] WARNING: C:/Jenkins/workspace/xgboost-
win64_release_0.90/src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of
reg:squarederror.
[18:47:14] WARNING: C:/Jenkins/workspace/xgboost-
win64_release_0.90/src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of
reg:squarederror.
[18:48:15] WARNING: C:/Jenkins/workspace/xgboost-
win64_release_0.90/src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of
reg:squarederror.
[19:06:13] WARNING: C:/Jenkins/workspace/xgboost-
win64_release_0.90/src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of
reg:squarederror.
[19:24:29] WARNING: C:/Jenkins/workspace/xgboost-
win64_release_0.90/src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of
```

```
reg:squarederror.  
[19:42:37] WARNING: C:/Jenkins/workspace/xgboost-  
win64_release_0.90/src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of  
reg:squarederror.  
XGBRegressor(base_score=0.5, booster='gbtree', colsample_bylevel=1,  
             colsample_bynode=1, colsample_bytree=1, gamma=0,  
             importance_type='gain', learning_rate=0.1, max_delta_step=0,  
             max_depth=5, min_child_weight=1, missing=None, n_estimators=200,  
             n_jobs=1, nthread=None, objective='reg:linear', random_state=0,  
             reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,  
             silent=None, subsample=1, verbosity=1)
```

In [145]:

```
model = xgb.XGBRegressor(base_score=0.5, booster='gbtree', colsample_bylevel=1,  
                         colsample_bynode=1, colsample_bytree=1, gamma=0,  
                         importance_type='gain', learning_rate=0.1, max_delta_step=0,  
                         max_depth=5, min_child_weight=1, missing=None, n_estimators=200,  
                         n_jobs=1, nthread=None, objective='reg:linear', random_state=0,  
                         reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,  
                         silent=None, subsample=1, verbosity=1)  
model.fit(df_train_fr, tsne_train_output)
```

```
[19:46:52] WARNING: C:/Jenkins/workspace/xgboost-  
win64_release_0.90/src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of  
reg:squarederror.
```

Out [145]:

```
XGBRegressor(base_score=0.5, booster='gbtree', colsample_bylevel=1,  
             colsample_bynode=1, colsample_bytree=1, gamma=0,  
             importance_type='gain', learning_rate=0.1, max_delta_step=0,  
             max_depth=5, min_child_weight=1, missing=None, n_estimators=200,  
             n_jobs=1, nthread=None, objective='reg:linear', random_state=0,  
             reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,  
             silent=None, subsample=1, verbosity=1)
```

In [146]:

```
xg_pred_test = model.predict(df_test_fr)  
xgb_test_predictions = [round(value) for value in xg_pred_test]  
xg_pred_train = model.predict(df_train_fr)  
xgb_train_predictions = [round(value) for value in xg_pred_train]
```

In [147]:

```
train_MAPE_xgb = mean_absolute_error(tsne_train_output, xg_pred_train) / (sum(tsne_train_output)/len(tsne_train_output))  
train_MSE_xgb = mean_squared_error(tsne_train_output, xg_pred_train)  
  
test_MAPE_xgb = mean_absolute_error(tsne_test_output, xg_pred_test) / (sum(tsne_test_output)/len(tsne_test_output))  
test_MSE_xgb = mean_squared_error(tsne_test_output, xg_pred_test)
```

In [148]:

```
print('Mean absolute percentage train error for XGB Regressor =',train_MAPE_xgb)  
print('Mean absolute percentage test error for XGB Regressor =',test_MAPE_xgb)  
print('-'*50)  
print('Mean Squared train error for XGB Regressor =',train_MSE_xgb)  
print('Mean Squared test error for XGB Regressor =',test_MSE_xgb)
```

```
Mean absolute percentage train error for XGB Regressor = 0.1277173922558989  
Mean absolute percentage test error for XGB Regressor = 0.14267581646251726  
-----  
Mean Squared train error for XGB Regressor = 147.79842595660074  
Mean Squared test error for XGB Regressor = 218.44595716009079
```

## Prettytable

In [99]:

```
# Cluster size = 30
from prettytable import PrettyTable
x = PrettyTable()
x = PrettyTable(["Models", "MAPE train", "MAPE test"])
x.add_row(['Linear Regression', "0.119", "0.115"])
x.add_row(['Random Forest', "0.112", "0.113"])
x.add_row(['XGBOOST', "0.118", "0.115"])
print(x)
```

Models	MAPE train	MAPE test
Linear Regression	0.119	0.115
Random Forest	0.112	0.113
XGBOOST	0.118	0.115

In [97]:

```
# Cluster size = 40
from prettytable import PrettyTable
x = PrettyTable()
x = PrettyTable(["Models", "MAPE train", "MAPE test"])
x.add_row(['Linear Regression', train_MAPE_lr, test_MAPE_lr])
x.add_row(['Random Forest', train_MAPE_rf, test_MAPE_rf])
x.add_row(['XGBOOST', '0.127', "0.142"])
print(x)
```

Models	MAPE train	MAPE test
Linear Regression	0.1331421875328096	0.14242584380608697
Random Forest	0.12447930452348953	0.13236543170553486
XGBOOST	0.127	0.142

## Conclusions

1. We have used two cluster sizes 30 and 40.
2. First we prepare models for cluster size = 40. We added fourier features into our dataframe and after hyperparameter tuning, we calculated MAPE value which was > 12%
3. Then we added triple exponential smoothing features into our dataframe to decrease the value of MAPE, but MAPE value was still greater than 12%
4. Then we choose cluster size = 30 in which average number of clusters within vicinity are 26% which is greater than cluster size =40 and also inter-cluster distance is 0.51 which is almost same as cluster size =40.
5. So we prepared data and calculated MAPE values for Linear regression, random forest and xgboost models when cluster size = 30 and we decreased the value of MAPE which is < 12%
6. Random Forest model has best MAPE value which is 11.30%

In [ ]: